

# Event Streaming

With

# Apache Kafka

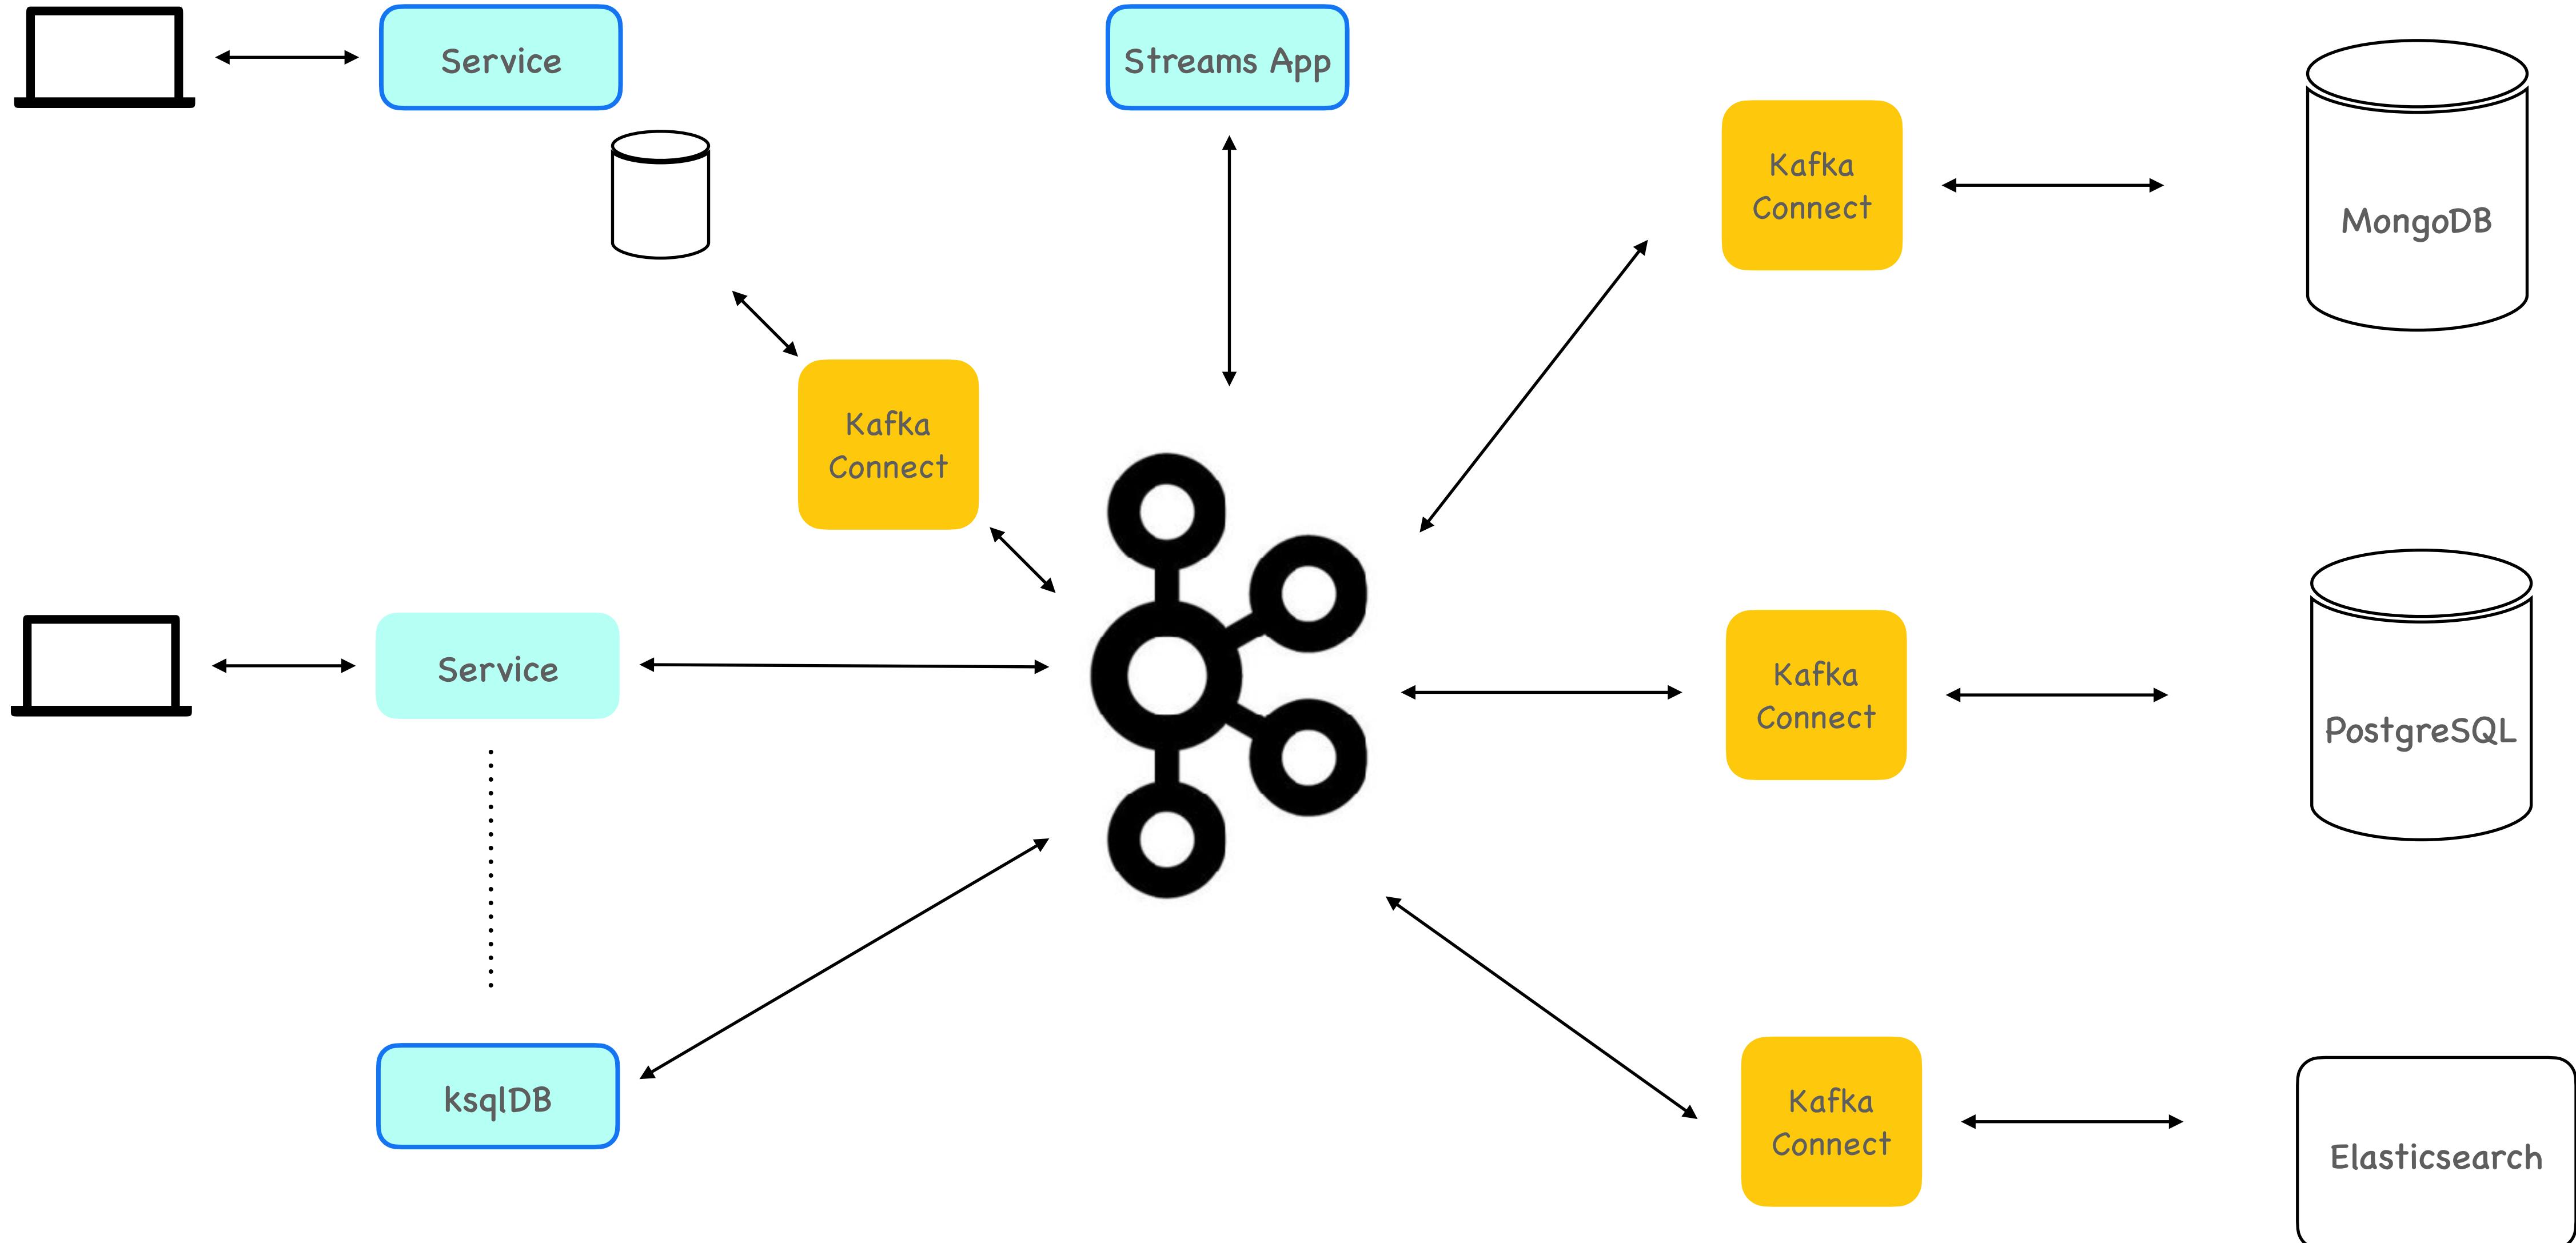
[dklein@confluent.io](mailto:dklein@confluent.io)

@daveklein

[linkedin.com/in/daveklein19](https://linkedin.com/in/daveklein19)

# What is Apache Kafka?

# Event Streaming Platform



# Event

## Notification

Order Placed

Temperature Read

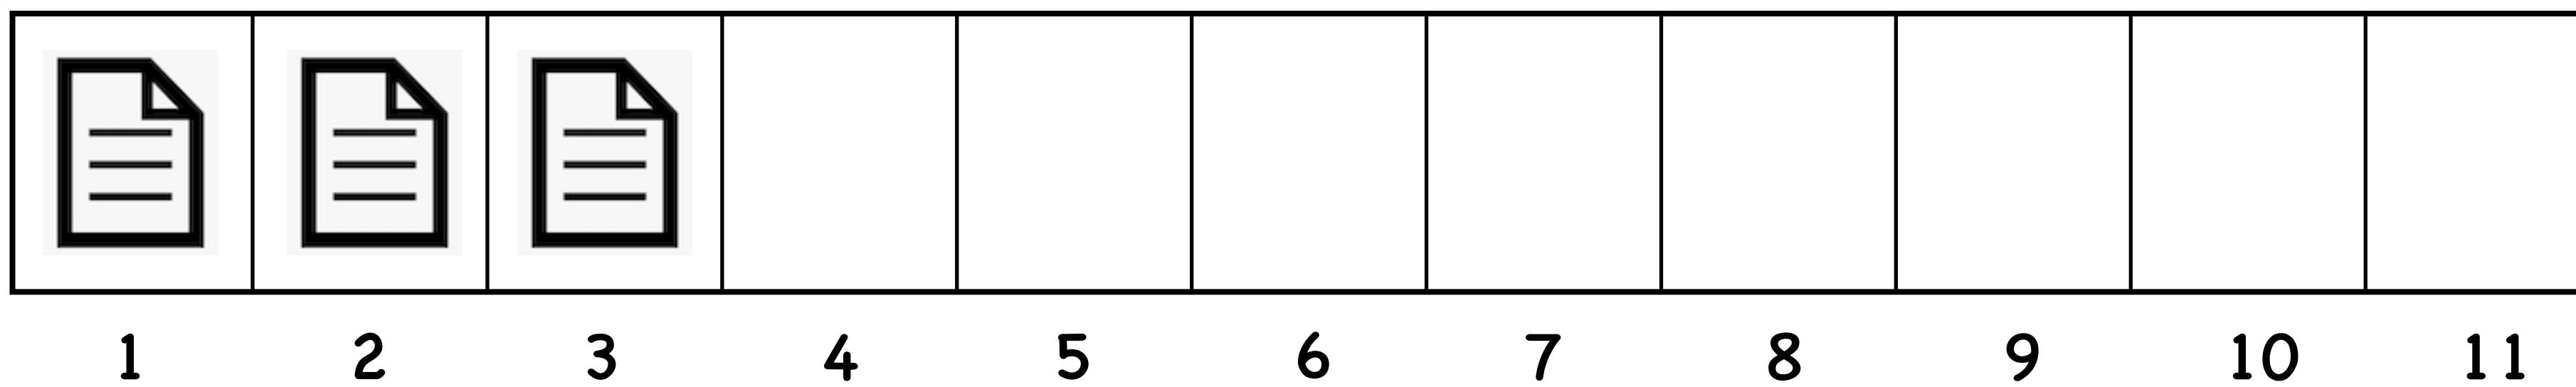
## State

{item:123, price: 29.95,  
qty: 2}

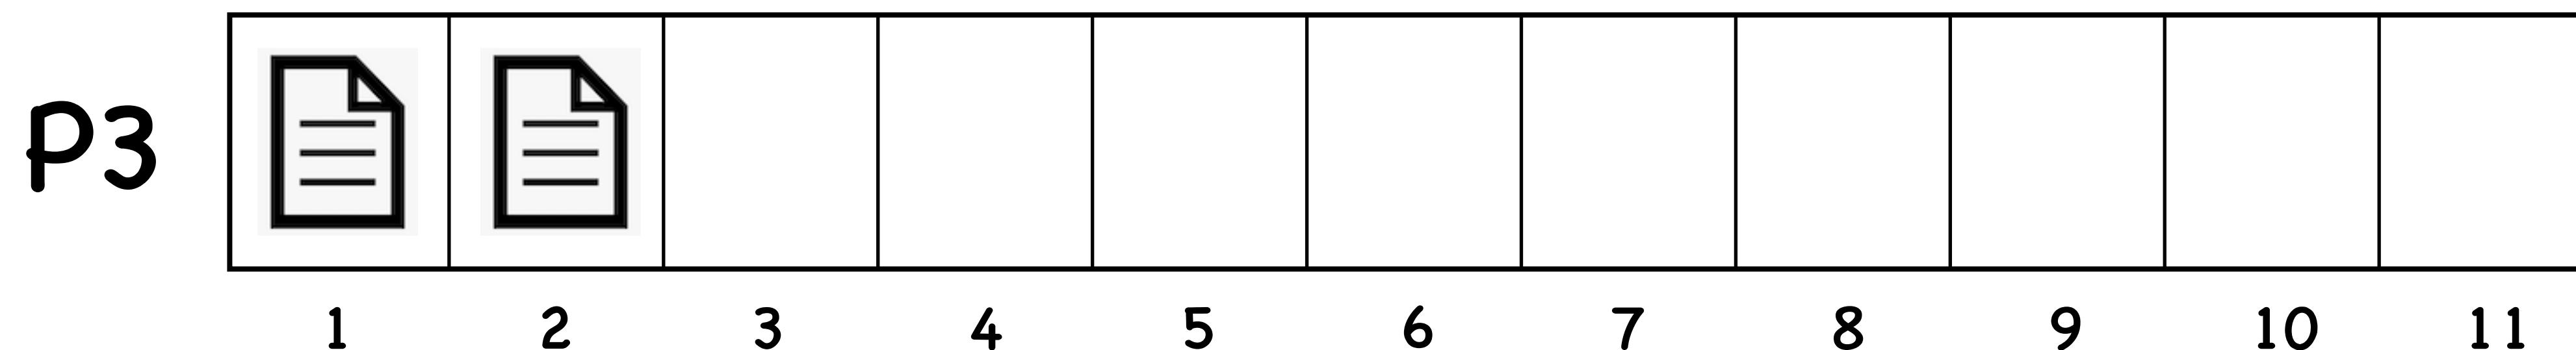
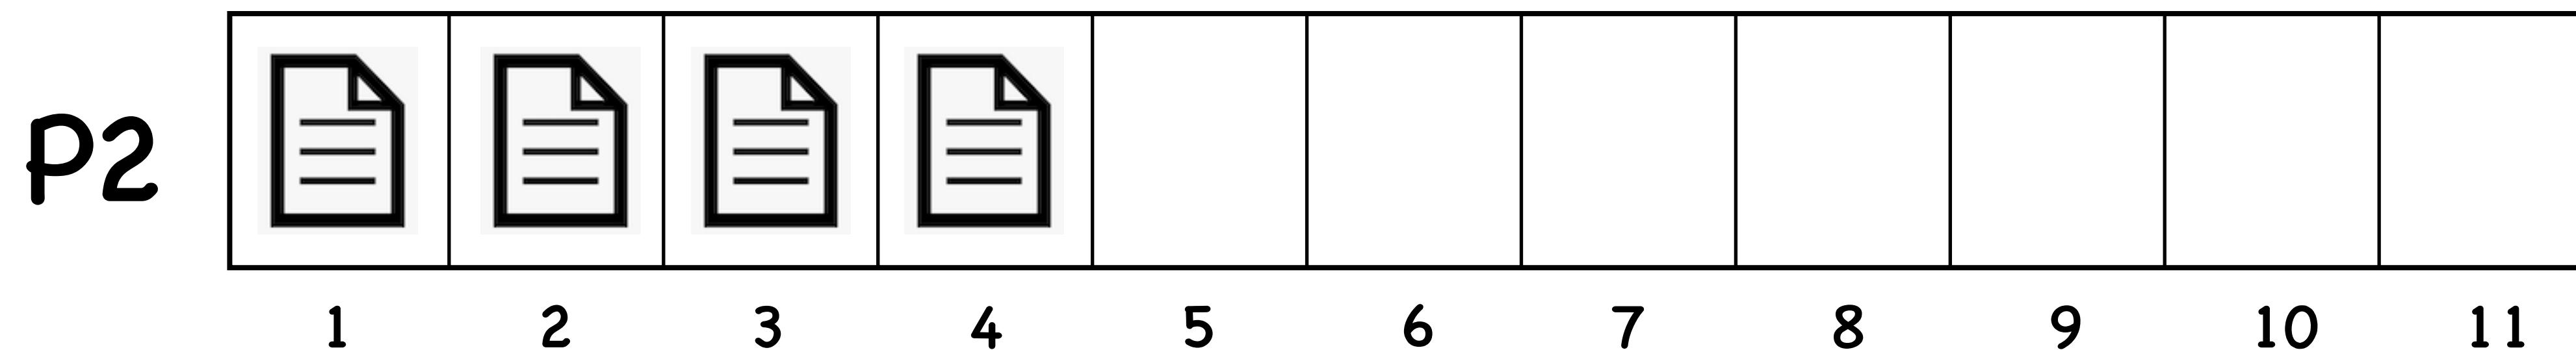
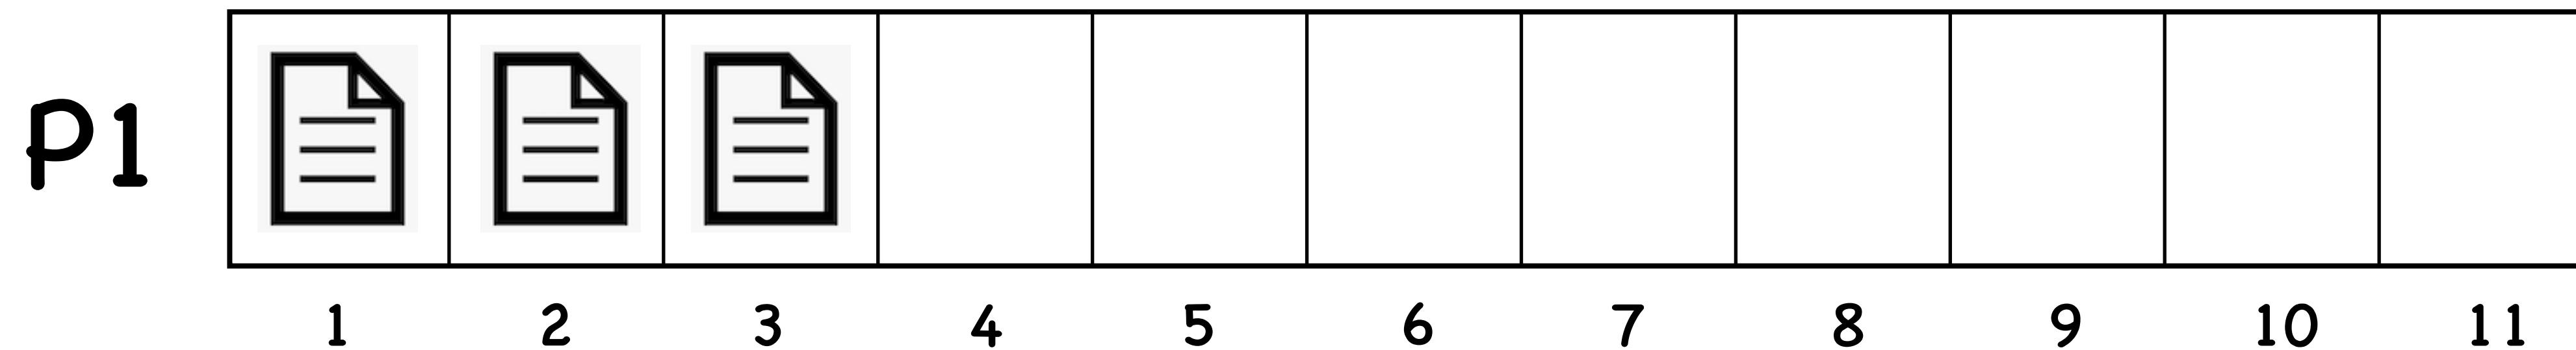
{temp:29, unit: F,  
time: 1606949836369}

```
"key": "a1",  
"value": {  
    "eventType": "added-to-cart",  
    "title": "Kafka Streams in Action",  
    "author": "Bill Bejeck",  
    "price": 44.99  
}
```

# Topic (log)

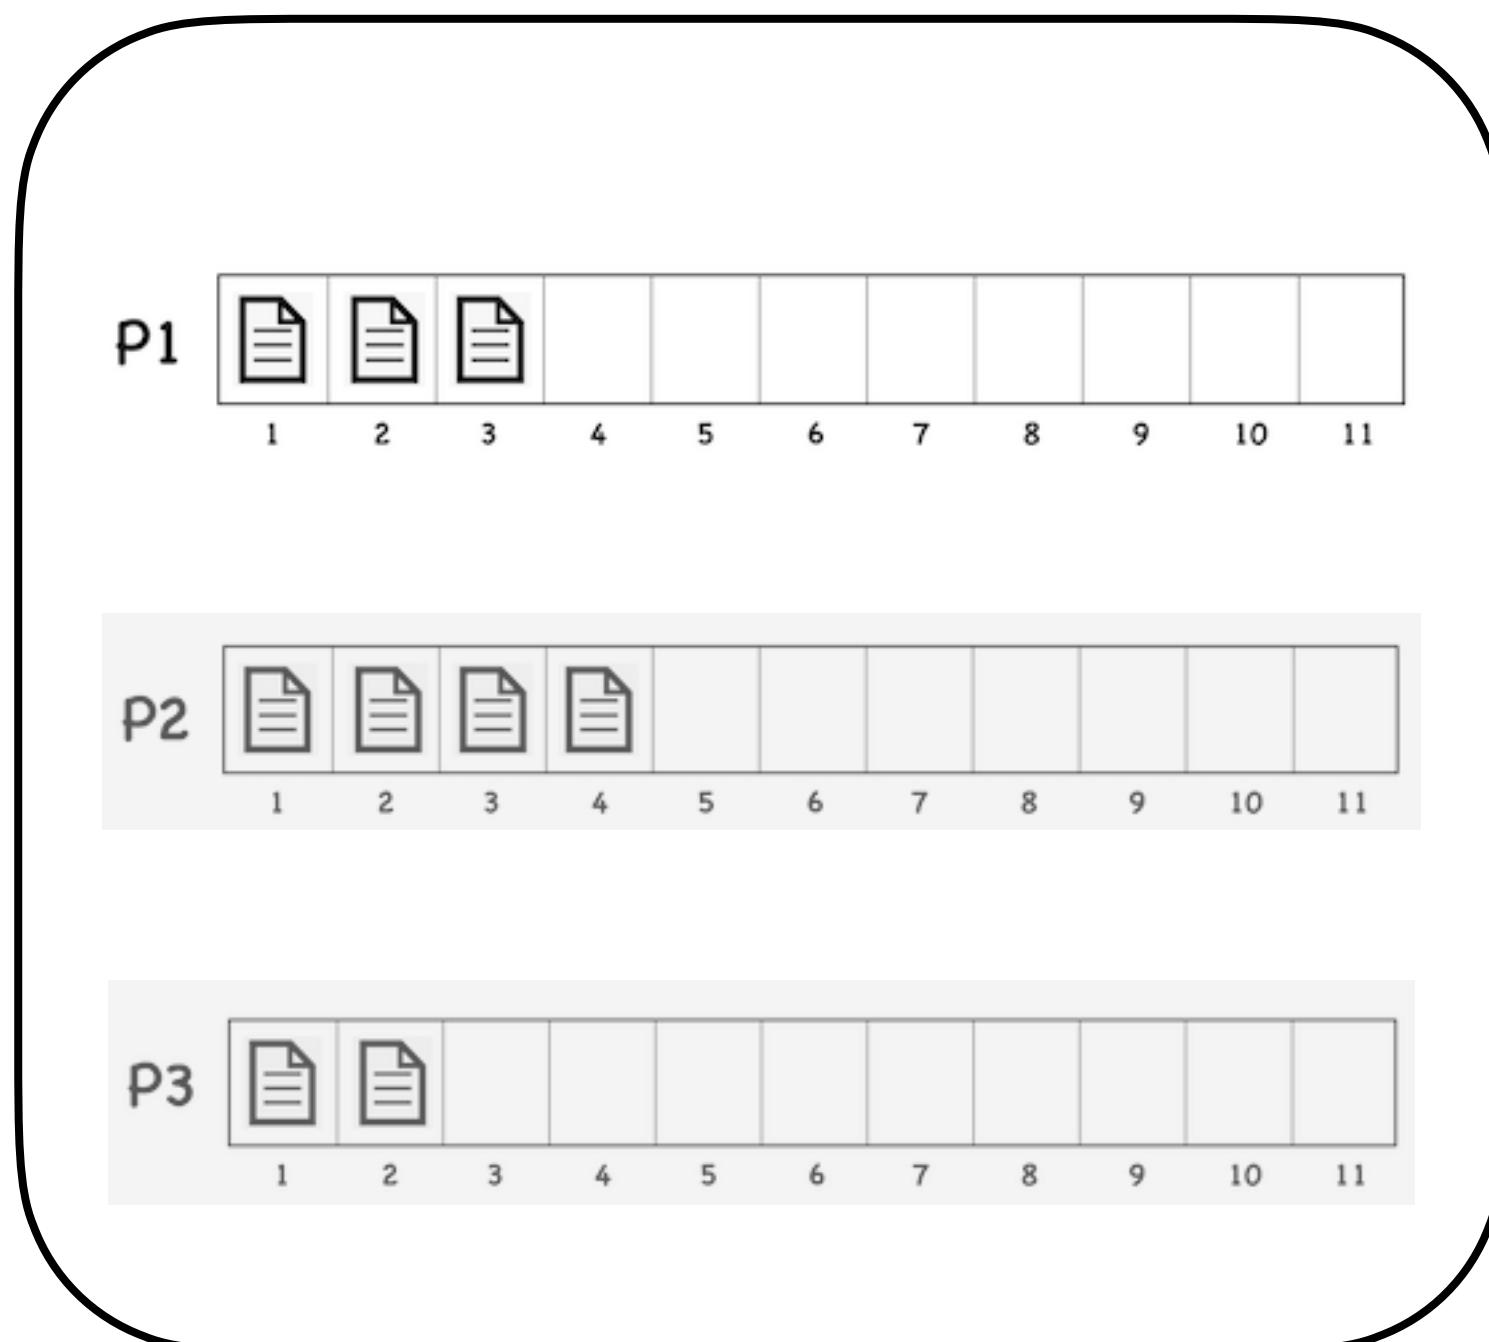


# Topic Partitions

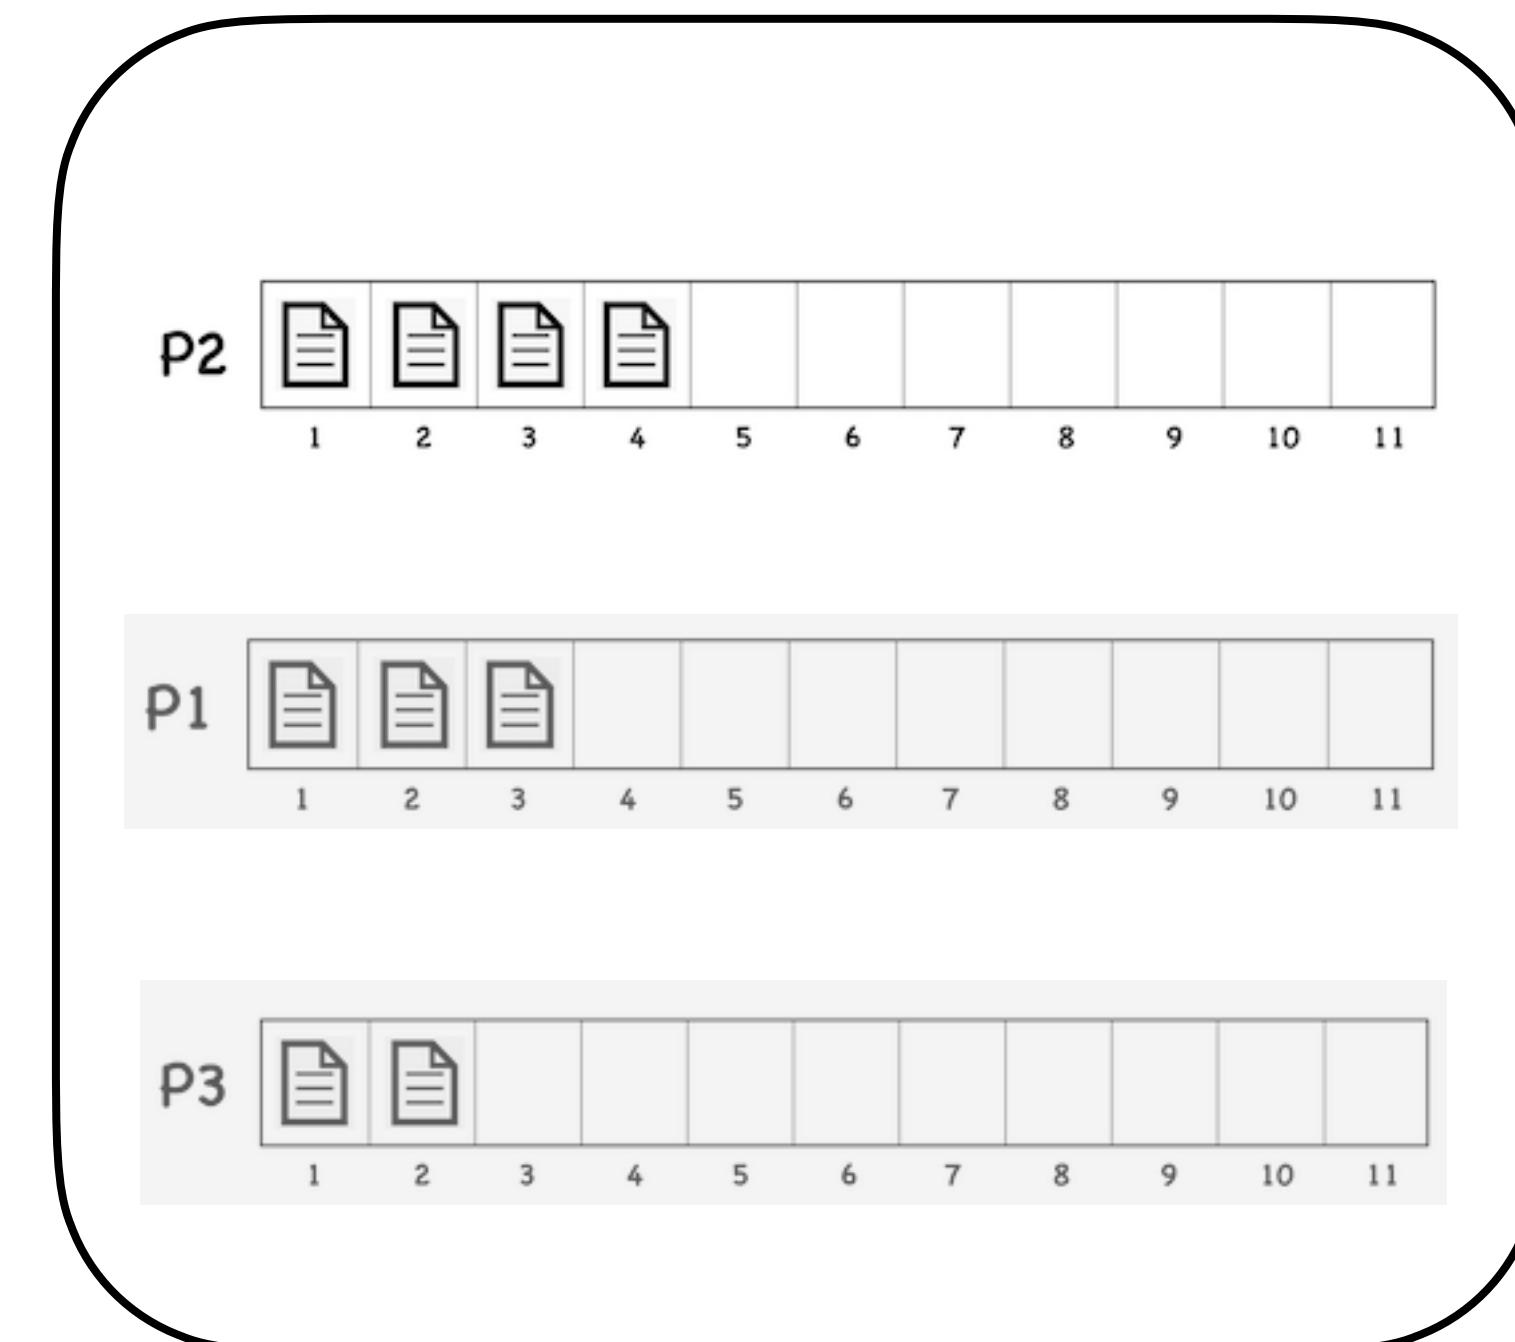


# Brokers and Replication

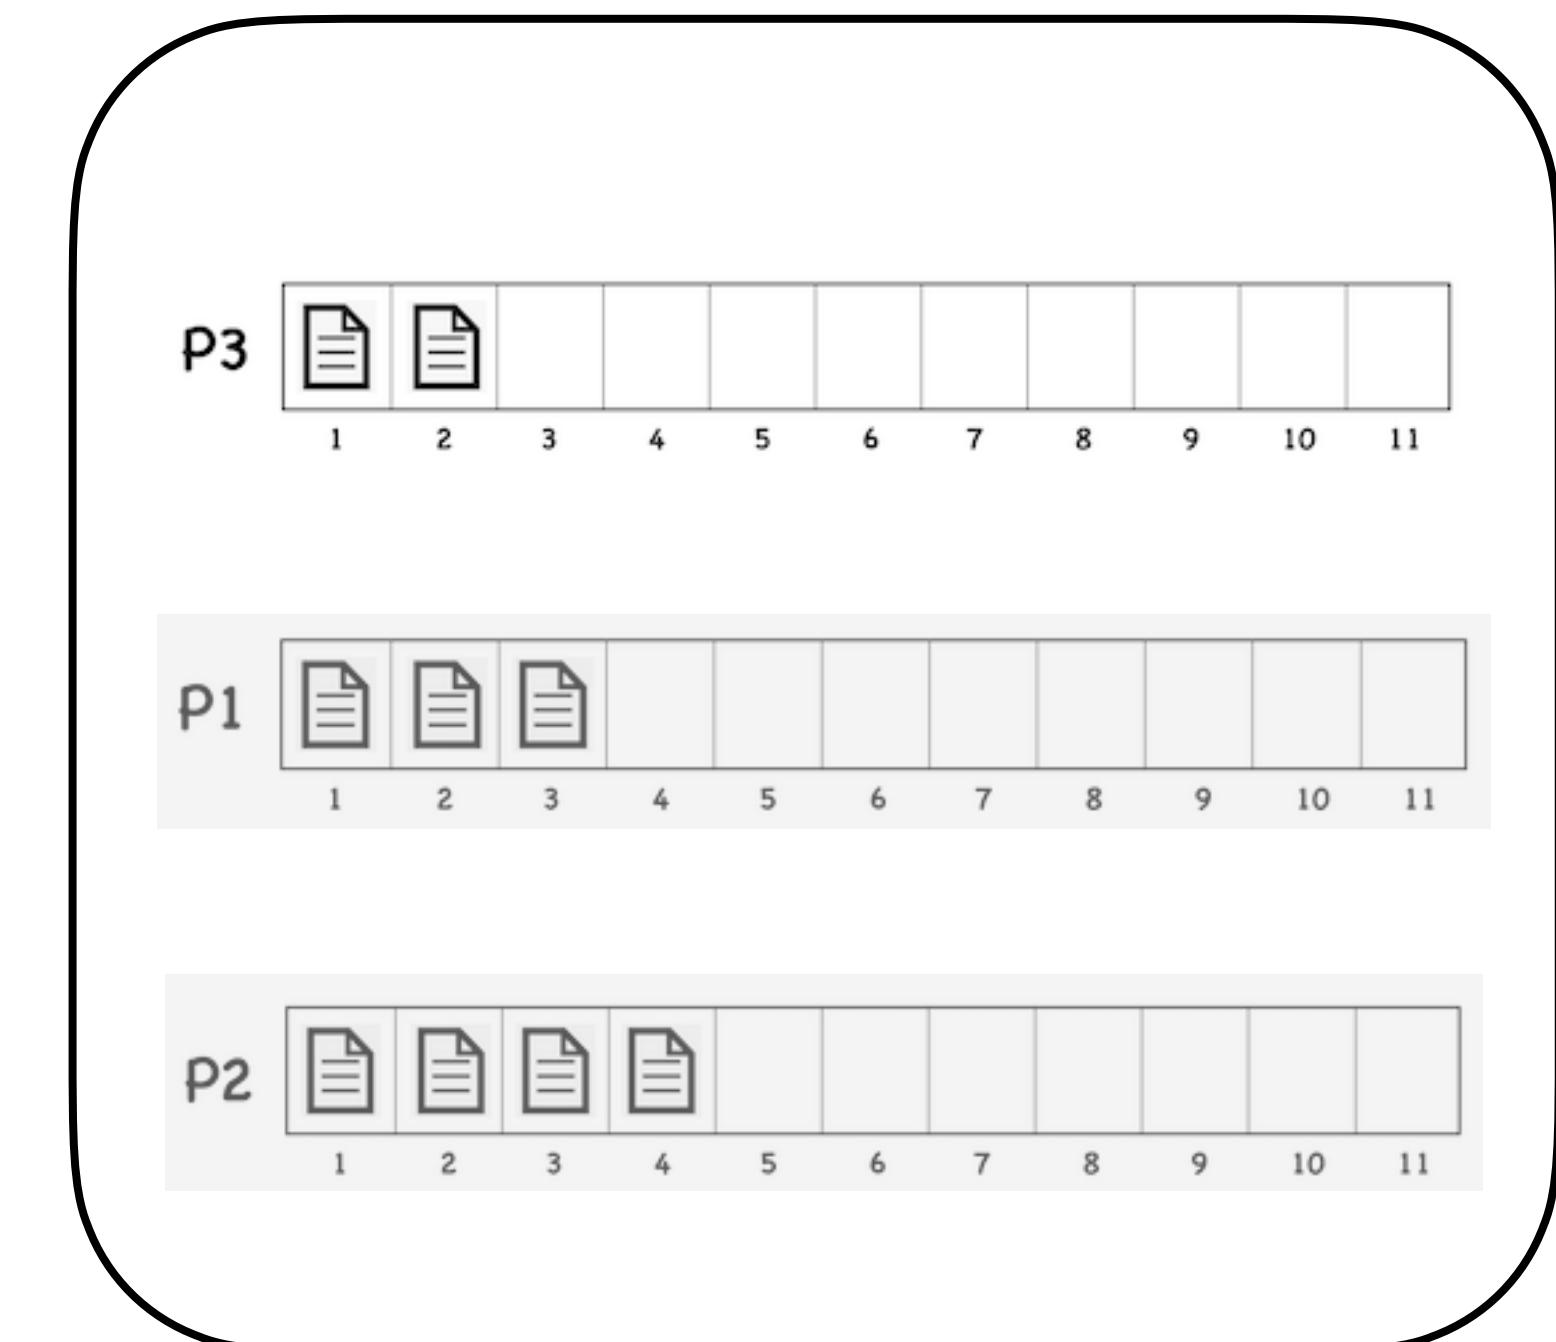
Broker 1



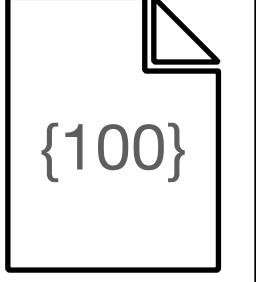
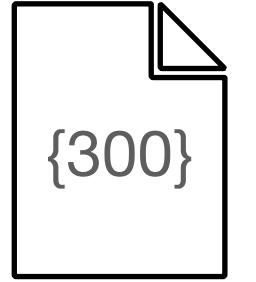
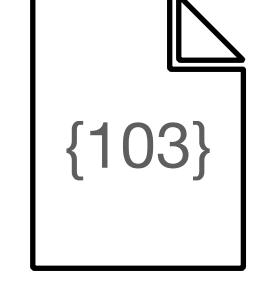
Broker 2



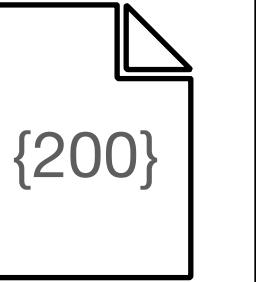
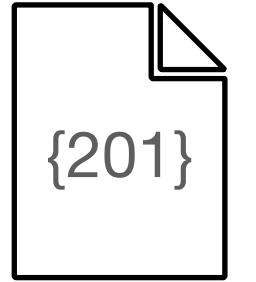
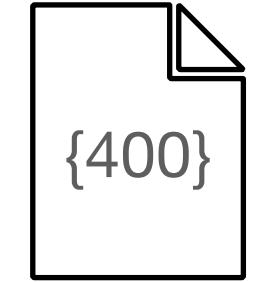
Broker 3



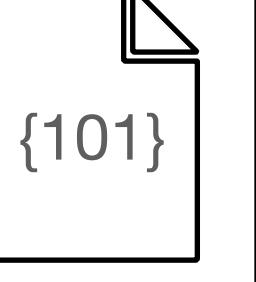
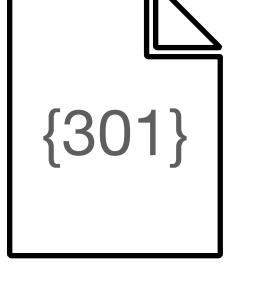
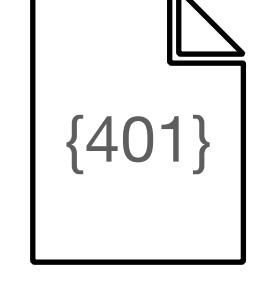
# Topic

 {100}	 {300}	 {103}								
--	---	---	--	--	--	--	--	--	--	--

1 2 3 4 5 6 7 8 9 10 11

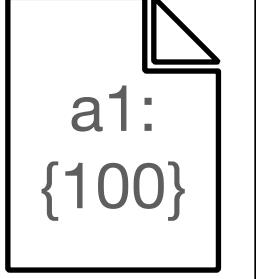
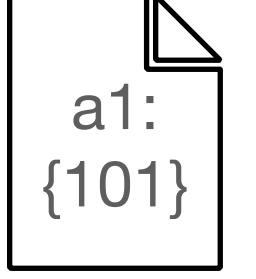
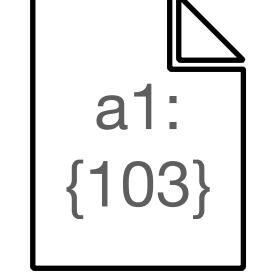
 {200}	 {201}	 {400}								
--	---	---	--	--	--	--	--	--	--	--

1 2 3 4 5 6 7 8 9 10 11

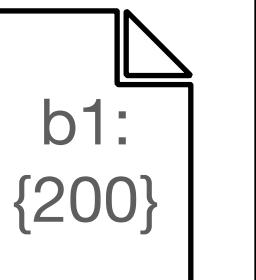
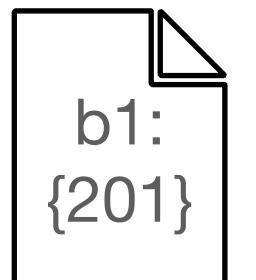
 {101}	 {301}	 {401}								
--	---	---	--	--	--	--	--	--	--	--

1 2 3 4 5 6 7 8 9 10 11

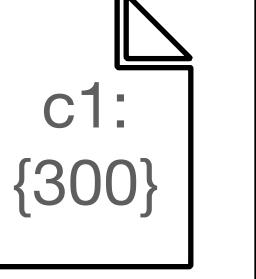
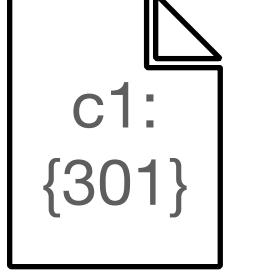
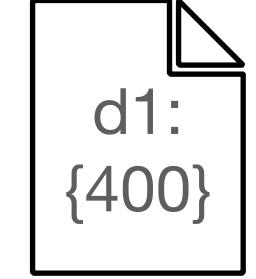
# Topic

 a1: {100}	 a1: {101}	 a1: {103}								
---	--	--	--	--	--	--	--	--	--	--

1 2 3 4 5 6 7 8 9 10 11

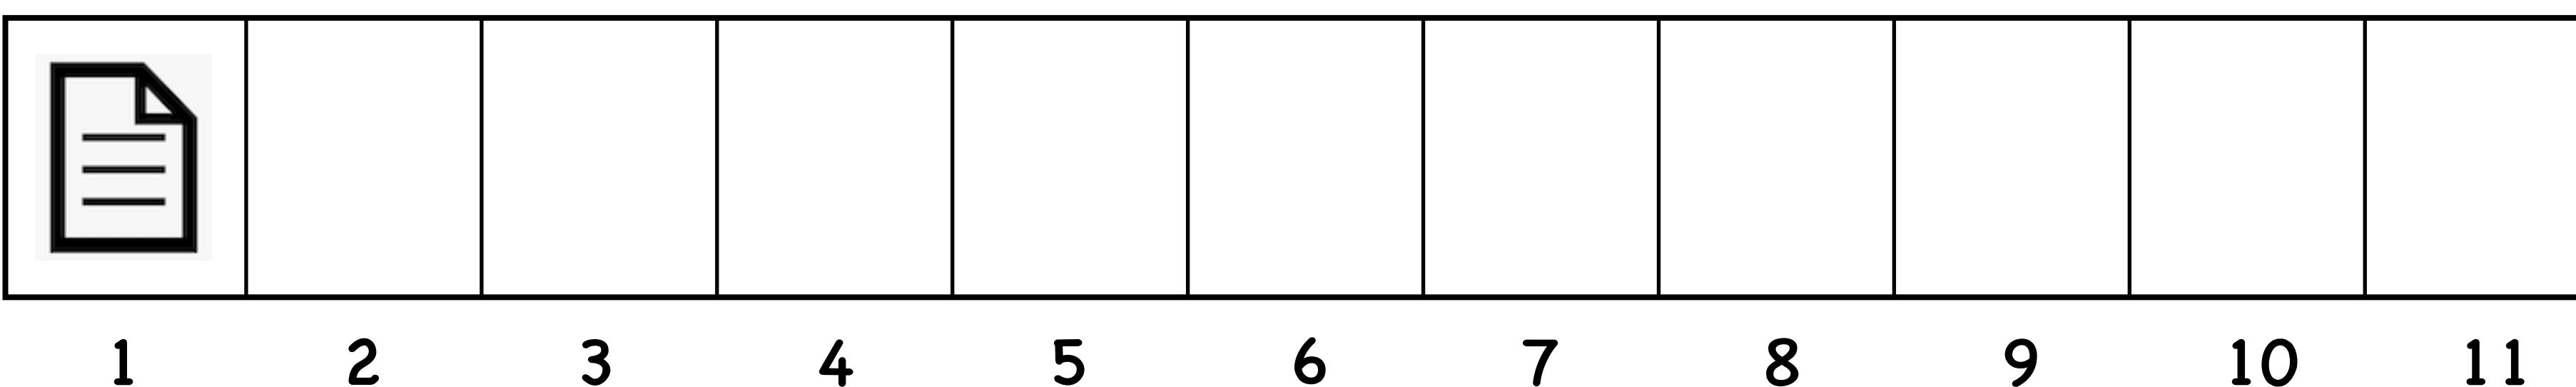
 b1: {200}	 b1: {201}									
---	--	--	--	--	--	--	--	--	--	--

1 2 3 4 5 6 7 8 9 10 11

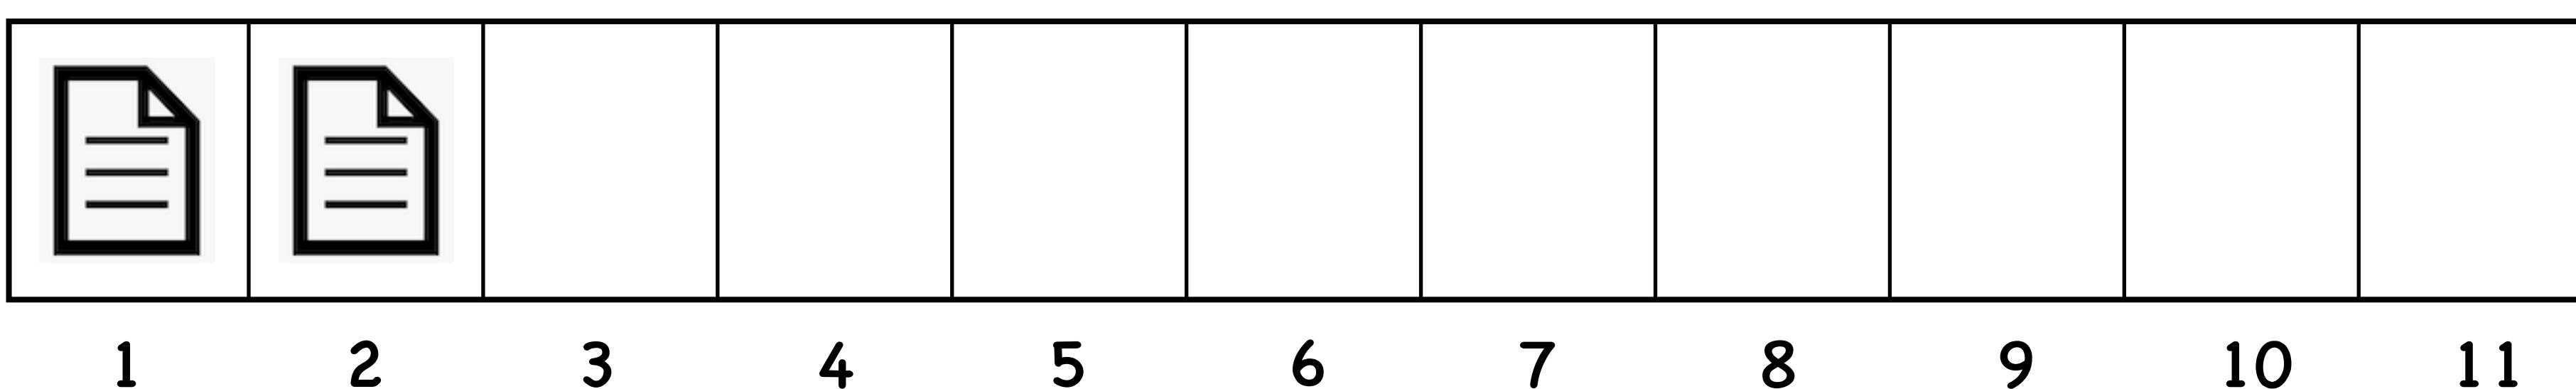
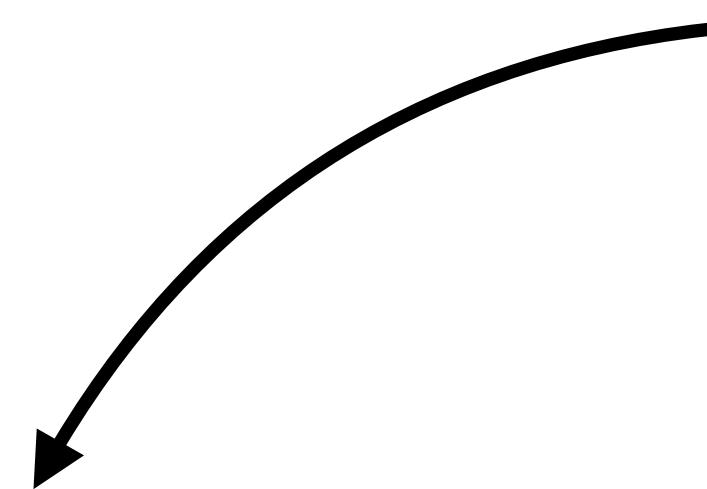
 c1: {300}	 c1: {301}	 d1: {400}	 d1: {401}							
---	--	--	--	--	--	--	--	--	--	--

1 2 3 4 5 6 7 8 9 10 11

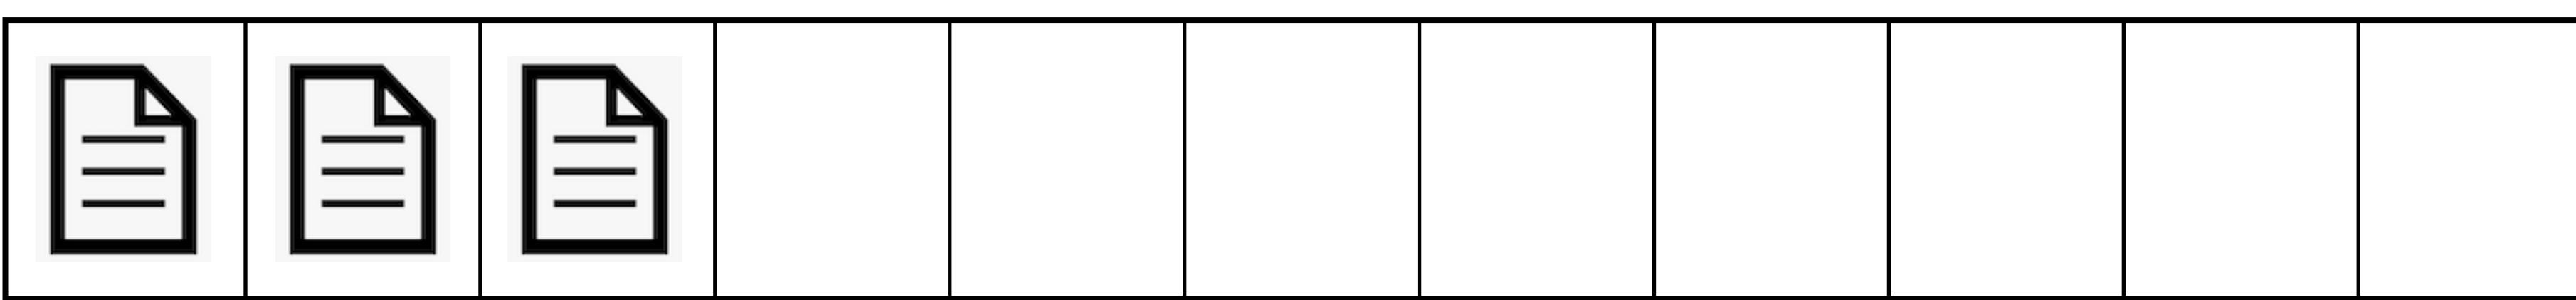
Producer



**Producer**

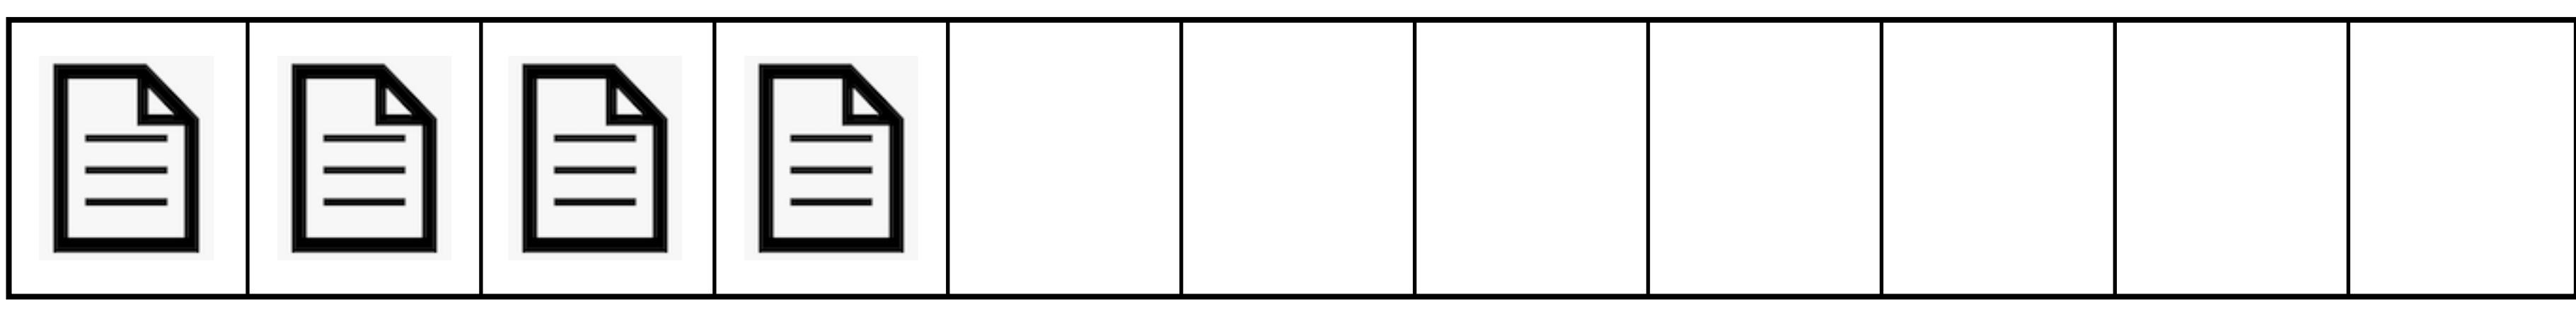


**Producer**



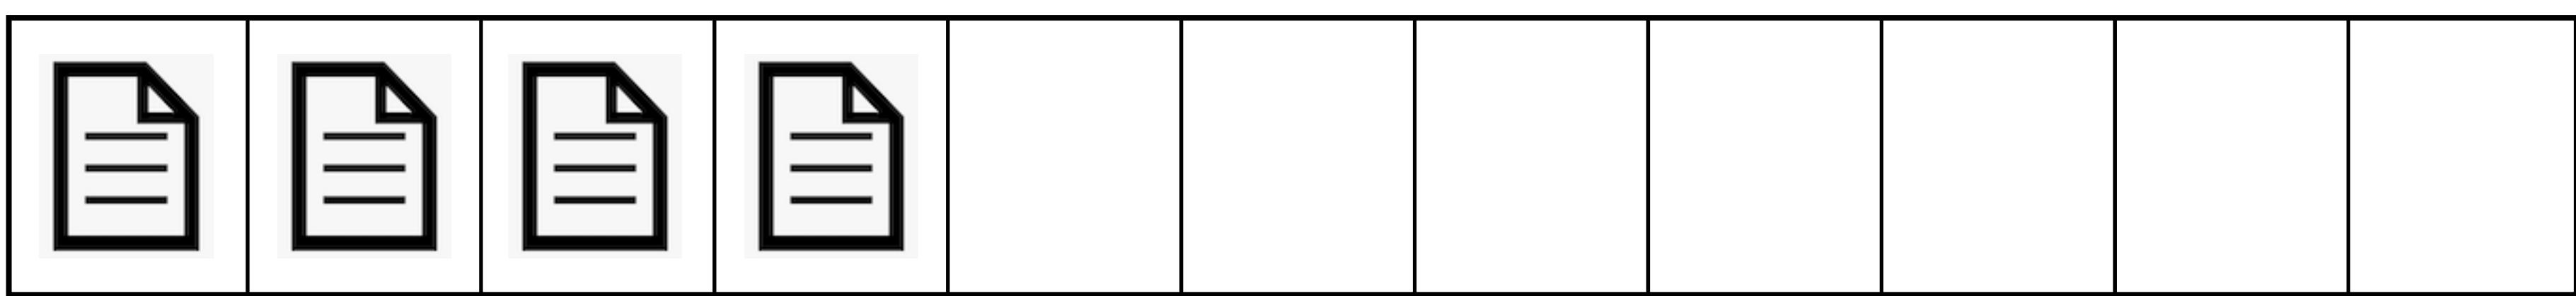
**Consumer**

Producer



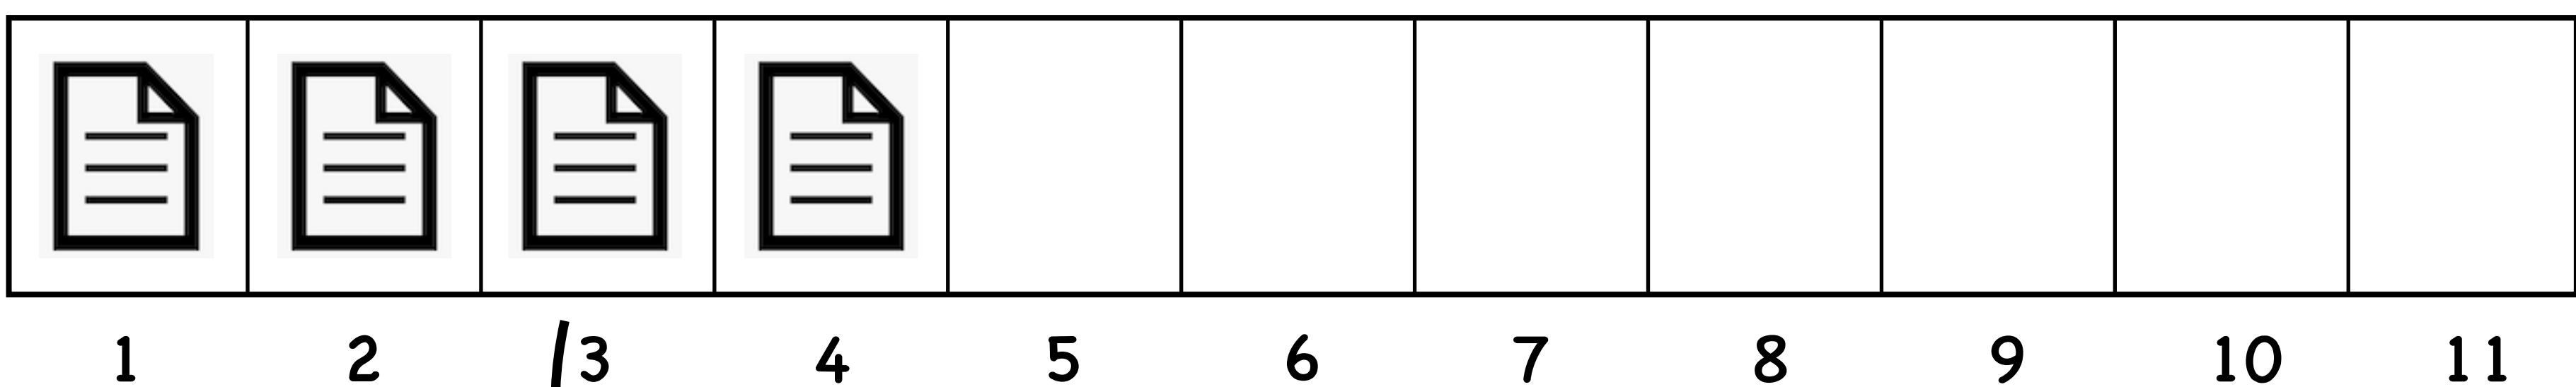
Consumer

**Producer**



**Consumer**

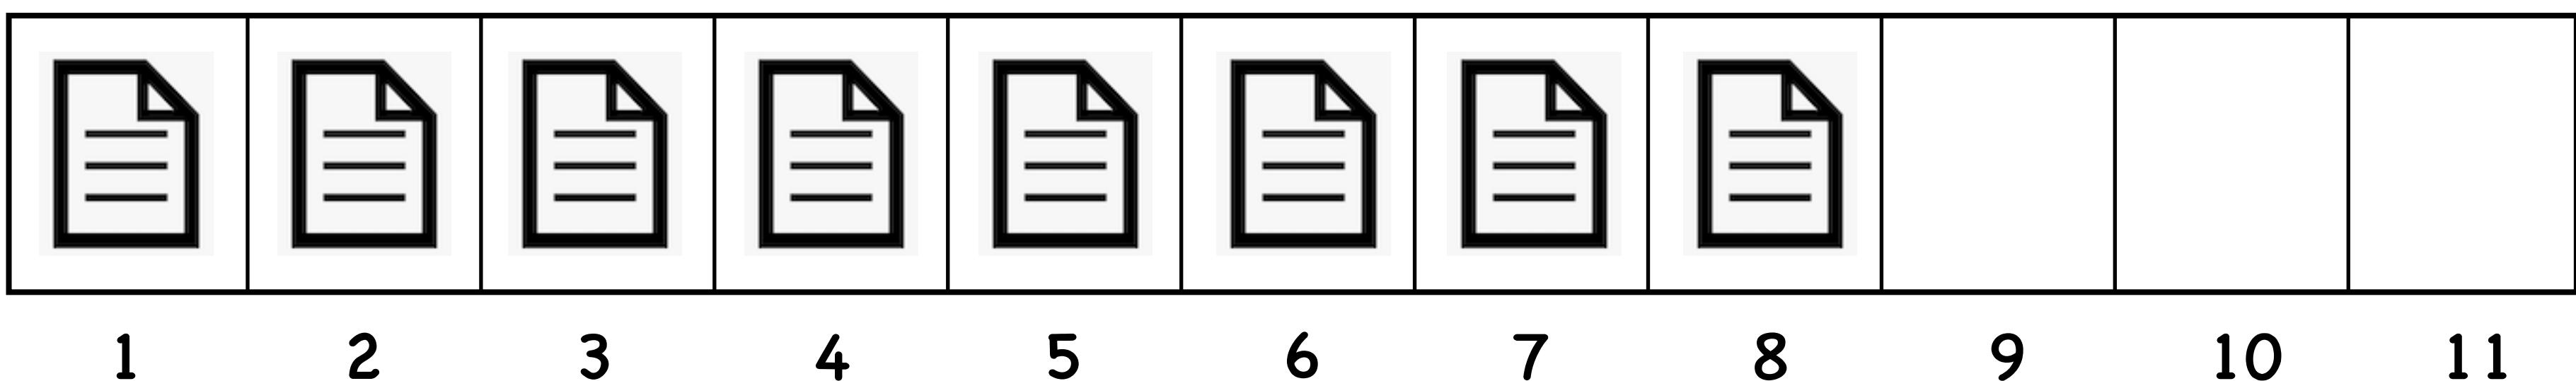
**Producer**



**Consumer**

Committed Offset: 3

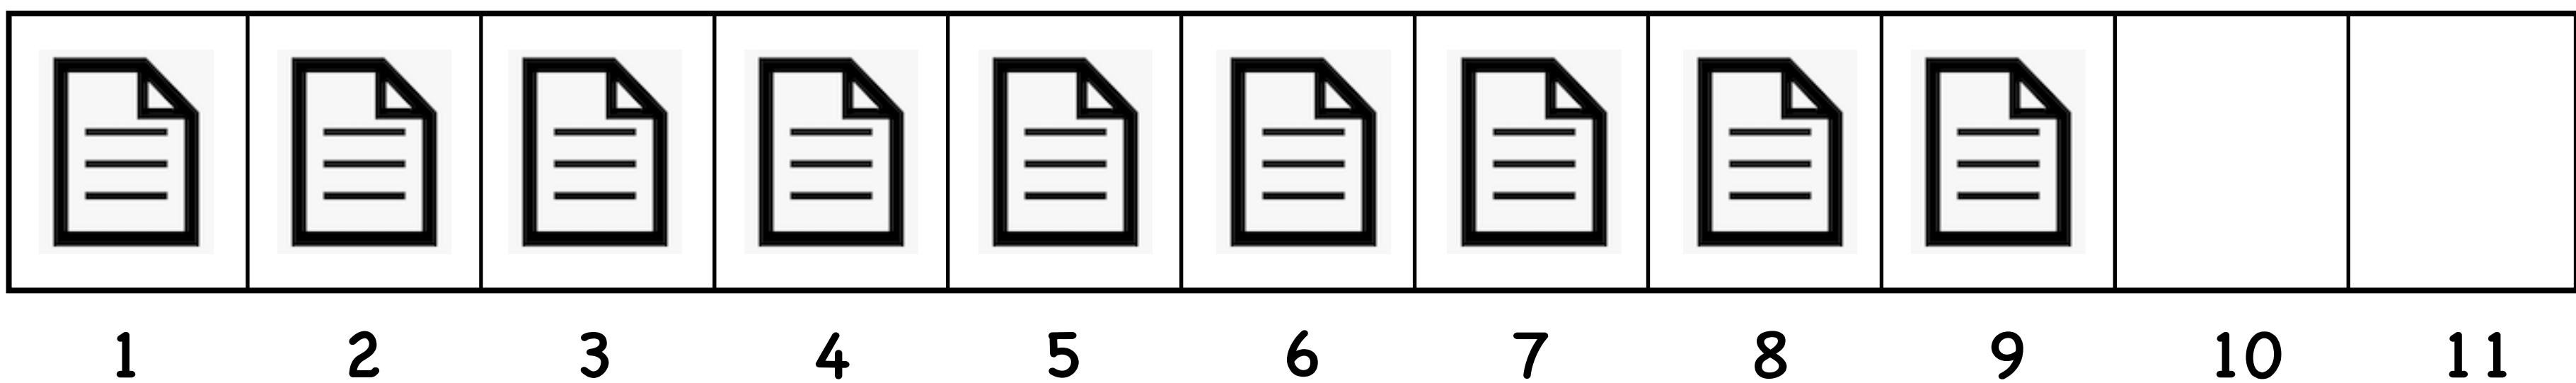
**Producer**



**Consumer**

Committed Offset: 3

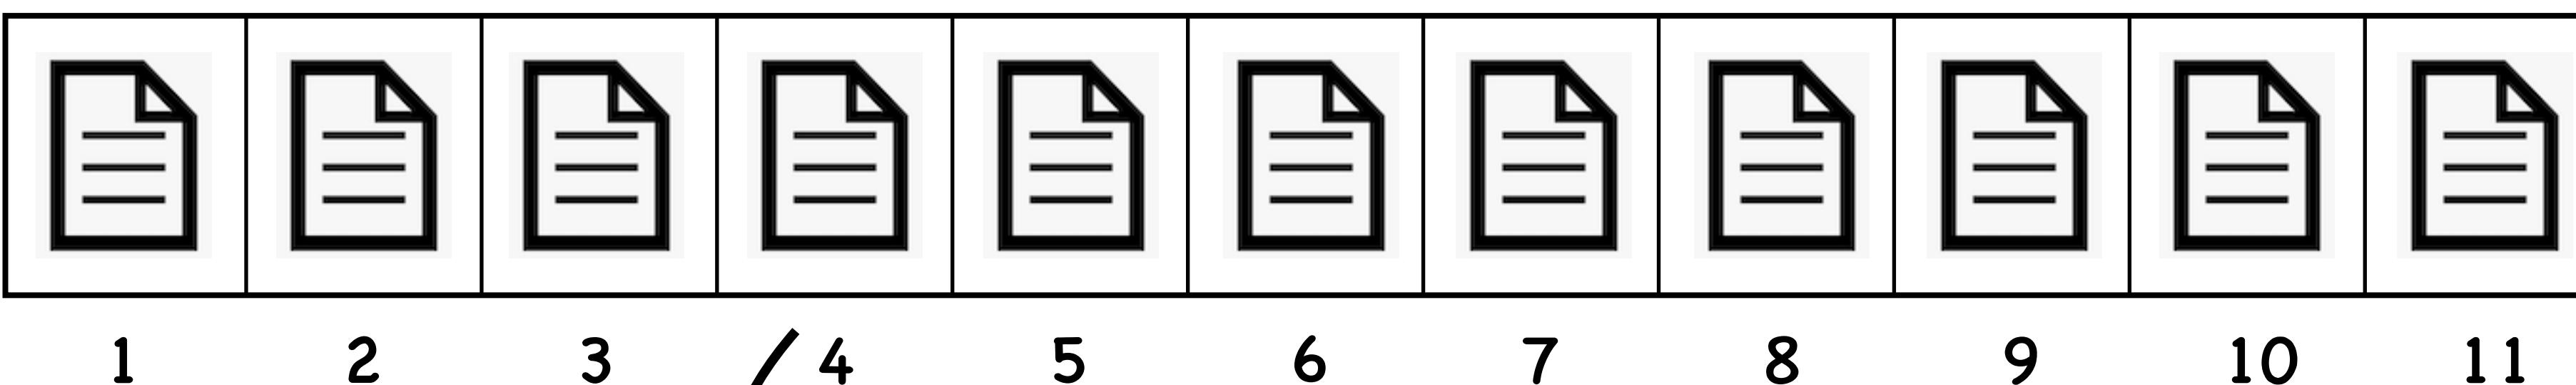
**Producer**



**Consumer**

Committed Offset: 3

**Producer**



**Consumer**

Committed Offset: 3

# Sample Kafka Producer

```
Producer<String, String> producer = new KafkaProducer<String, String>(props);

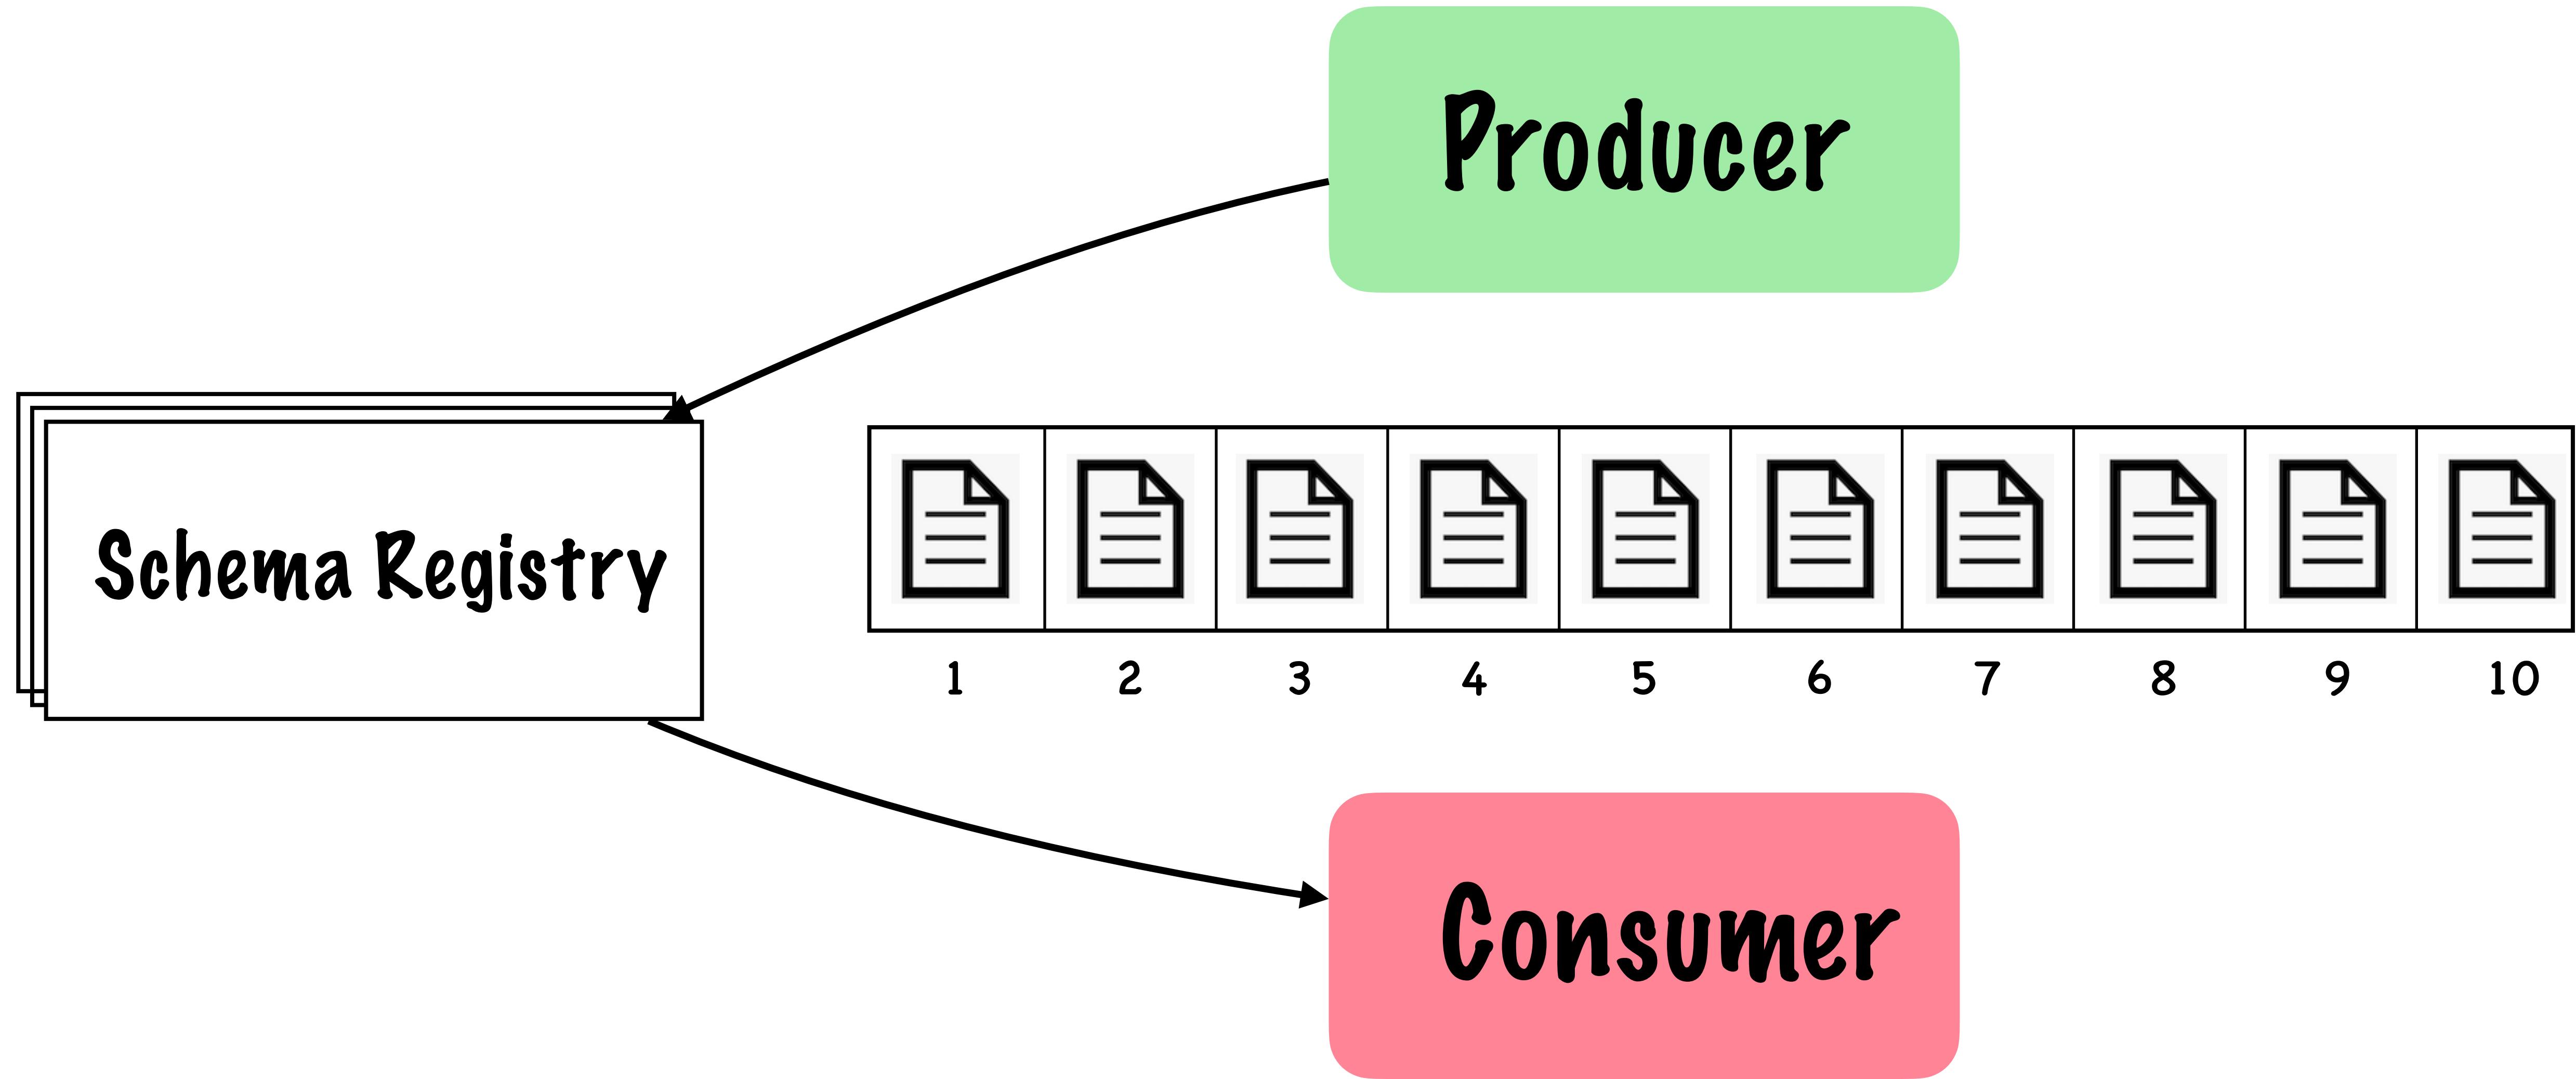
String key = "acme";
String value = "some info about acme, probably in JSON";
ProducerRecord record = new ProducerRecord<String, String>(topic, key, value);

producer.send(record, new Callback() {
    @Override
    public void onCompletion(RecordMetadata m, Exception e) {
        if (e != null) {
            e.printStackTrace();
        } else {
            System.out.printf("Produced record to topic %s partition [%d]",
                m.topic(), m.partition());
        }
    }
});
```

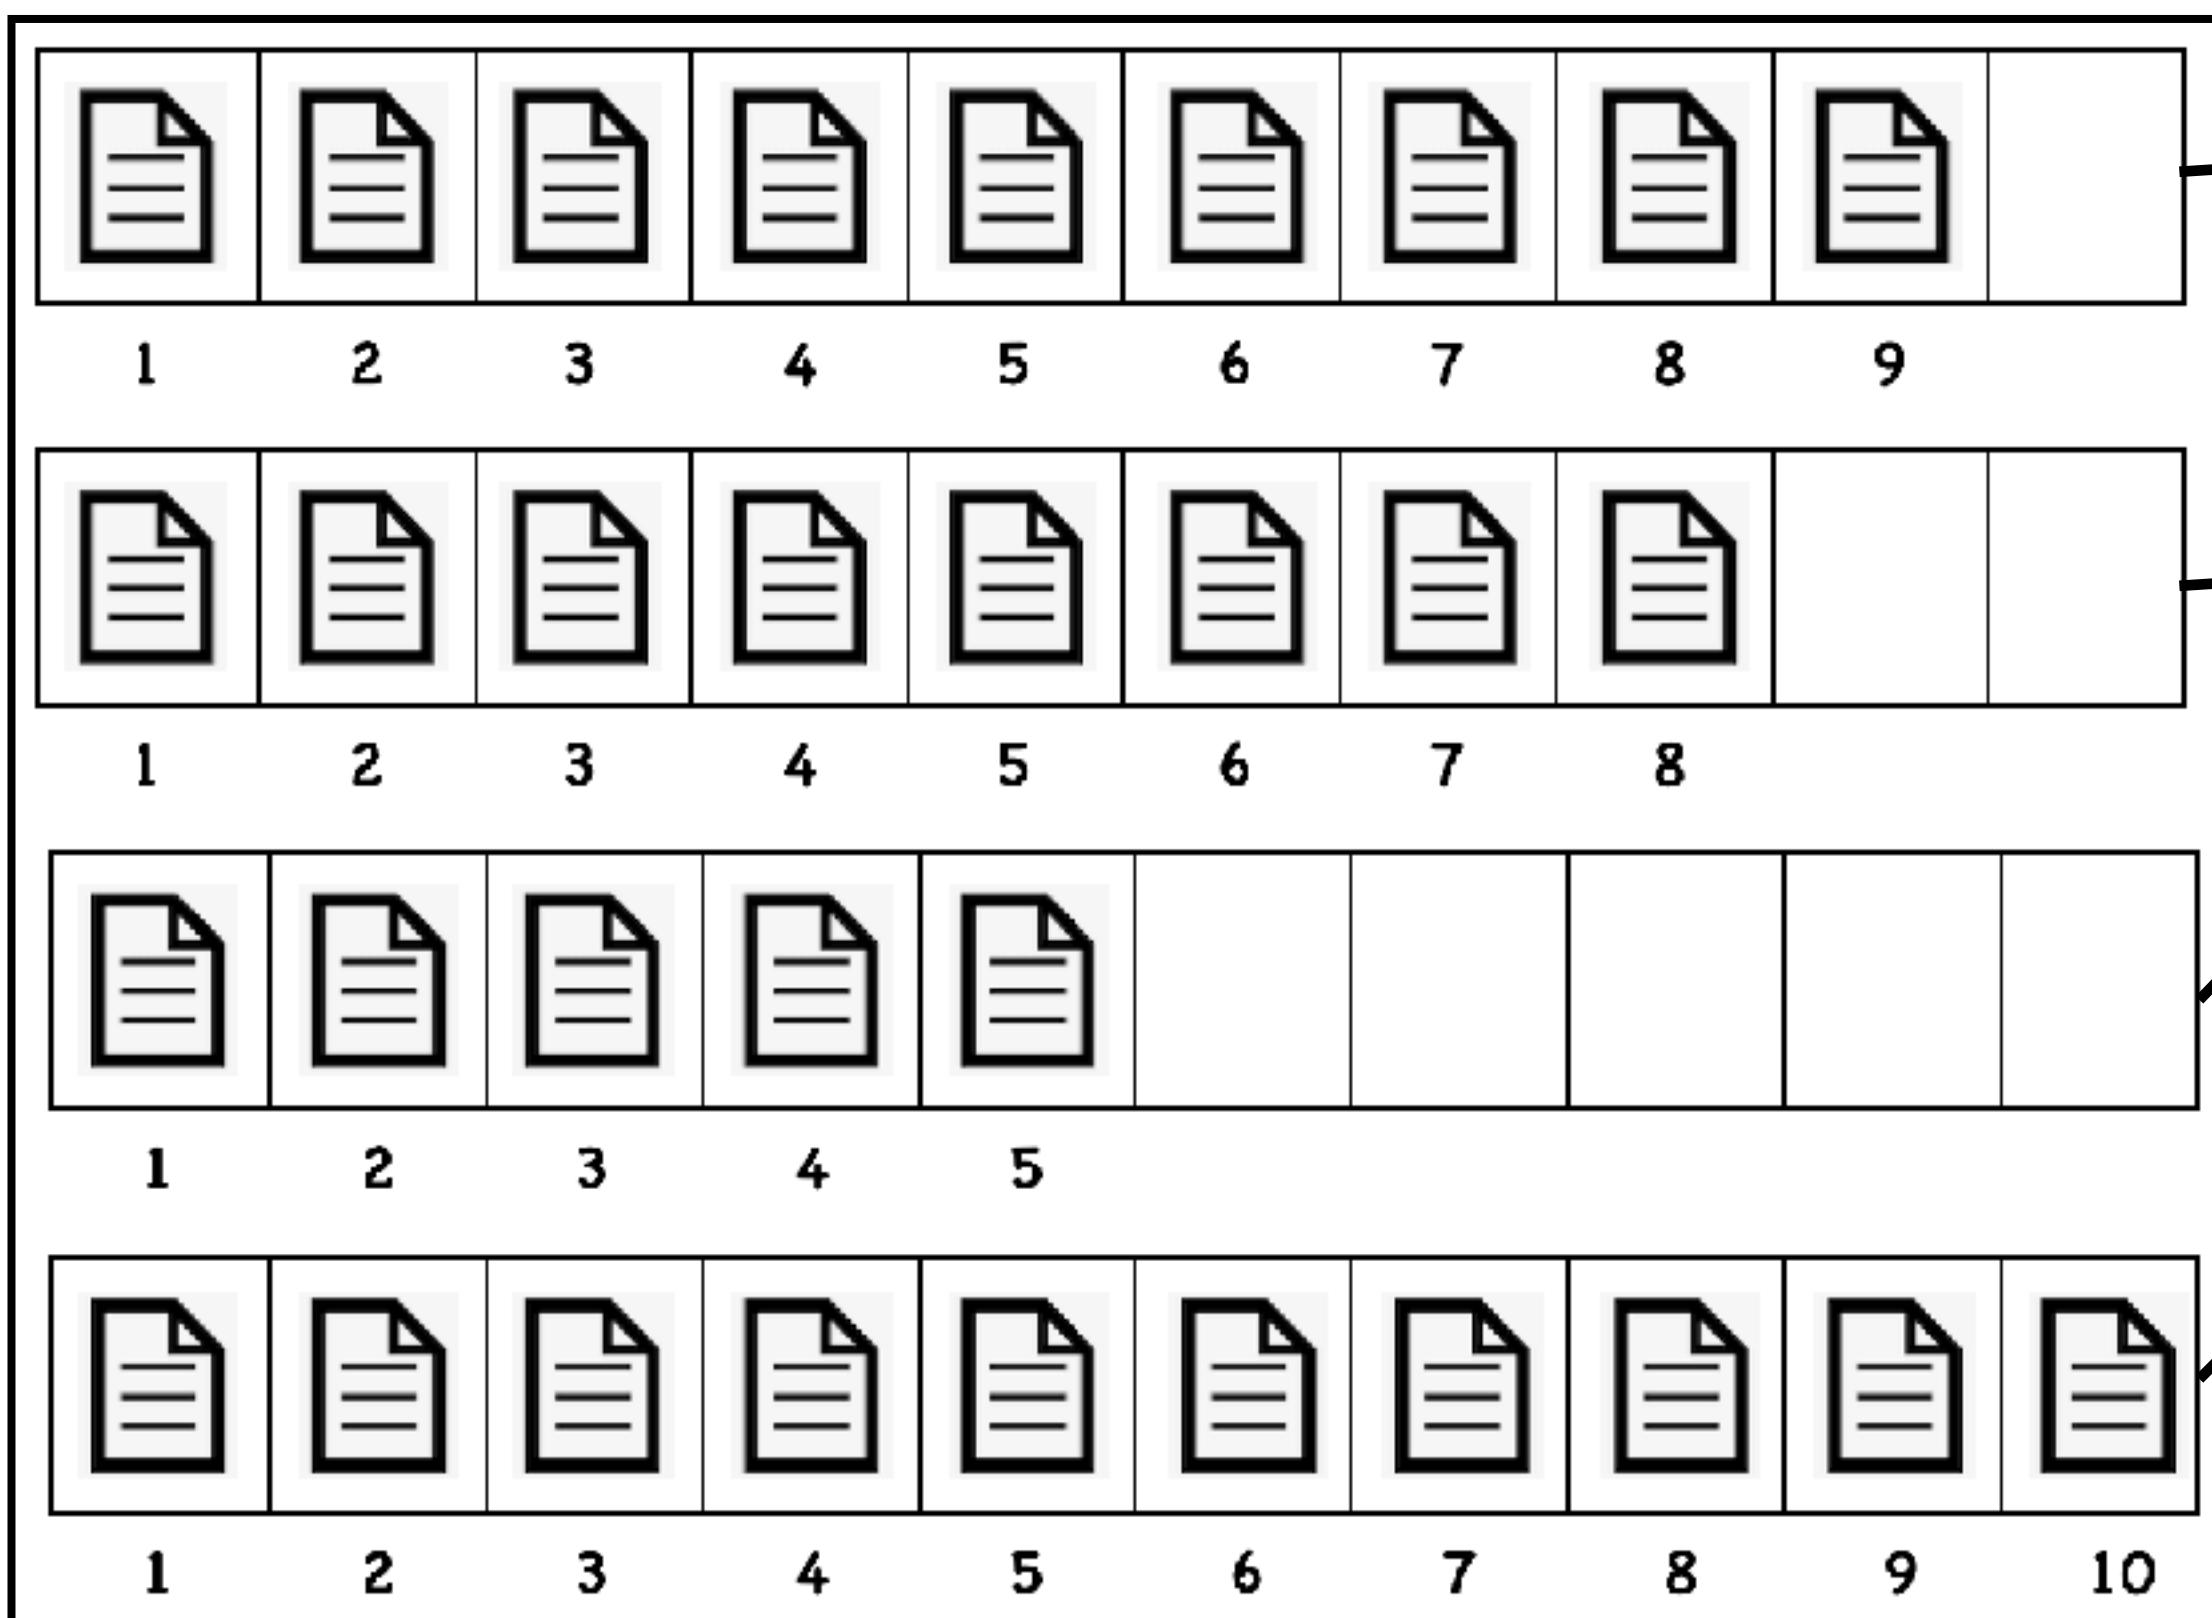
# Sample Kafka Consumer

```
Consumer<String, String> consumer = new KafkaConsumer<String, String>(props);
consumer.subscribe(Arrays.asList(topic));

try {
    while (true) {
        ConsumerRecords<String, String> records = consumer.poll(100);
        for (ConsumerRecord<String, String> record : records) {
            String key = record.key();
            String value = record.value();
            System.out.printf("Consumed record - key: %s and value: %s", key, value);
        }
    }
} finally {
    consumer.close();
}
```



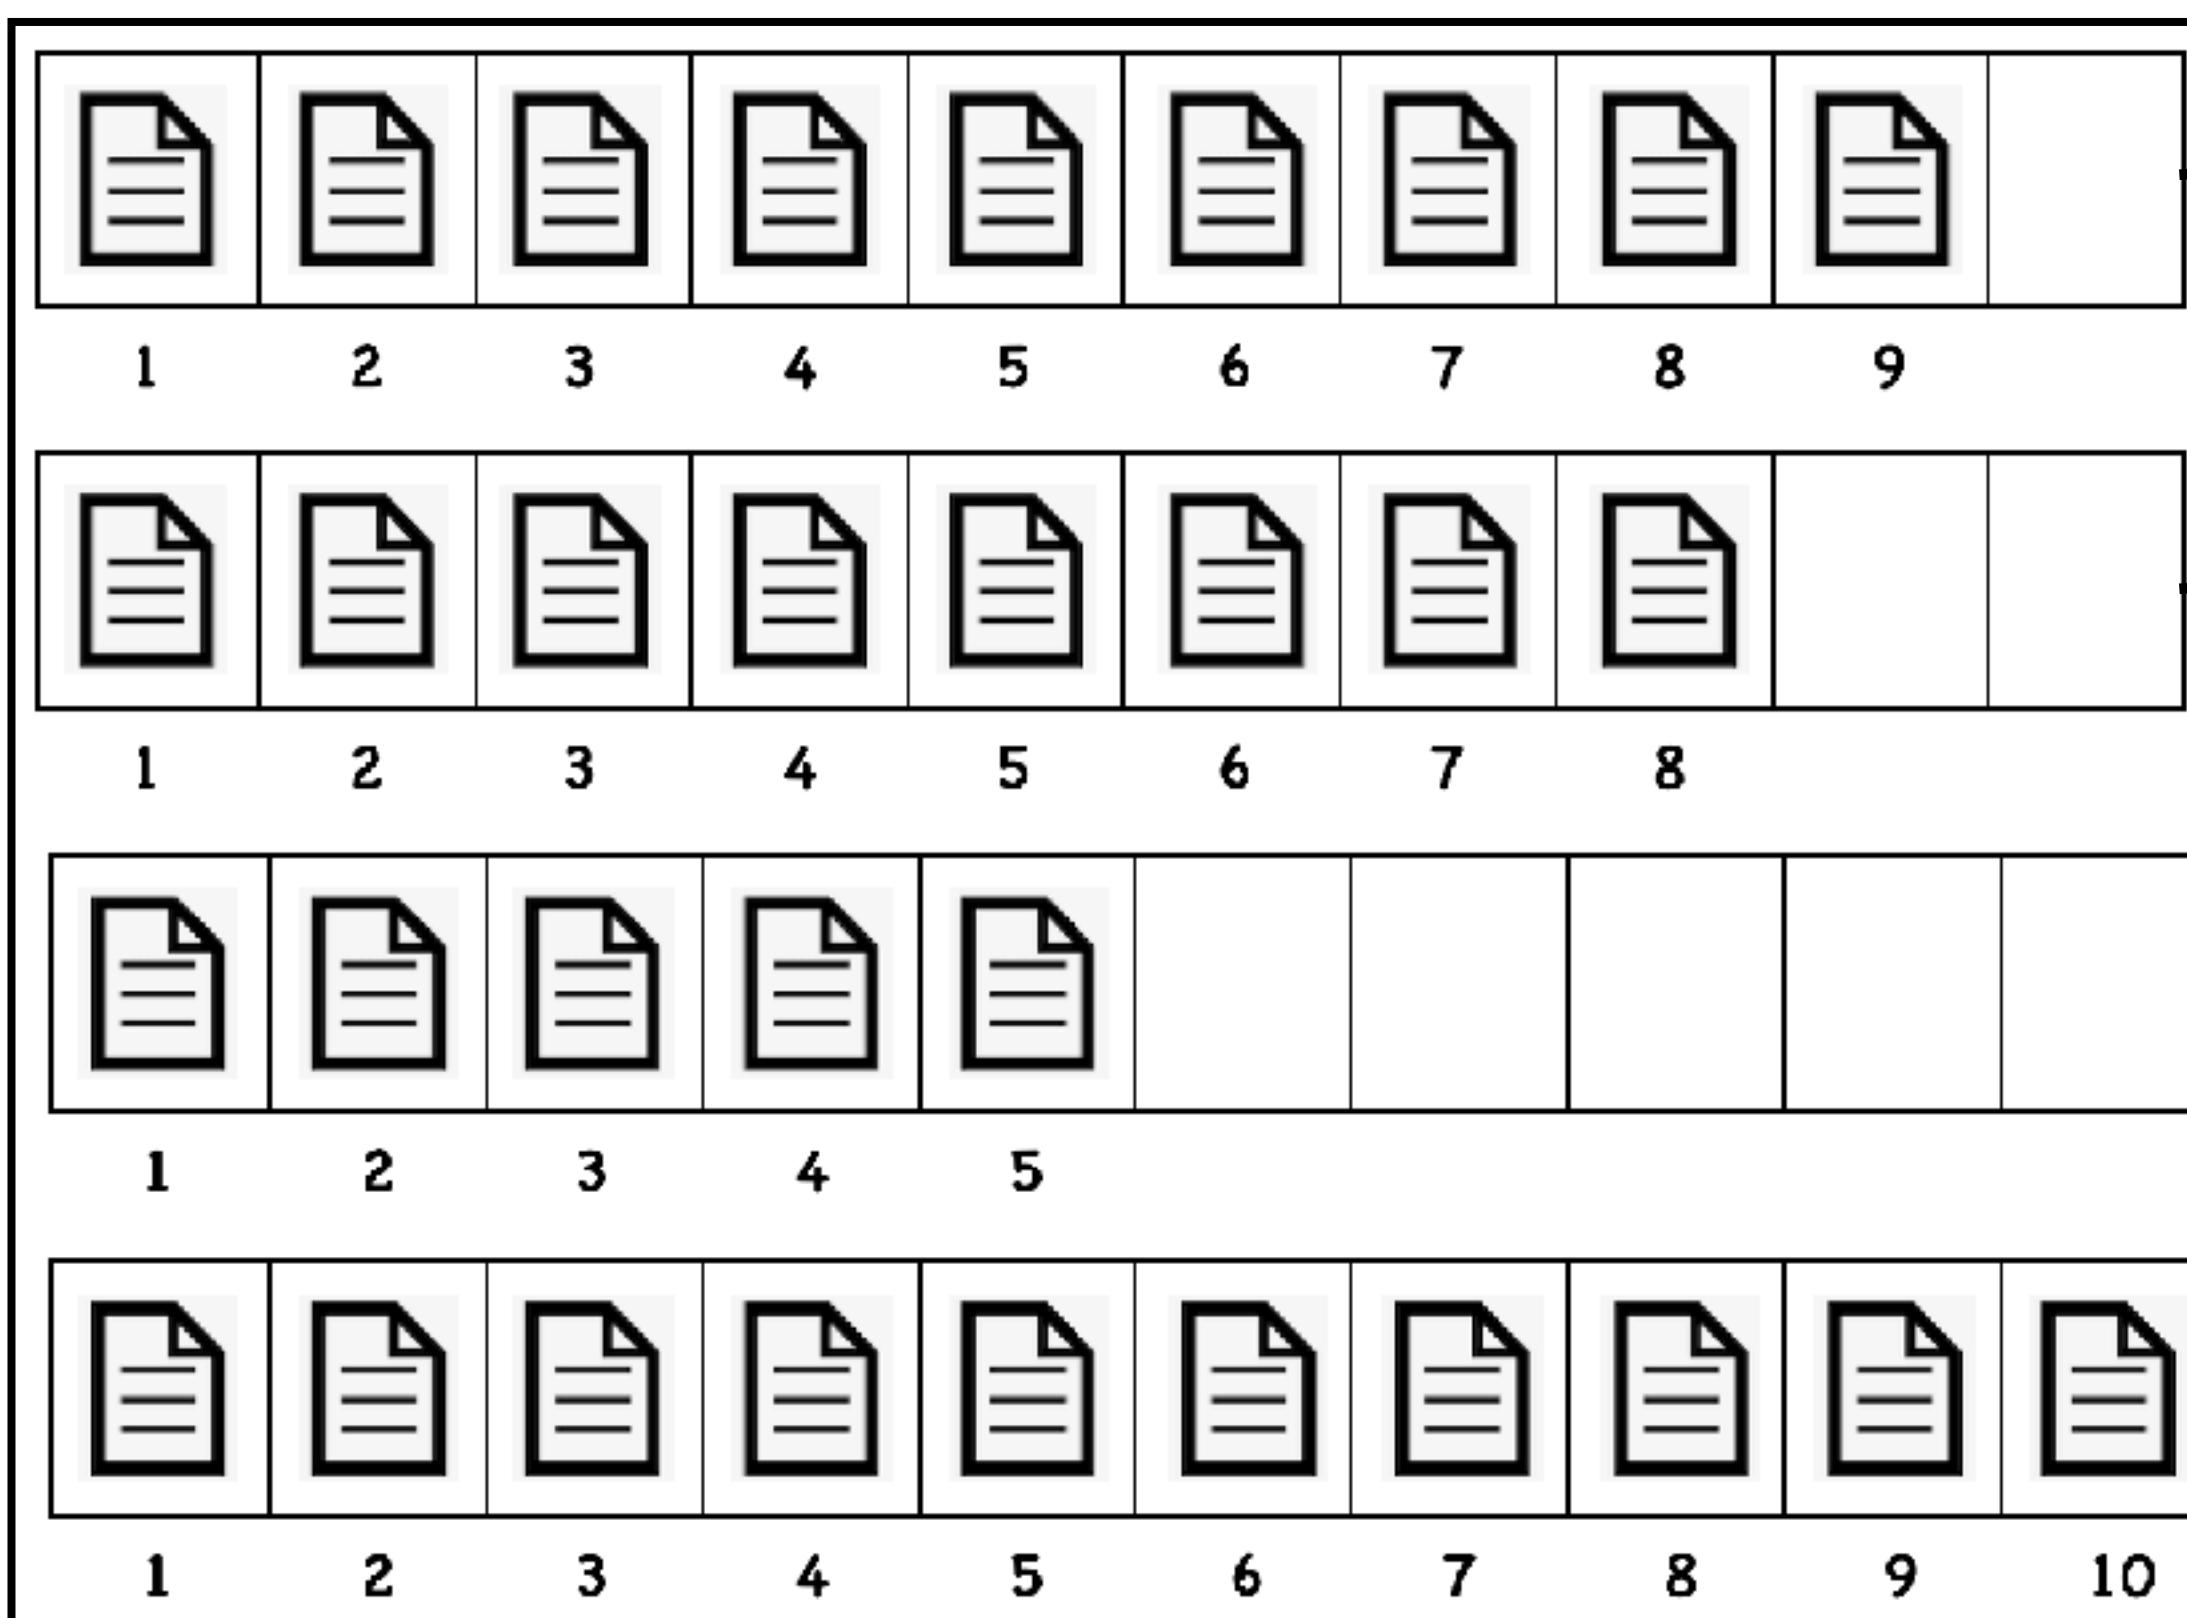
## Consumer Group A



AppInstance01

AppInstance02

## Consumer Group A

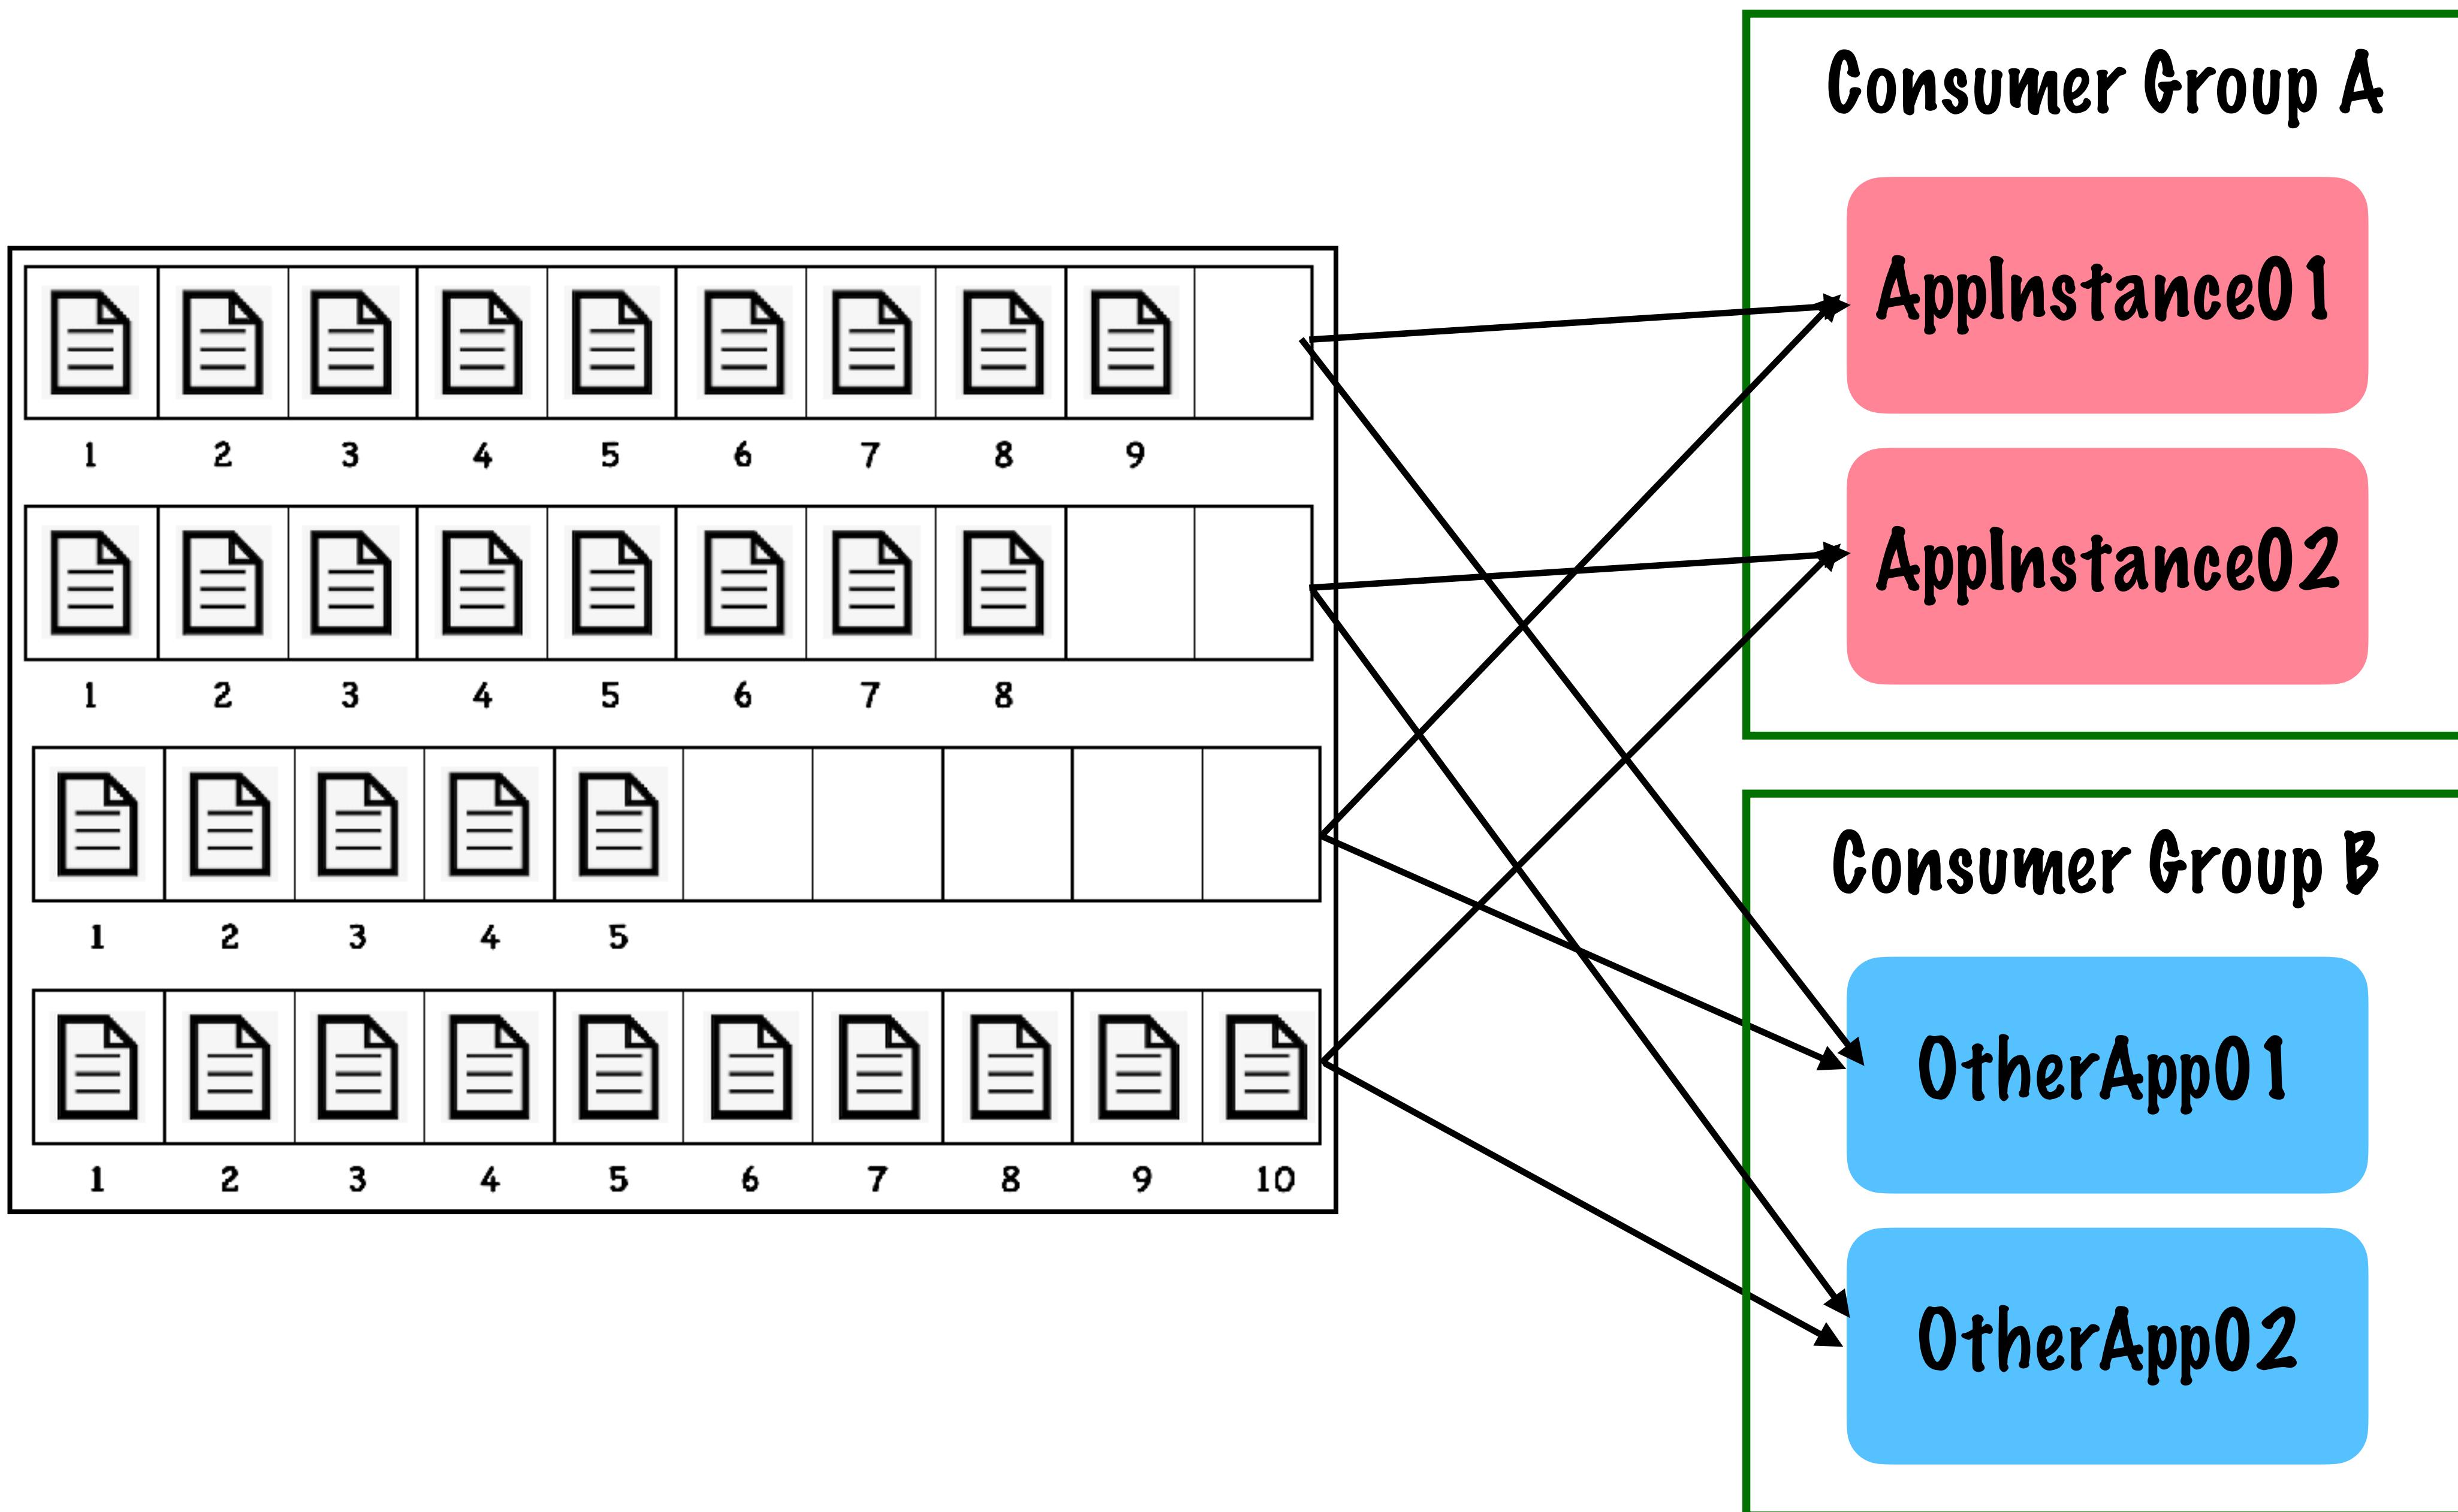


AppInstance01

AppInstance02

AppInstance03

AppInstance04



# Exercises Part 1

[https://github.com/daveklein/kafka-](https://github.com/daveklein/kafka-workshop)

workshop

# Kafka Streams

**What is it?**

# What is it Good For?

Customer 360

Recommendation Engines

Factory Automation

Inventory Management

Fraud Detection

Customer Loyalty Programs

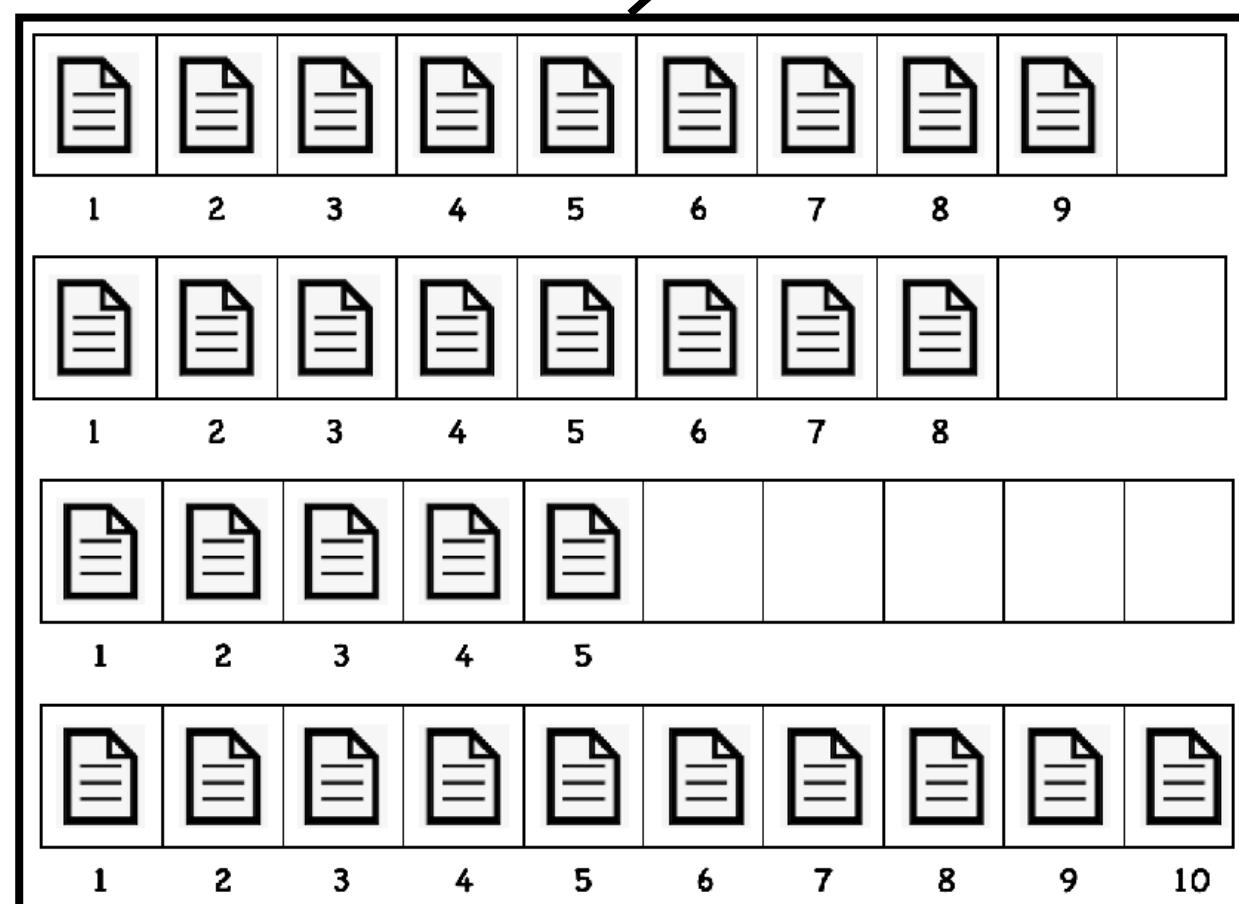
Fleet Management

Real-time Payments

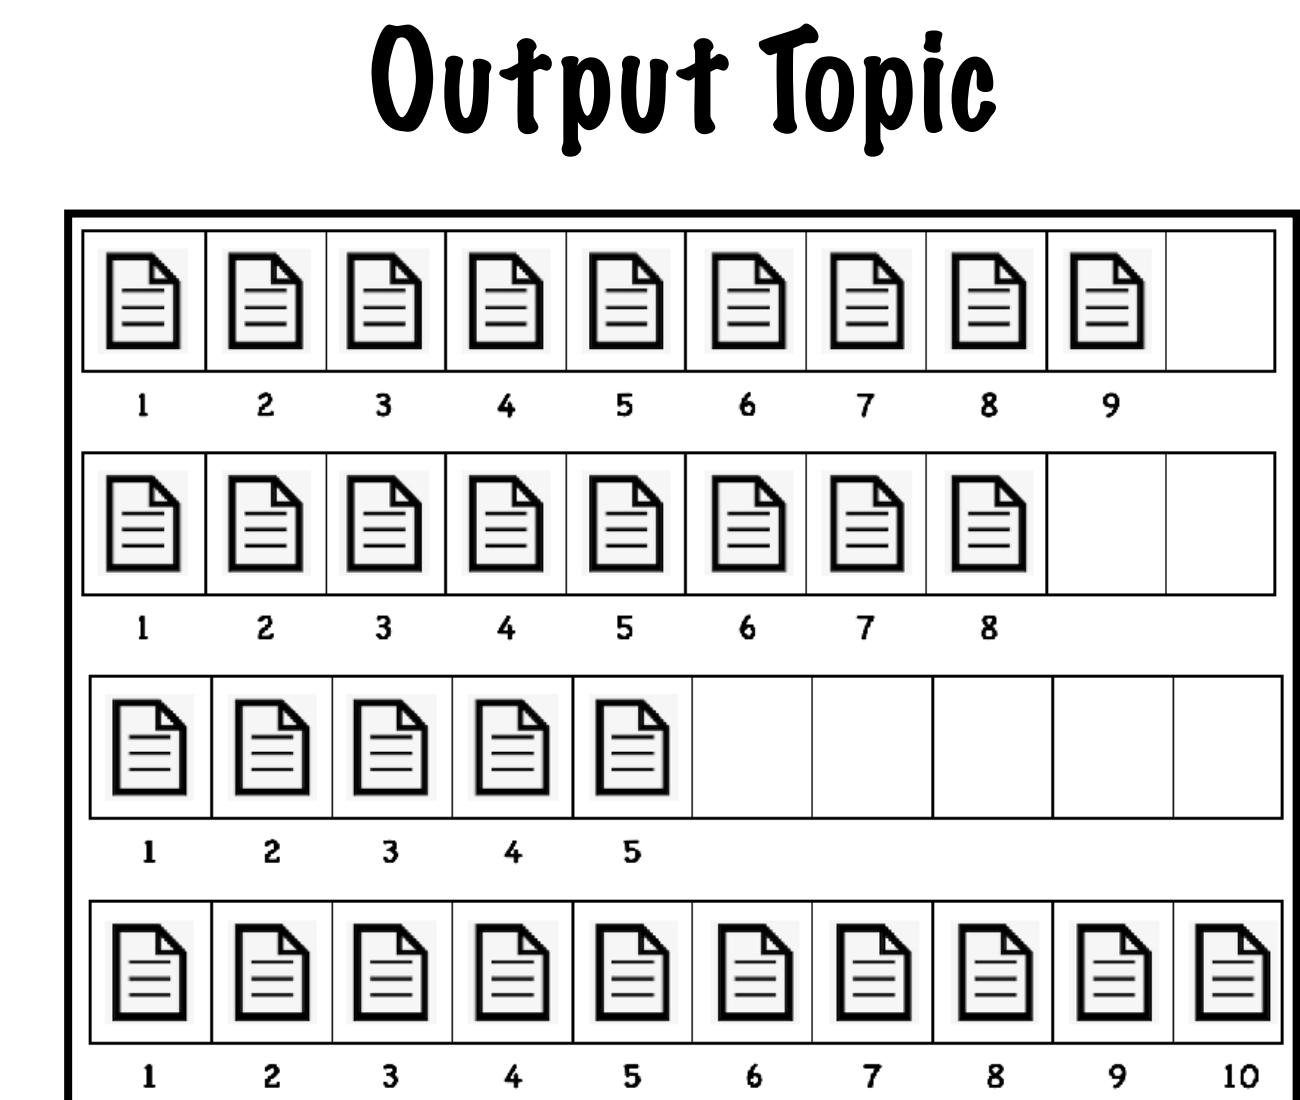
# Kafka Streams DSL

- Java Library (use with any JVM language)
- Define a processor topology (DAG) using a fluent interface
- Stateless operations (Branch, Filter, Map)
- Stateful operations (Count, Aggregate, Reduce)
- Manages state using RocksDB, backed by Kafka topics
- Supports Streams and Tables
- Join Stream -> Stream, Table -> Table, or Stream -> Table

# Kafka Streams

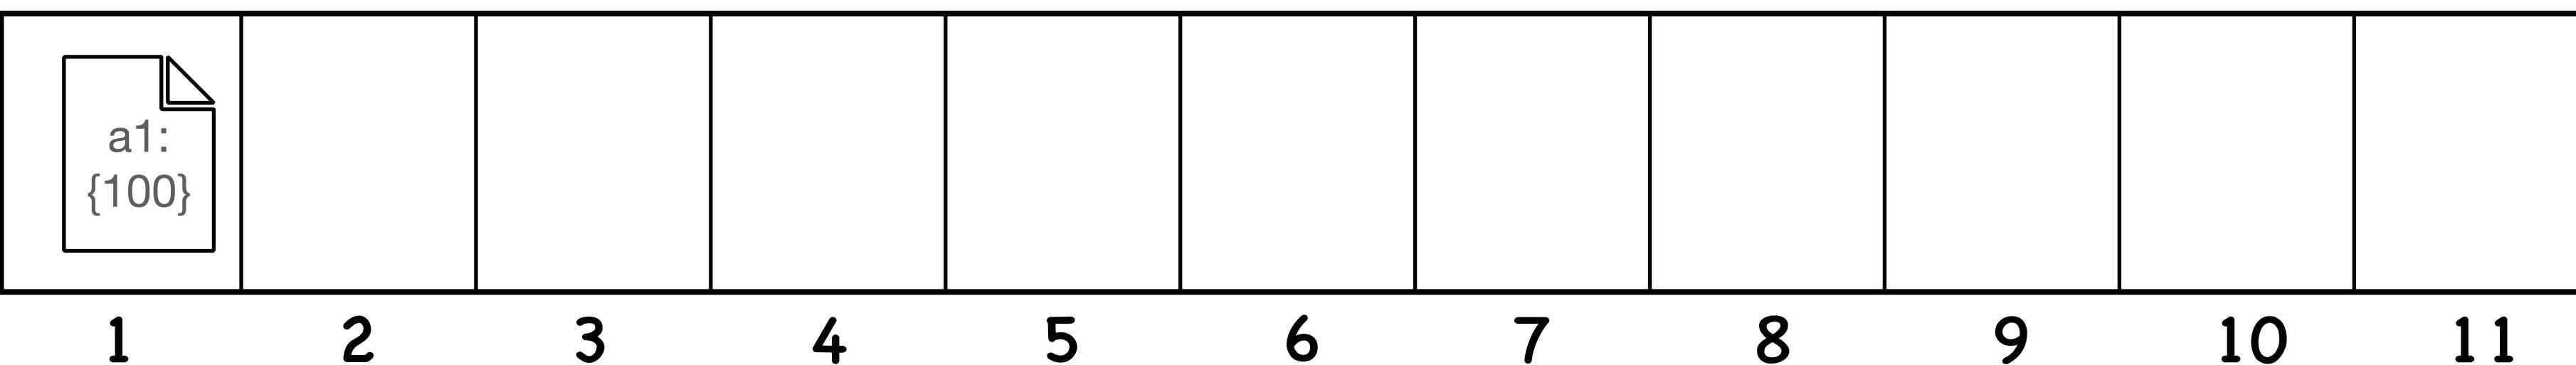


Input Topic



Output Topic

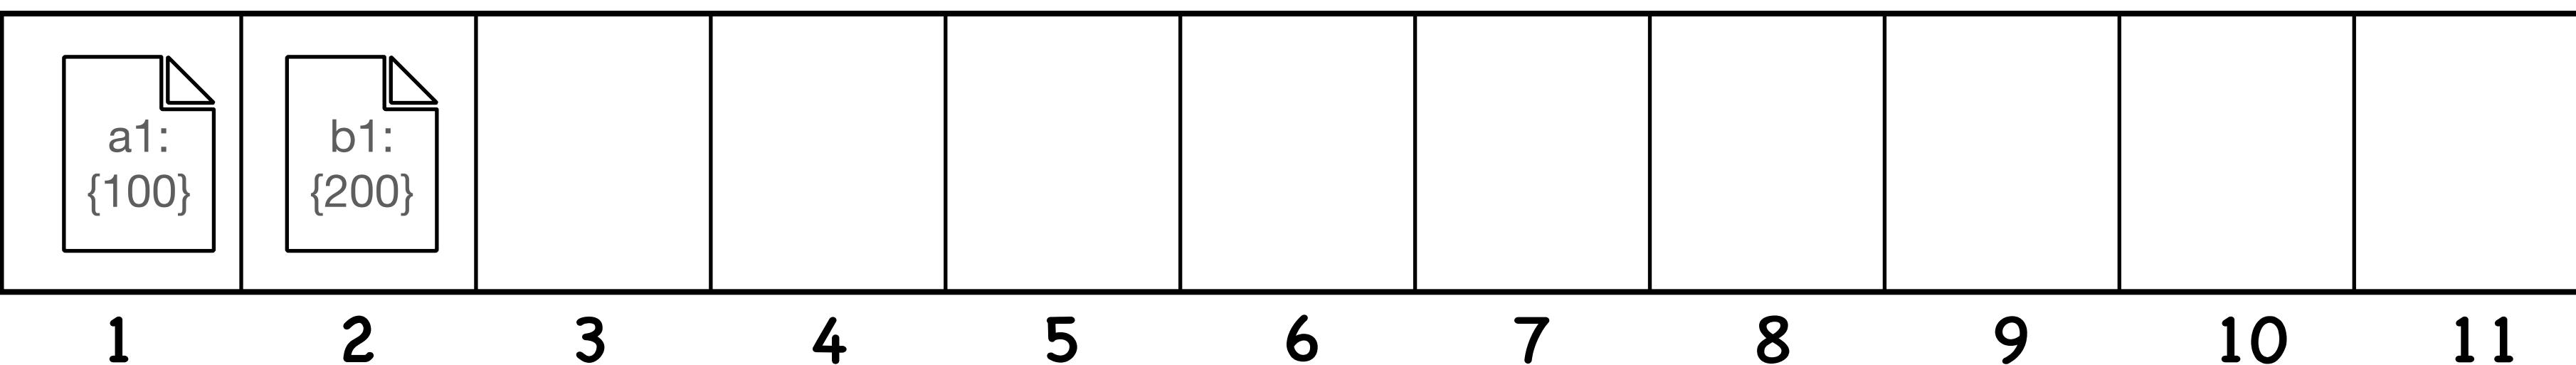
# Stream



# Table

a1	{100}

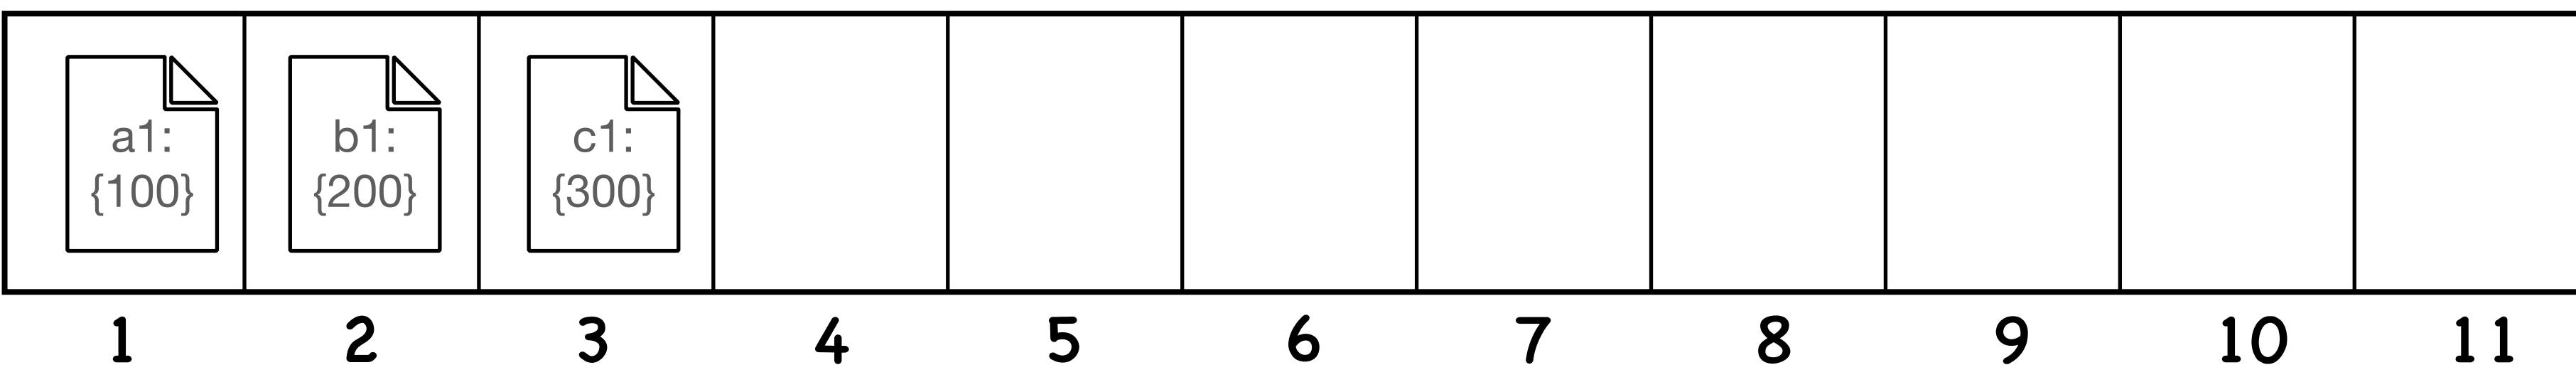
# Stream



# Table

a1	{100}
b1	{200}

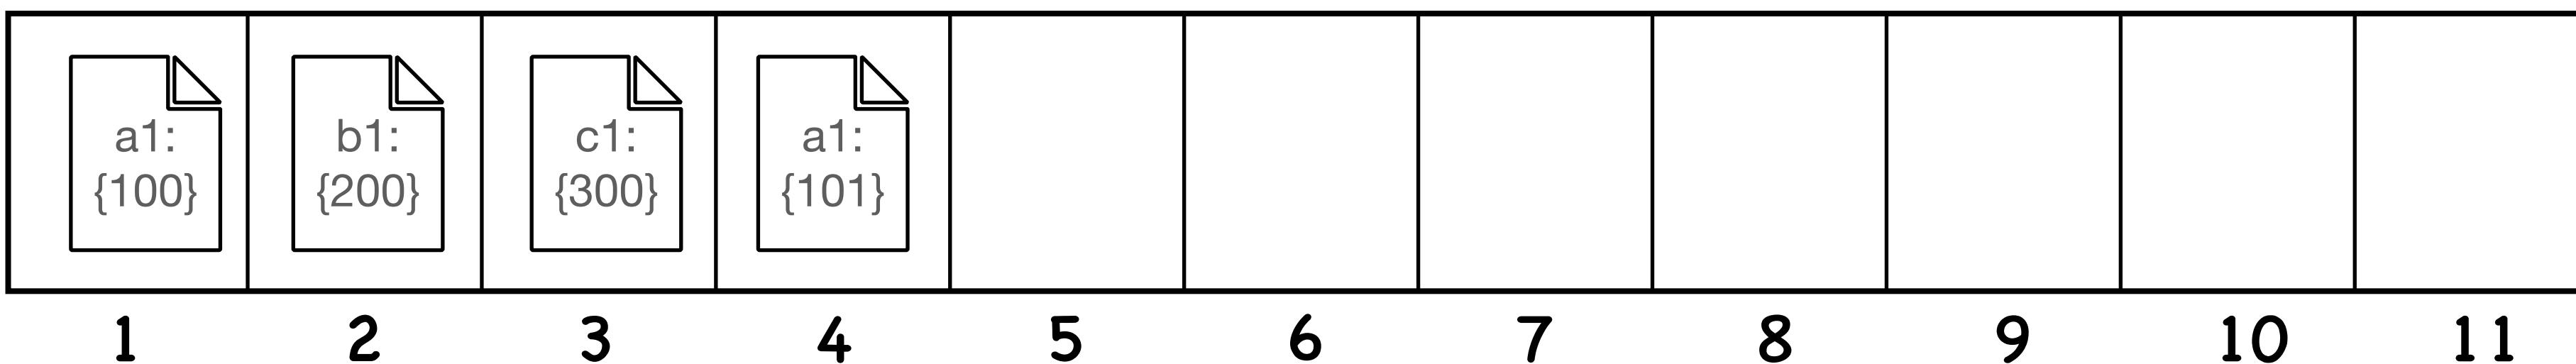
# Stream



# Table

a1	{100}
b1	{200}
c1	{300}

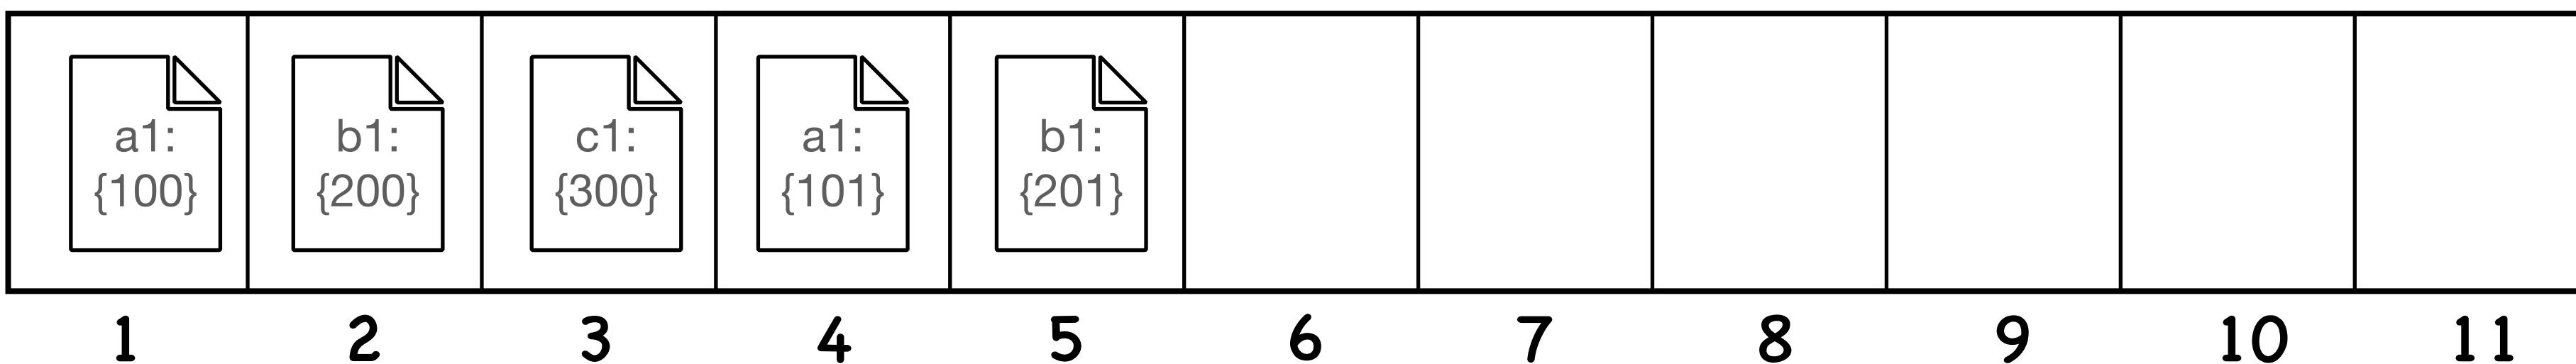
# Stream



# Table

a1	{101}
b1	{200}
c1	{300}

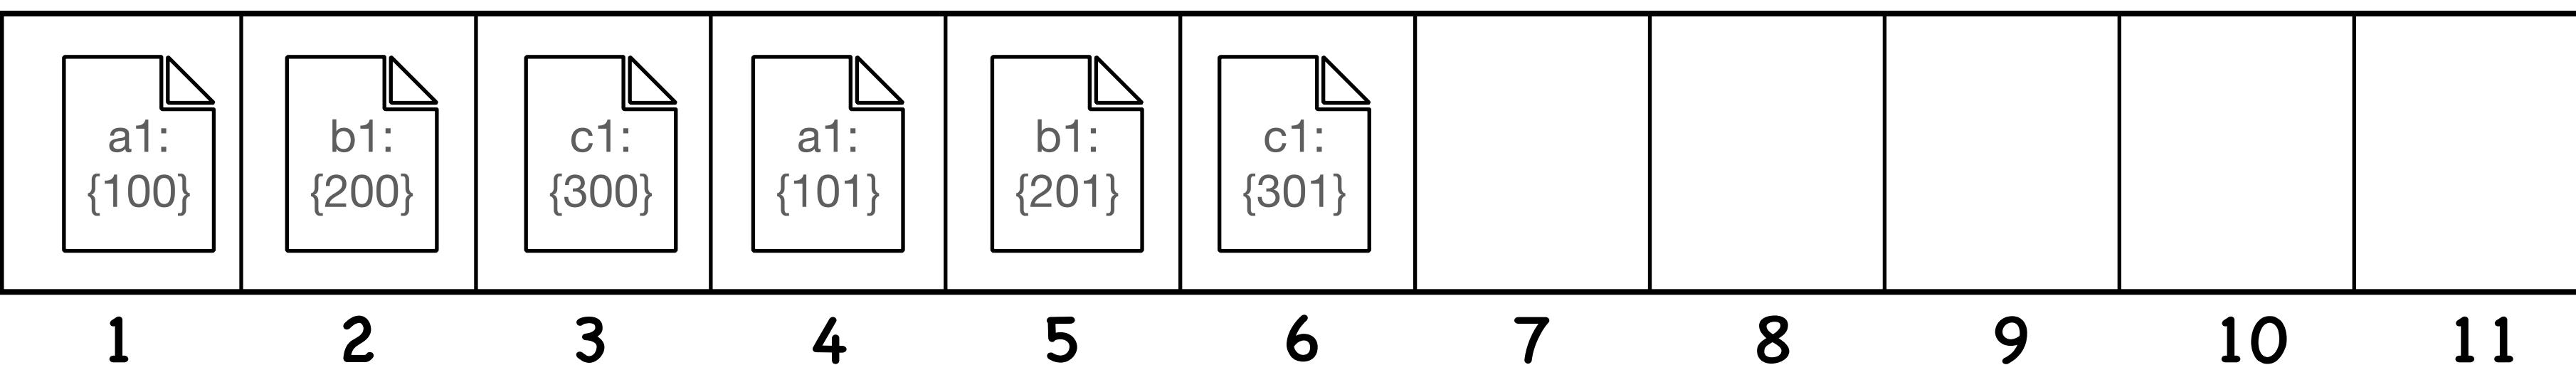
# Stream



# Table

a1	{101}
b1	{201}
c1	{300}

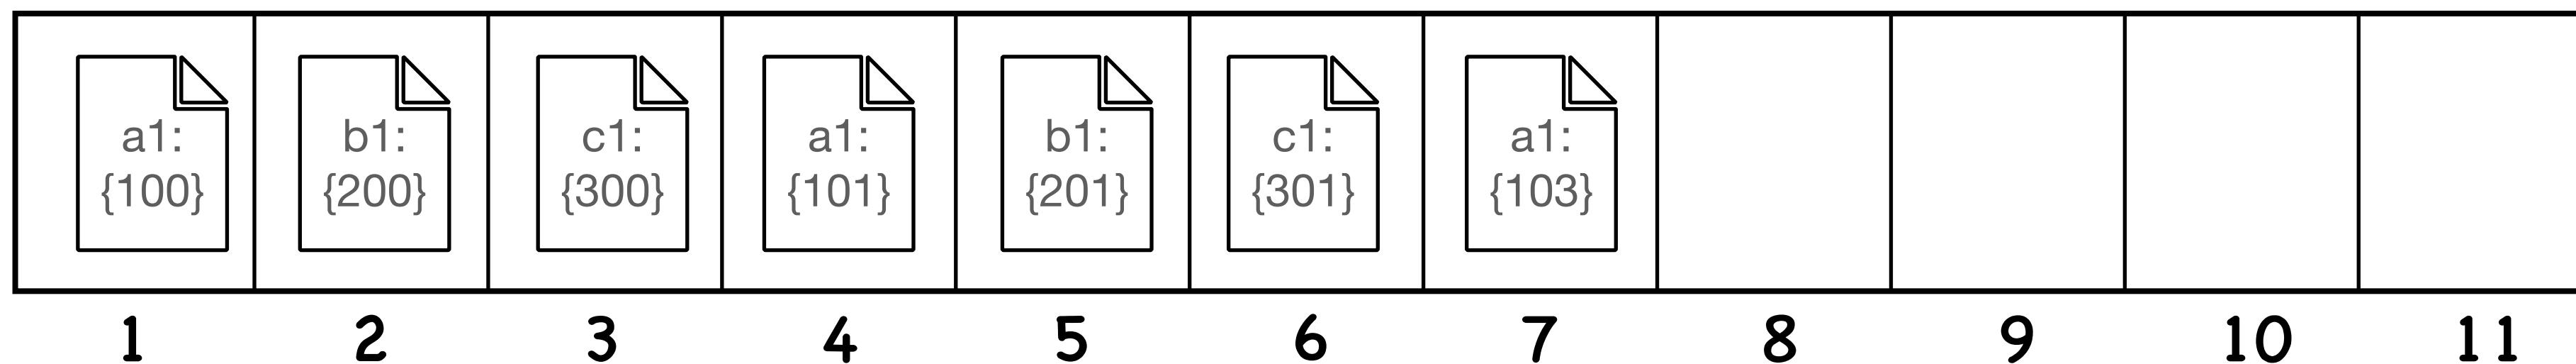
# Stream



# Table

a1	{101}
b1	{201}
c1	{301}

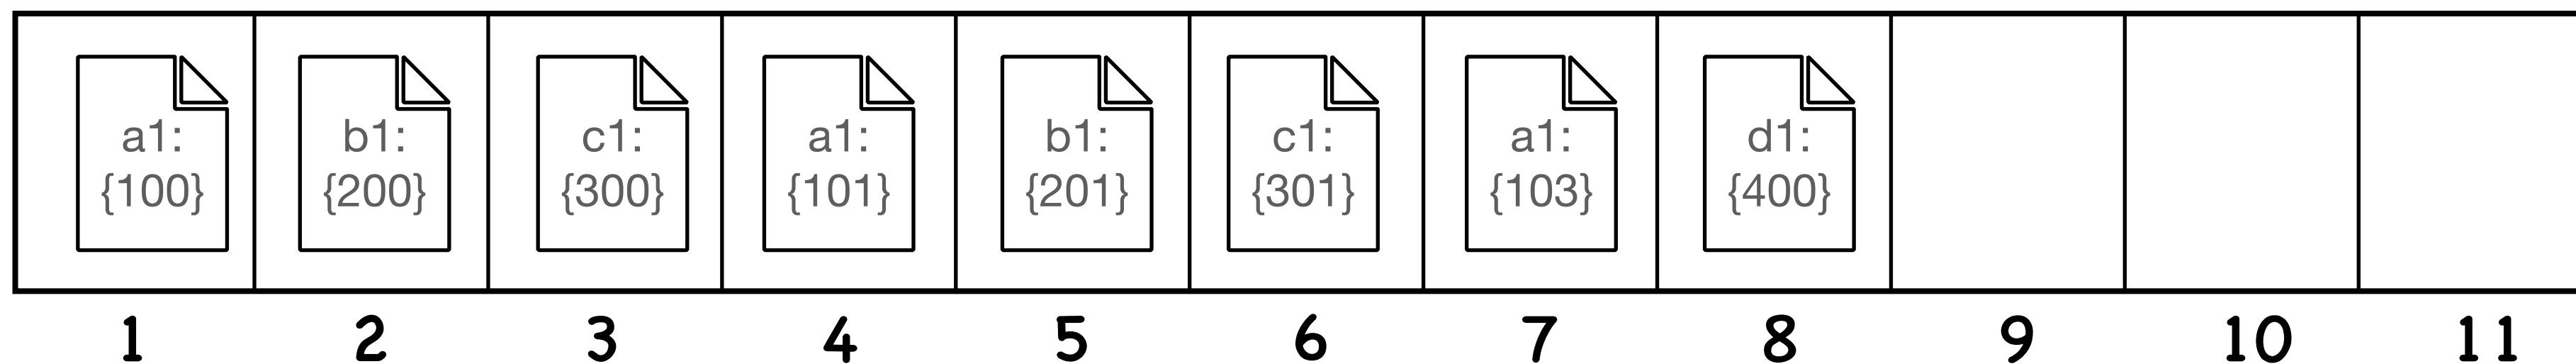
# Stream



# Table

a1	{103}
b1	{201}
c1	{301}

# Stream



# Table

a1	{103}
b1	{201}
c1	{301}
d1	{400}

**Streams  
record history**



**Tables  
represent state**



“The sequence of moves”

“The state of the board”

# Kafka Streams DSL Key Classes

- KafkaStreams - Encapsulates Producer & Consumer, Manages topology execution
- StreamsBuilder - Provides a fluent interface for defining topologies
- StreamsConfig - Helps to accurately assign configuration properties
- KStream - Represents of an unbounded stream of events (key optional)
- KGroupedStream - A stream of events grouped for aggregation
- KTable - Holds the latest value for each key in an event stream (key required)
- Global KTable - KTable that is replicated for each application instance

# Kafka Streams DSL Example

```
final StreamsBuilder builder = new StreamsBuilder();
final Topology topology;

builder.stream(inputTopic, Consumed.with(Serdes.String(), citySerde))
    .filterNot((key, val) -> val.getState() == "NY")
    .map((key, val) -> KeyValue.pair(val.getState(), val.getPopulation()))
    .groupByKey(Grouped.with(Serdes.String(), Serdes.Integer()))
    .reduce((sumVal, val) -> sumVal + val)
    .toStream()
    .to(outputTopic, Produced.with(Serdes.String(), Serdes.Integer()));

topology = builder.build();
```

```
System.out.println(topology.describe().toString());
```

```
KafkaStreams streams = new KafkaStreams(topology, props);
streams.start();
```

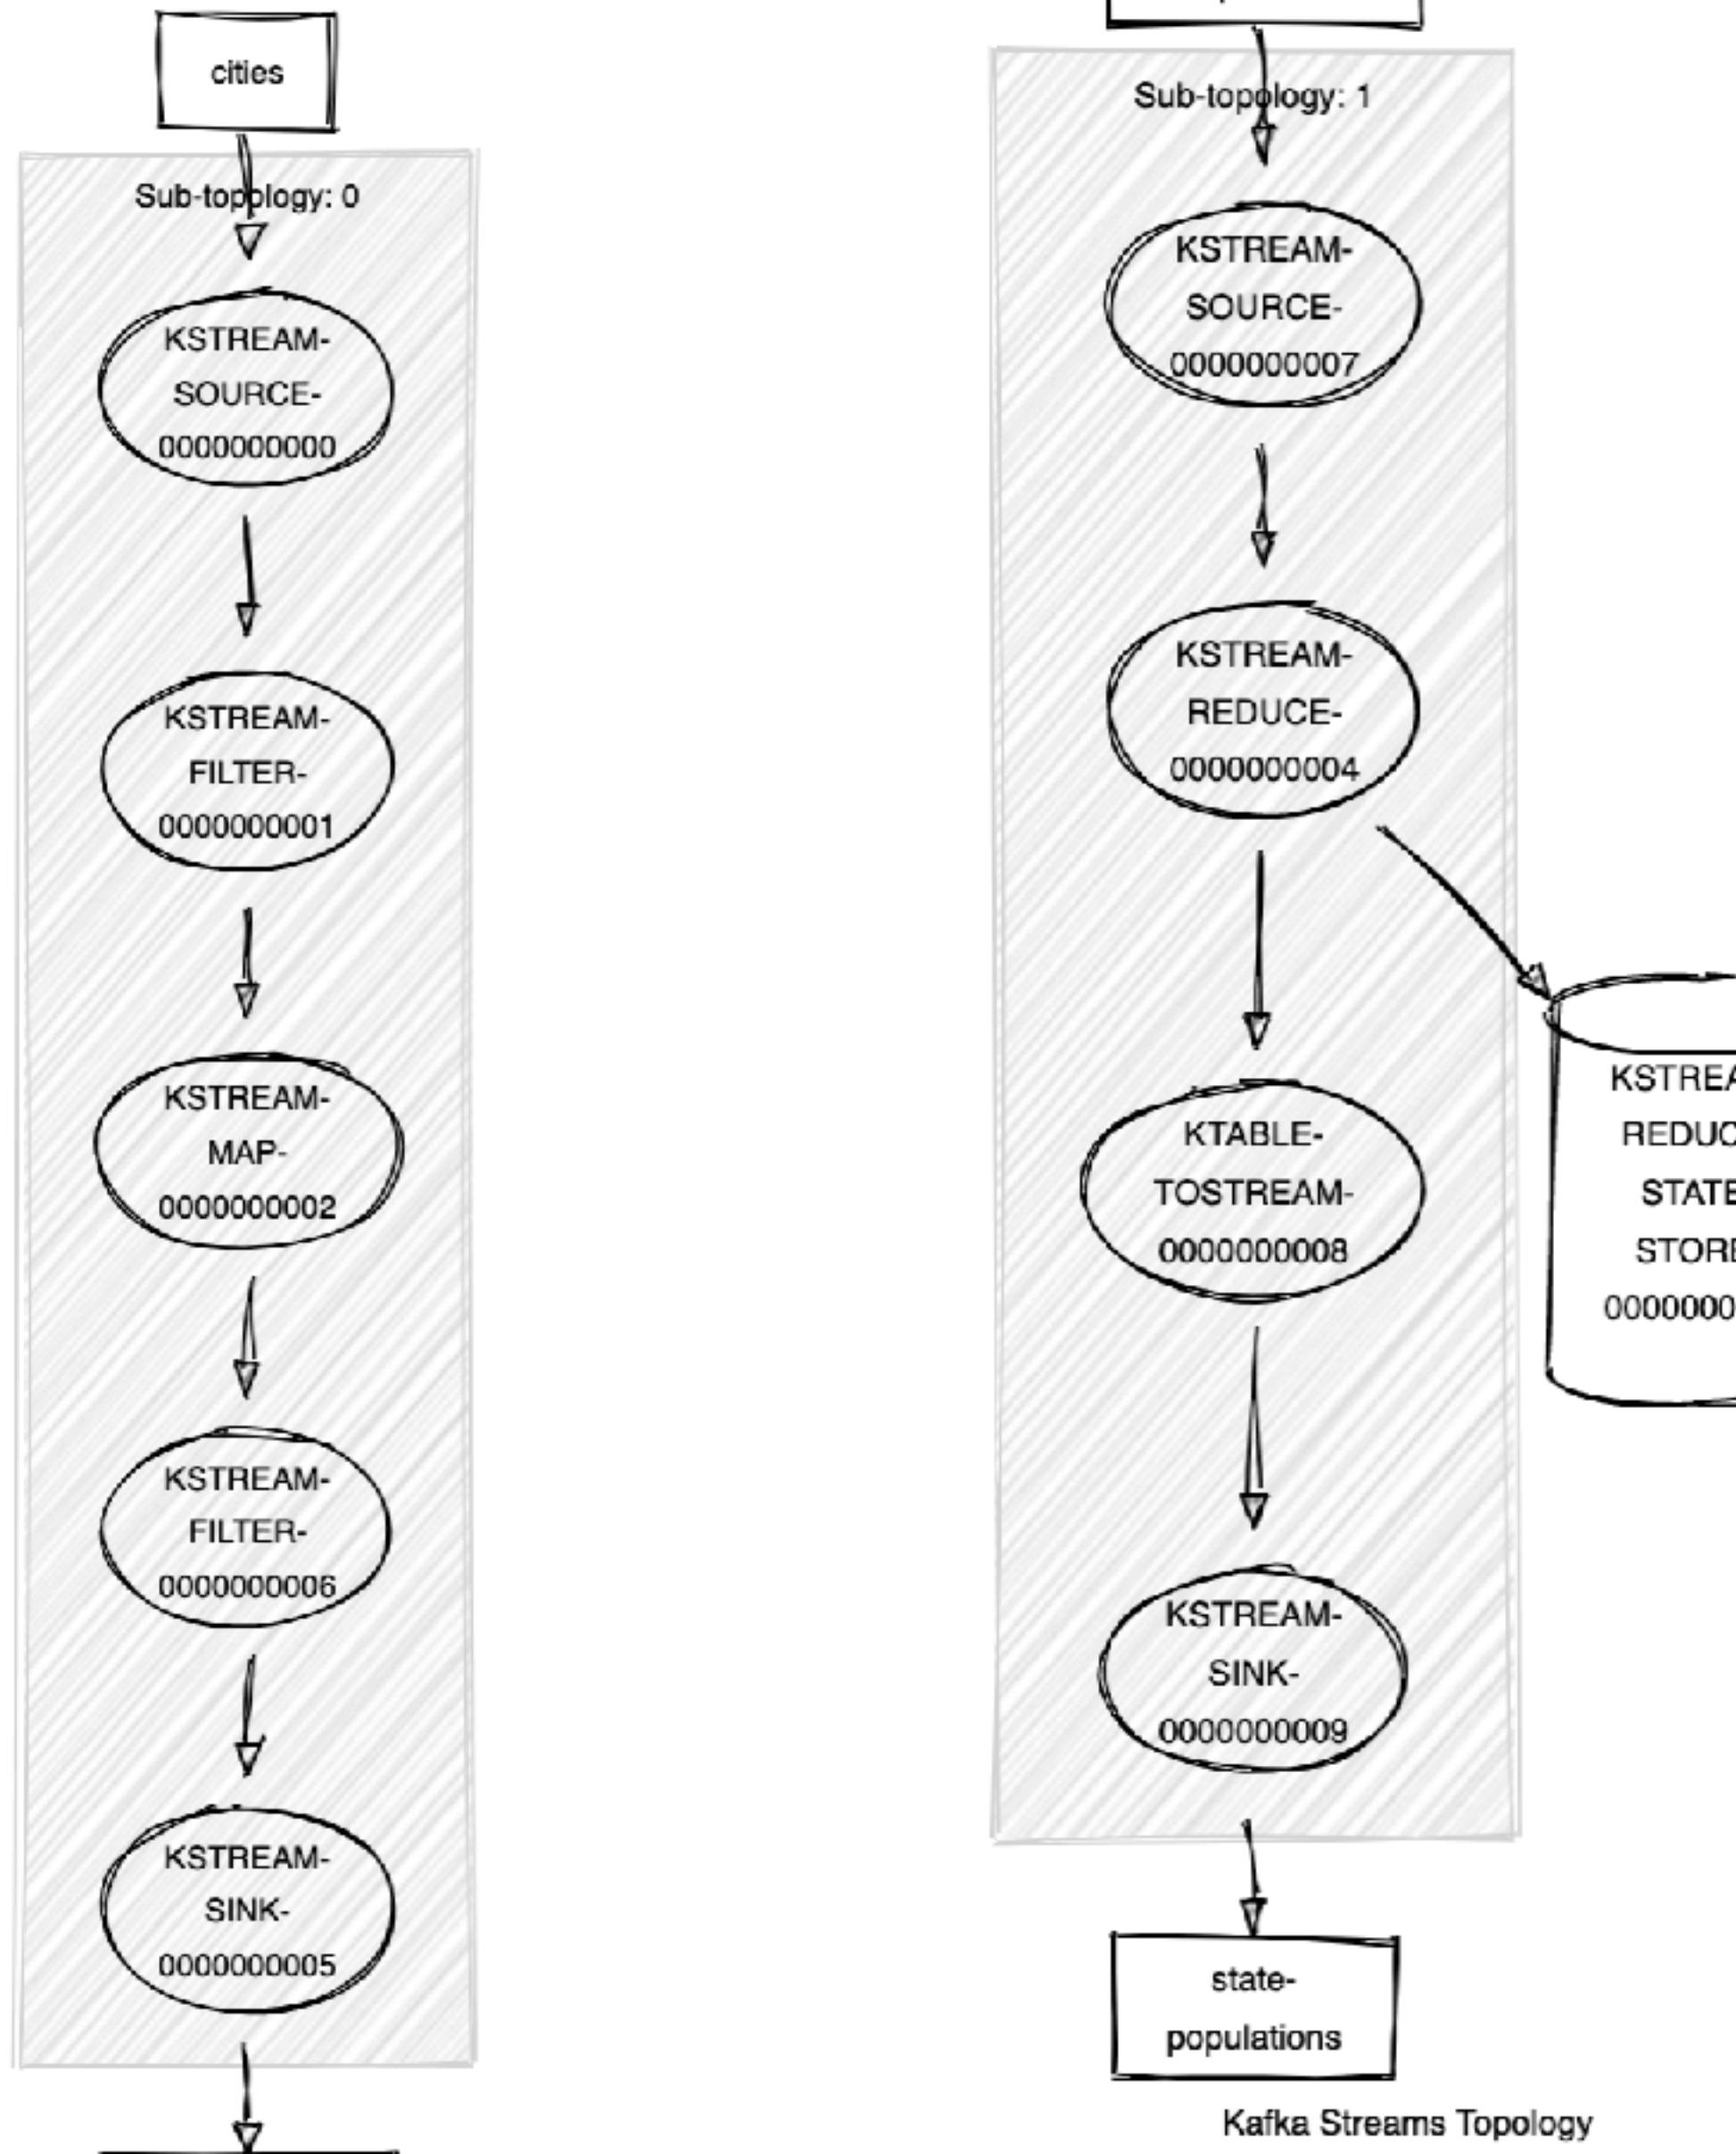
# Kafka Streams Topology Visualizer

Converts an ASCII Kafka Topology description into a hand drawn diagram. [Github link.](#)

Input Kafka Topology [update](#)

```
Topologies:  
Sub-topology: 0  
Source: KSTREAM-SOURCE-0000000000 (topics: [cities])  
--> KSTREAM-FILTER-0000000001  
Processor: KSTREAM-FILTER-0000000001 (stores: [])  
--> KSTREAM-MAP-0000000002  
<-> KSTREAM-SOURCE-0000000000  
Processor: KSTREAM-MAP-0000000002 (stores: [])  
--> KSTREAM-FILTER-0000000006  
<-> KSTREAM-FILTER-0000000001  
Processor: KSTREAM-FILTER-0000000006 (stores: [])  
--> KSTREAM-SINK-0000000005  
<-> KSTREAM-MAP-0000000002  
Sink: KSTREAM-SINK-0000000005 (topic: KSTREAM-REDUCE-STATE-STORE-0000000003-repartition)  
<-> KSTREAM-FILTER-0000000006  
  
Sub-topology: 1  
Source: KSTREAM-SOURCE-0000000007 (topics: [KSTREAM-REDUCE-STATE-STORE-0000000003-repartition])  
--> KSTREAM-REDUCE-0000000004  
Processor: KSTREAM-REDUCE-0000000004 (stores: [KSTREAM-REDUCE-STATE-STORE-0000000003])  
--> KTABLE-TOSTREAM-0000000008  
<-> KSTREAM-SOURCE-0000000007  
Processor: KTABLE-TOSTREAM-0000000008 (stores: [])
```

Output Sketch Diagram



<https://zz85.github.io/kafka-streams-viz/>

# Count Events in a Stream

```
builder.stream(inputTopic, Consumed.with(Serdes.String(), ticketSaleSerde))
    .map((k, v) -> new KeyValue<>(v.getTitle(), v))
    .groupByKey(Grouped.with(Serdes.String(), ticketSaleSerde))
    .count()
    .toStream()
    .mapValues(v -> v.toString() + " tickets sold")
    .to(outputTopic, Produced.with(Serdes.String(), Serdes.String()));
```

# Split Events Into Multiple Streams

```
KStream<String, ActingEvent>[] branches =  
    builder.<String,ActingEvent>stream(inputTopic)  
    .branch( (key, appearance) -> "drama".equals(appearance.getGenre()) ,  
             (key, appearance) -> "fantasy".equals(appearance.getGenre()) ,  
             (key, appearance) -> true) ;  
  
branches[0].to(allProps.getProperty("output.drama.topic.name")) ;  
branches[1].to(allProps.getProperty("output.fantasy.topic.name")) ;  
branches[2].to(allProps.getProperty("output.other.topic.name")) ;
```

# Join a Stream and a Table

```
KStream<String, Movie> movieStream = builder.<String, Movie>stream(movieTopic)
    .map( (key, movie) -> new KeyValue<>(String.valueOf(movie.getId()), movie));
movieStream.to(rekeyedMovieTopic);

KTable<String, Movie> movies = builder.table(rekeyedMovieTopic);

KStream<String, Rating> ratings = builder.<String, Rating>stream(ratingTopic)
    .map( (key, rating) -> new KeyValue<>(String.valueOf(rating.getId()), rating));
KStream<String, RatedMovie> ratedMovie = ratings.join(movies, joiner);
    .to(ratedMoviesTopic, Produced.with(Serdes.String(),
        ratedMovieAvroSerde(allProps)));

```

# Join a Stream and a Table

## (ValueJoiner)

```
public class MovieRatingJoiner implements ValueJoiner<Rating, Movie, RatedMovie> {  
  
    public RatedMovie apply(Rating rating, Movie movie) {  
        return RatedMovie.newBuilder()  
            .setId(movie.getId())  
            .setTitle(movie.getTitle())  
            .setReleaseYear(movie.getReleaseYear())  
            .setRating(rating.getRating())  
            .build();  
    }  
}
```

# Count Events by Time Window

```
builder.<String, Rating>stream(ratingTopic)
    .map((key, rating) -> new KeyValue<>(rating.getTitle(), rating))
    .groupByKey()
    .windowedBy(TimeWindows.of(Duration.ofMinutes(10)))
    .count()
    .toStream()
    .map((Windowed<String> key, Long count) -> new KeyValue<>(key.key(),
                                                                count.toString()))
    .to(ratingCountTopic, Produced.with(Serdes.String(), Serdes.String()));
```

# Exactly Once Semantics

In Kafka Streams app configuration:

```
processing.guarantee = "exactly_once"
```

Producer config:

```
enable.idempotence = true  
max.inflight.requests.per.connection = 1  
retries = Integer.MAX_VALUE  
transactional.id
```

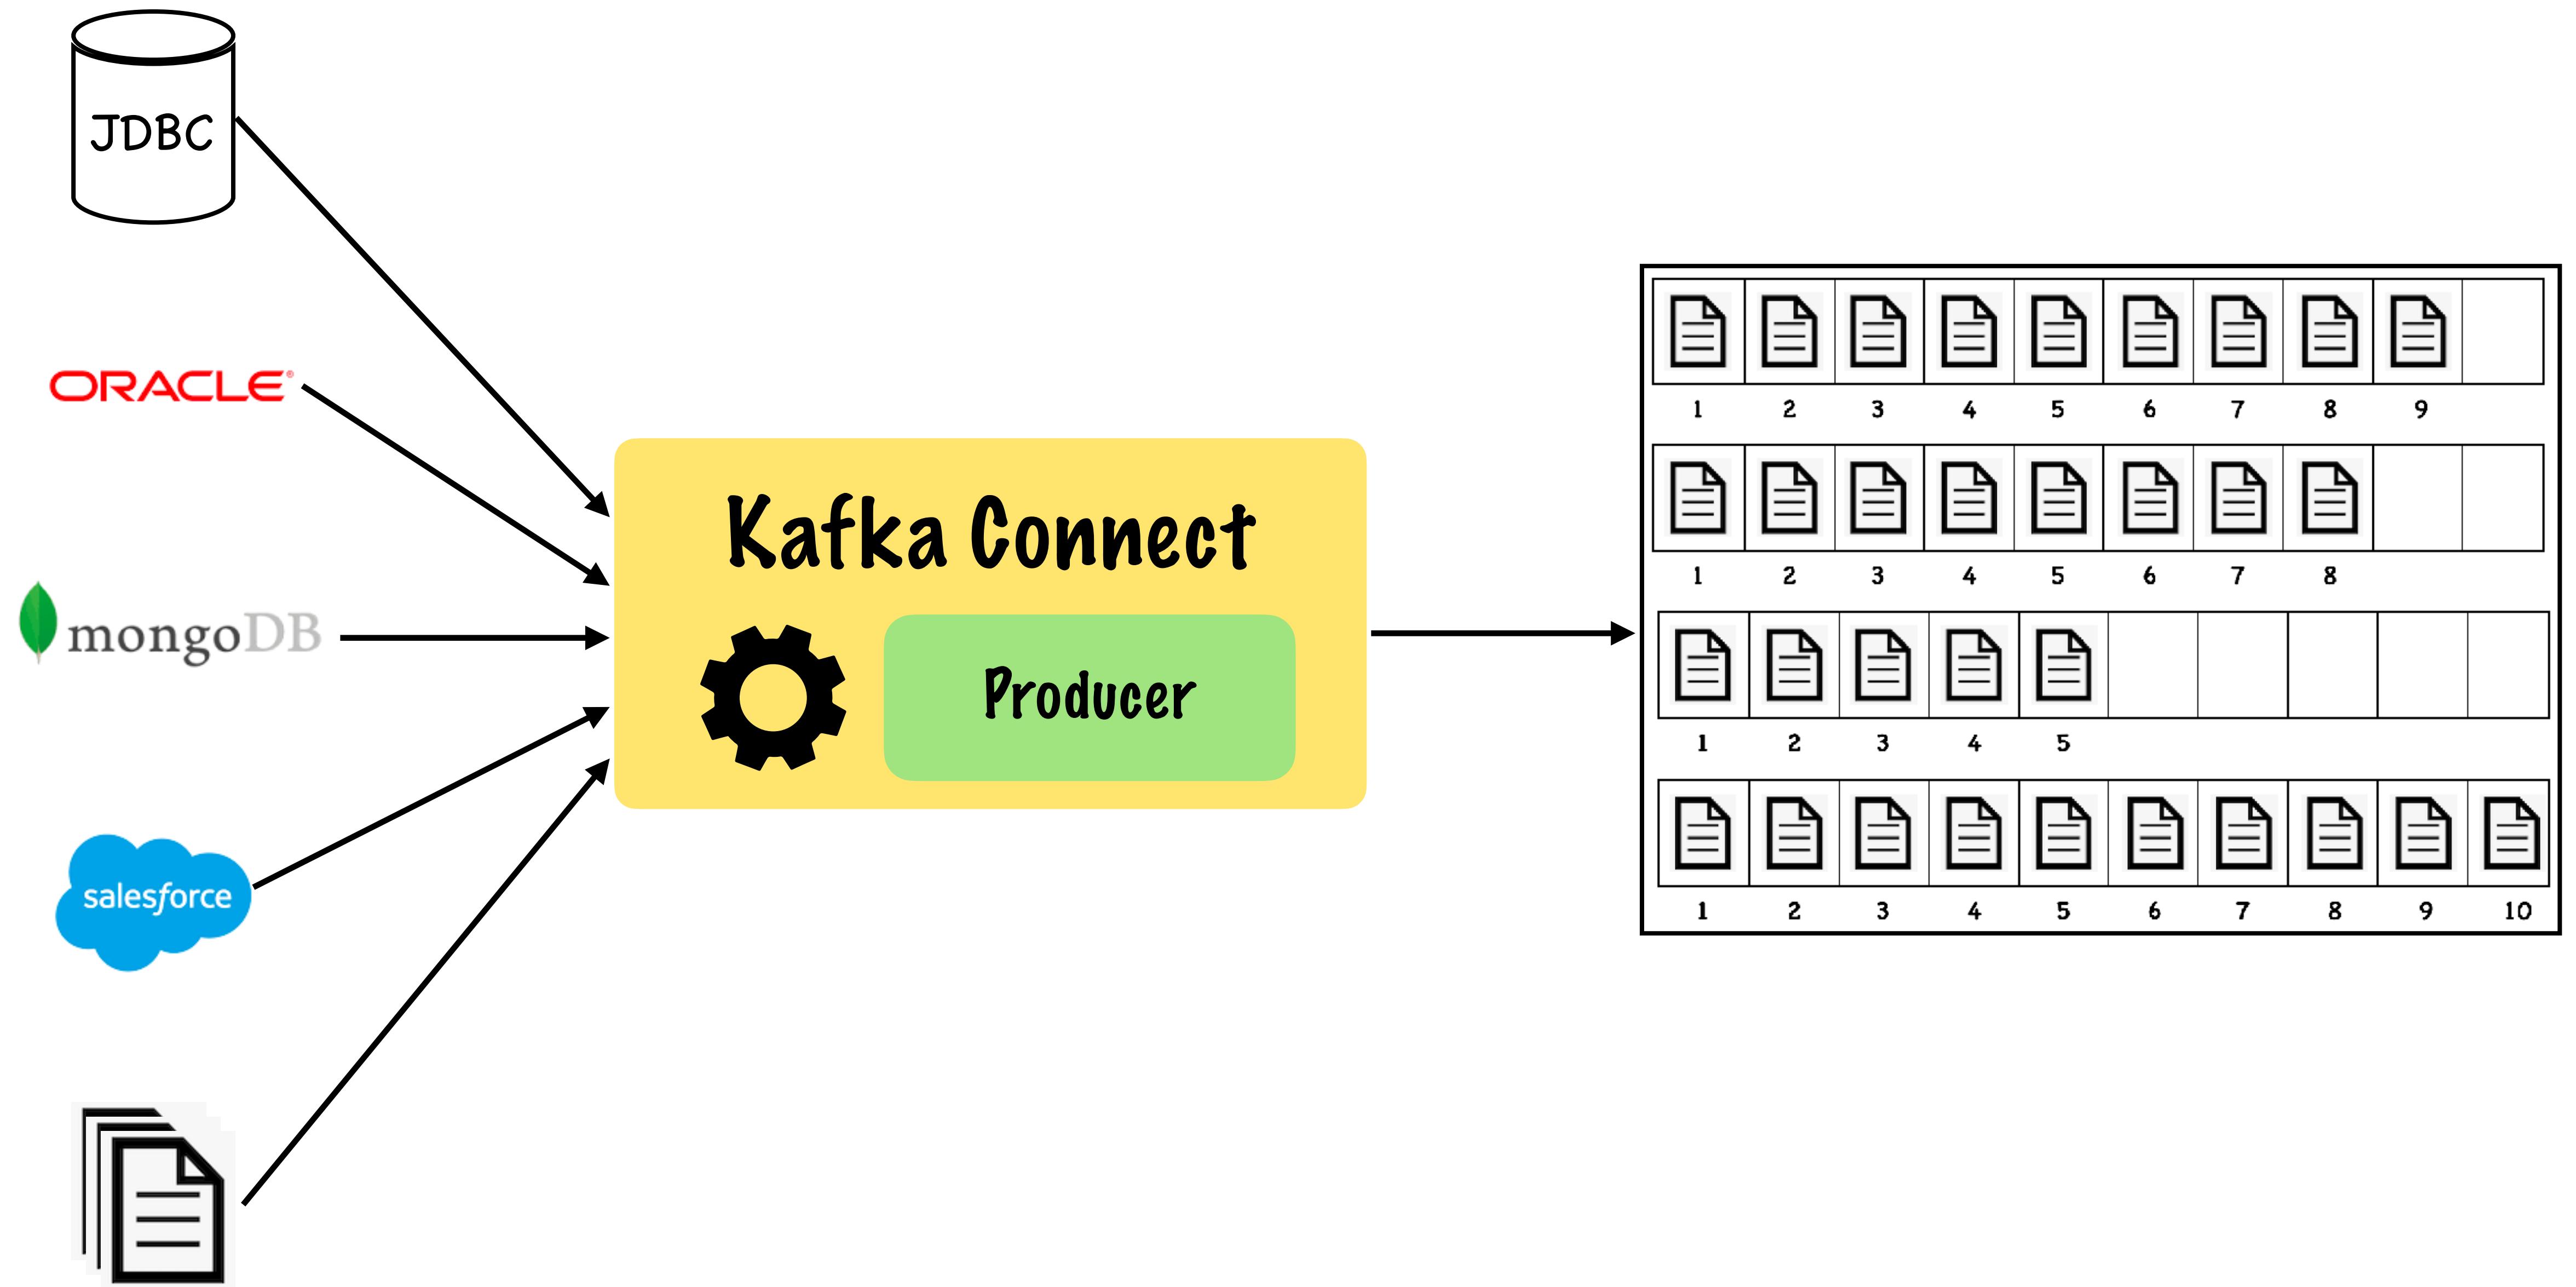
Consumer config:

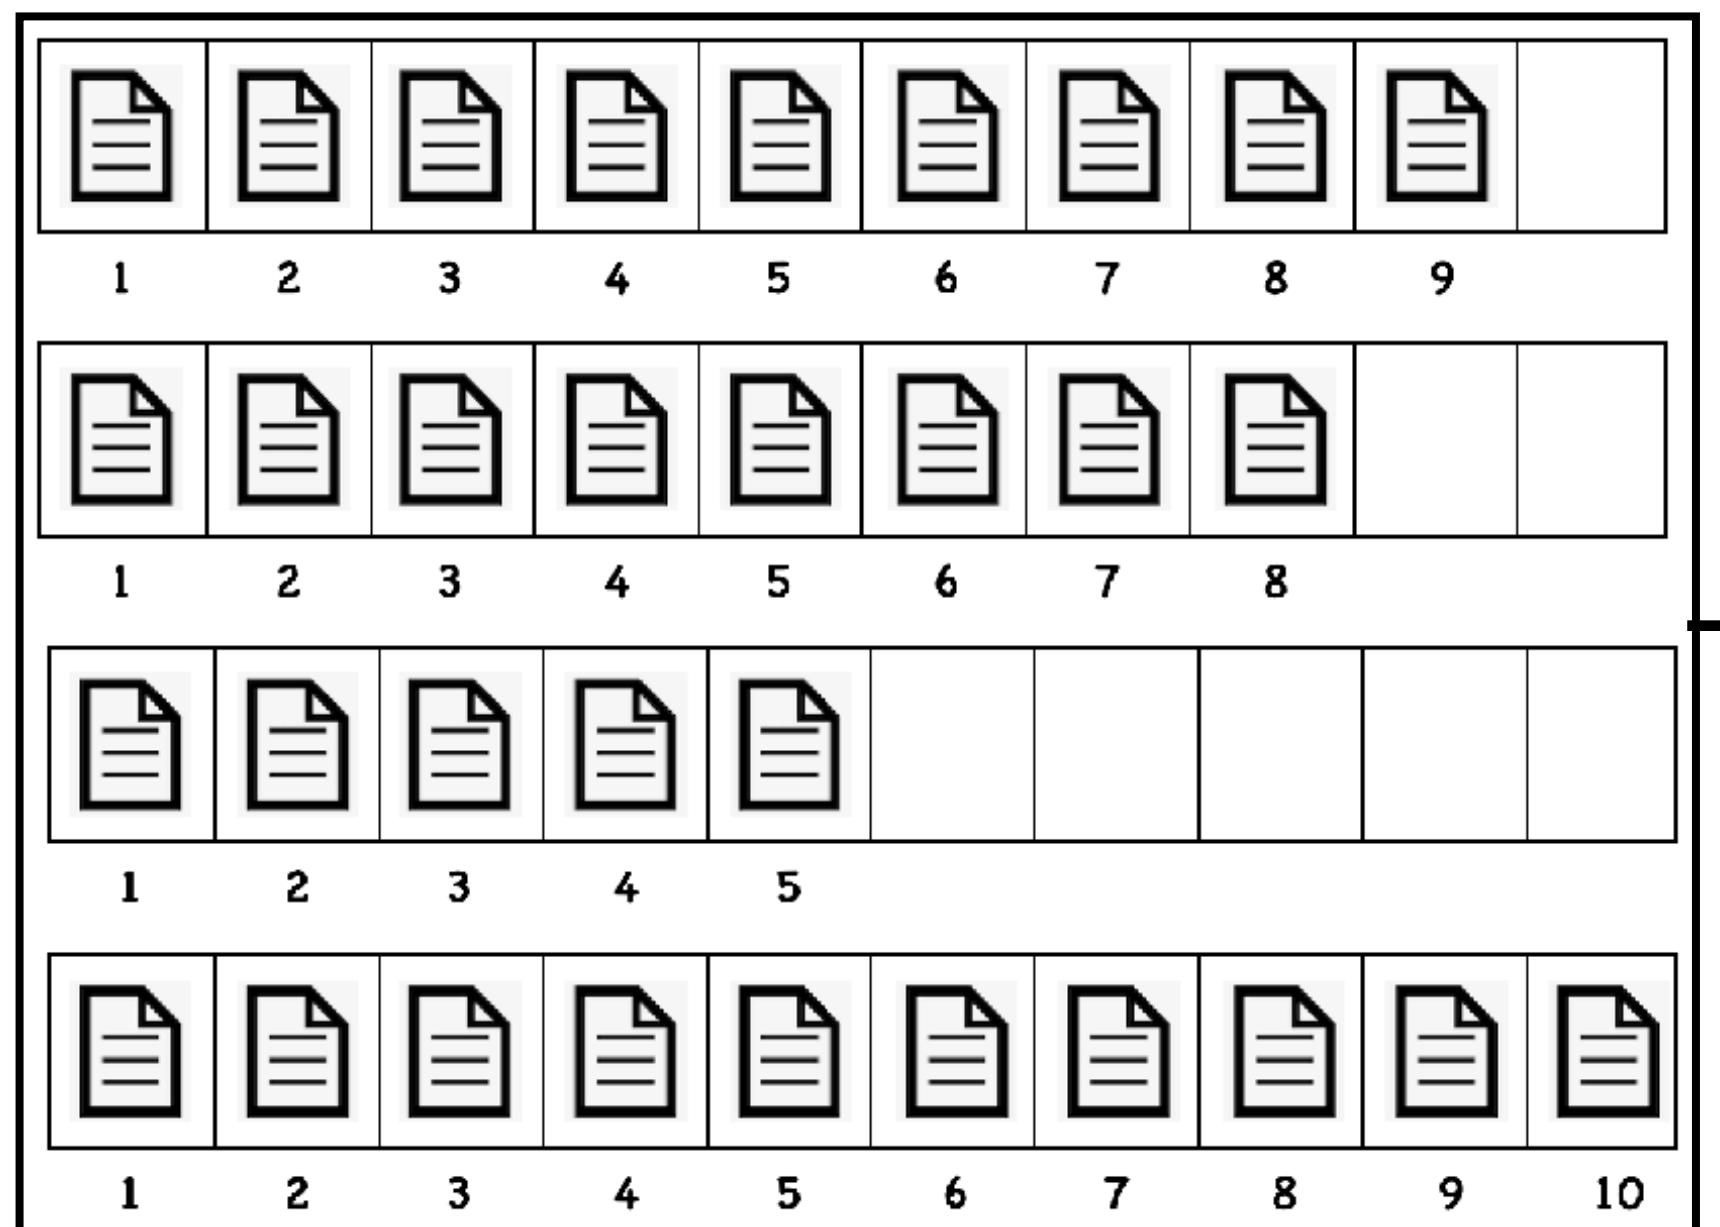
```
isolation.level = "read_committed"
```

<https://www.confluent.io/blog/enabling-exactly-once-kafka-streams>

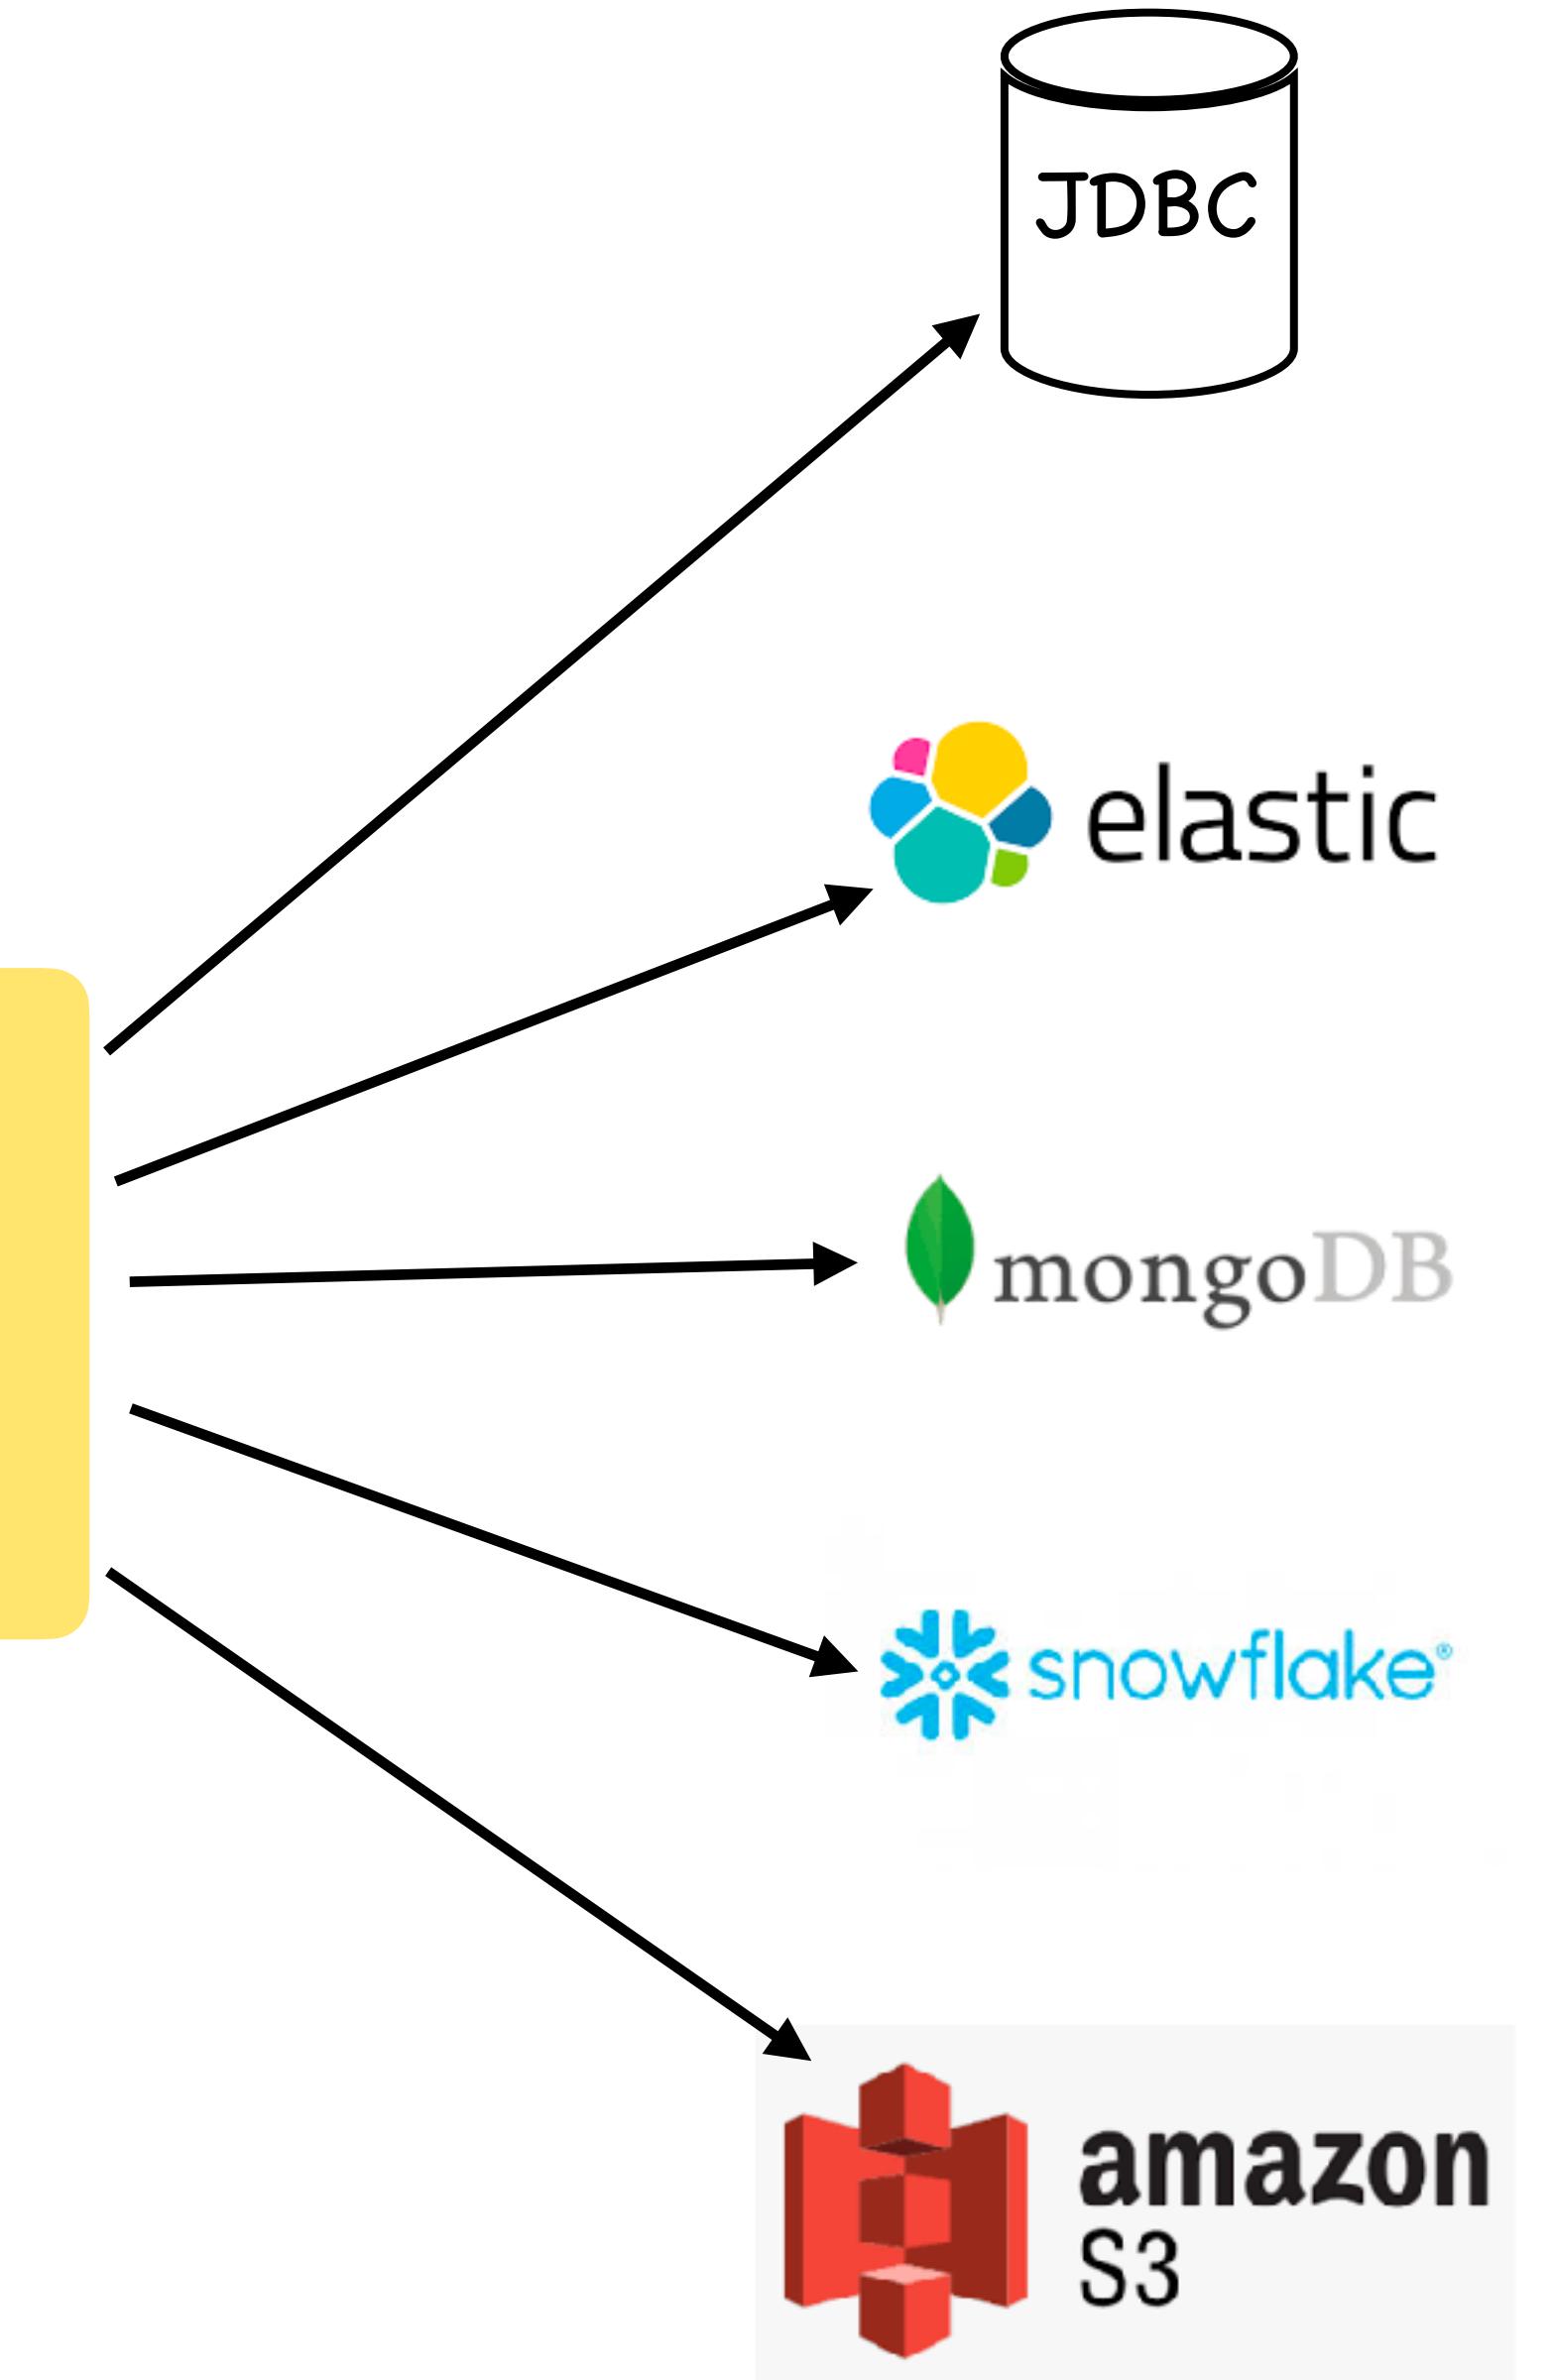
# Exercises Part 2

# More Kafka Tools and Resources





[confluent.io/hub](https://confluent.io/hub)





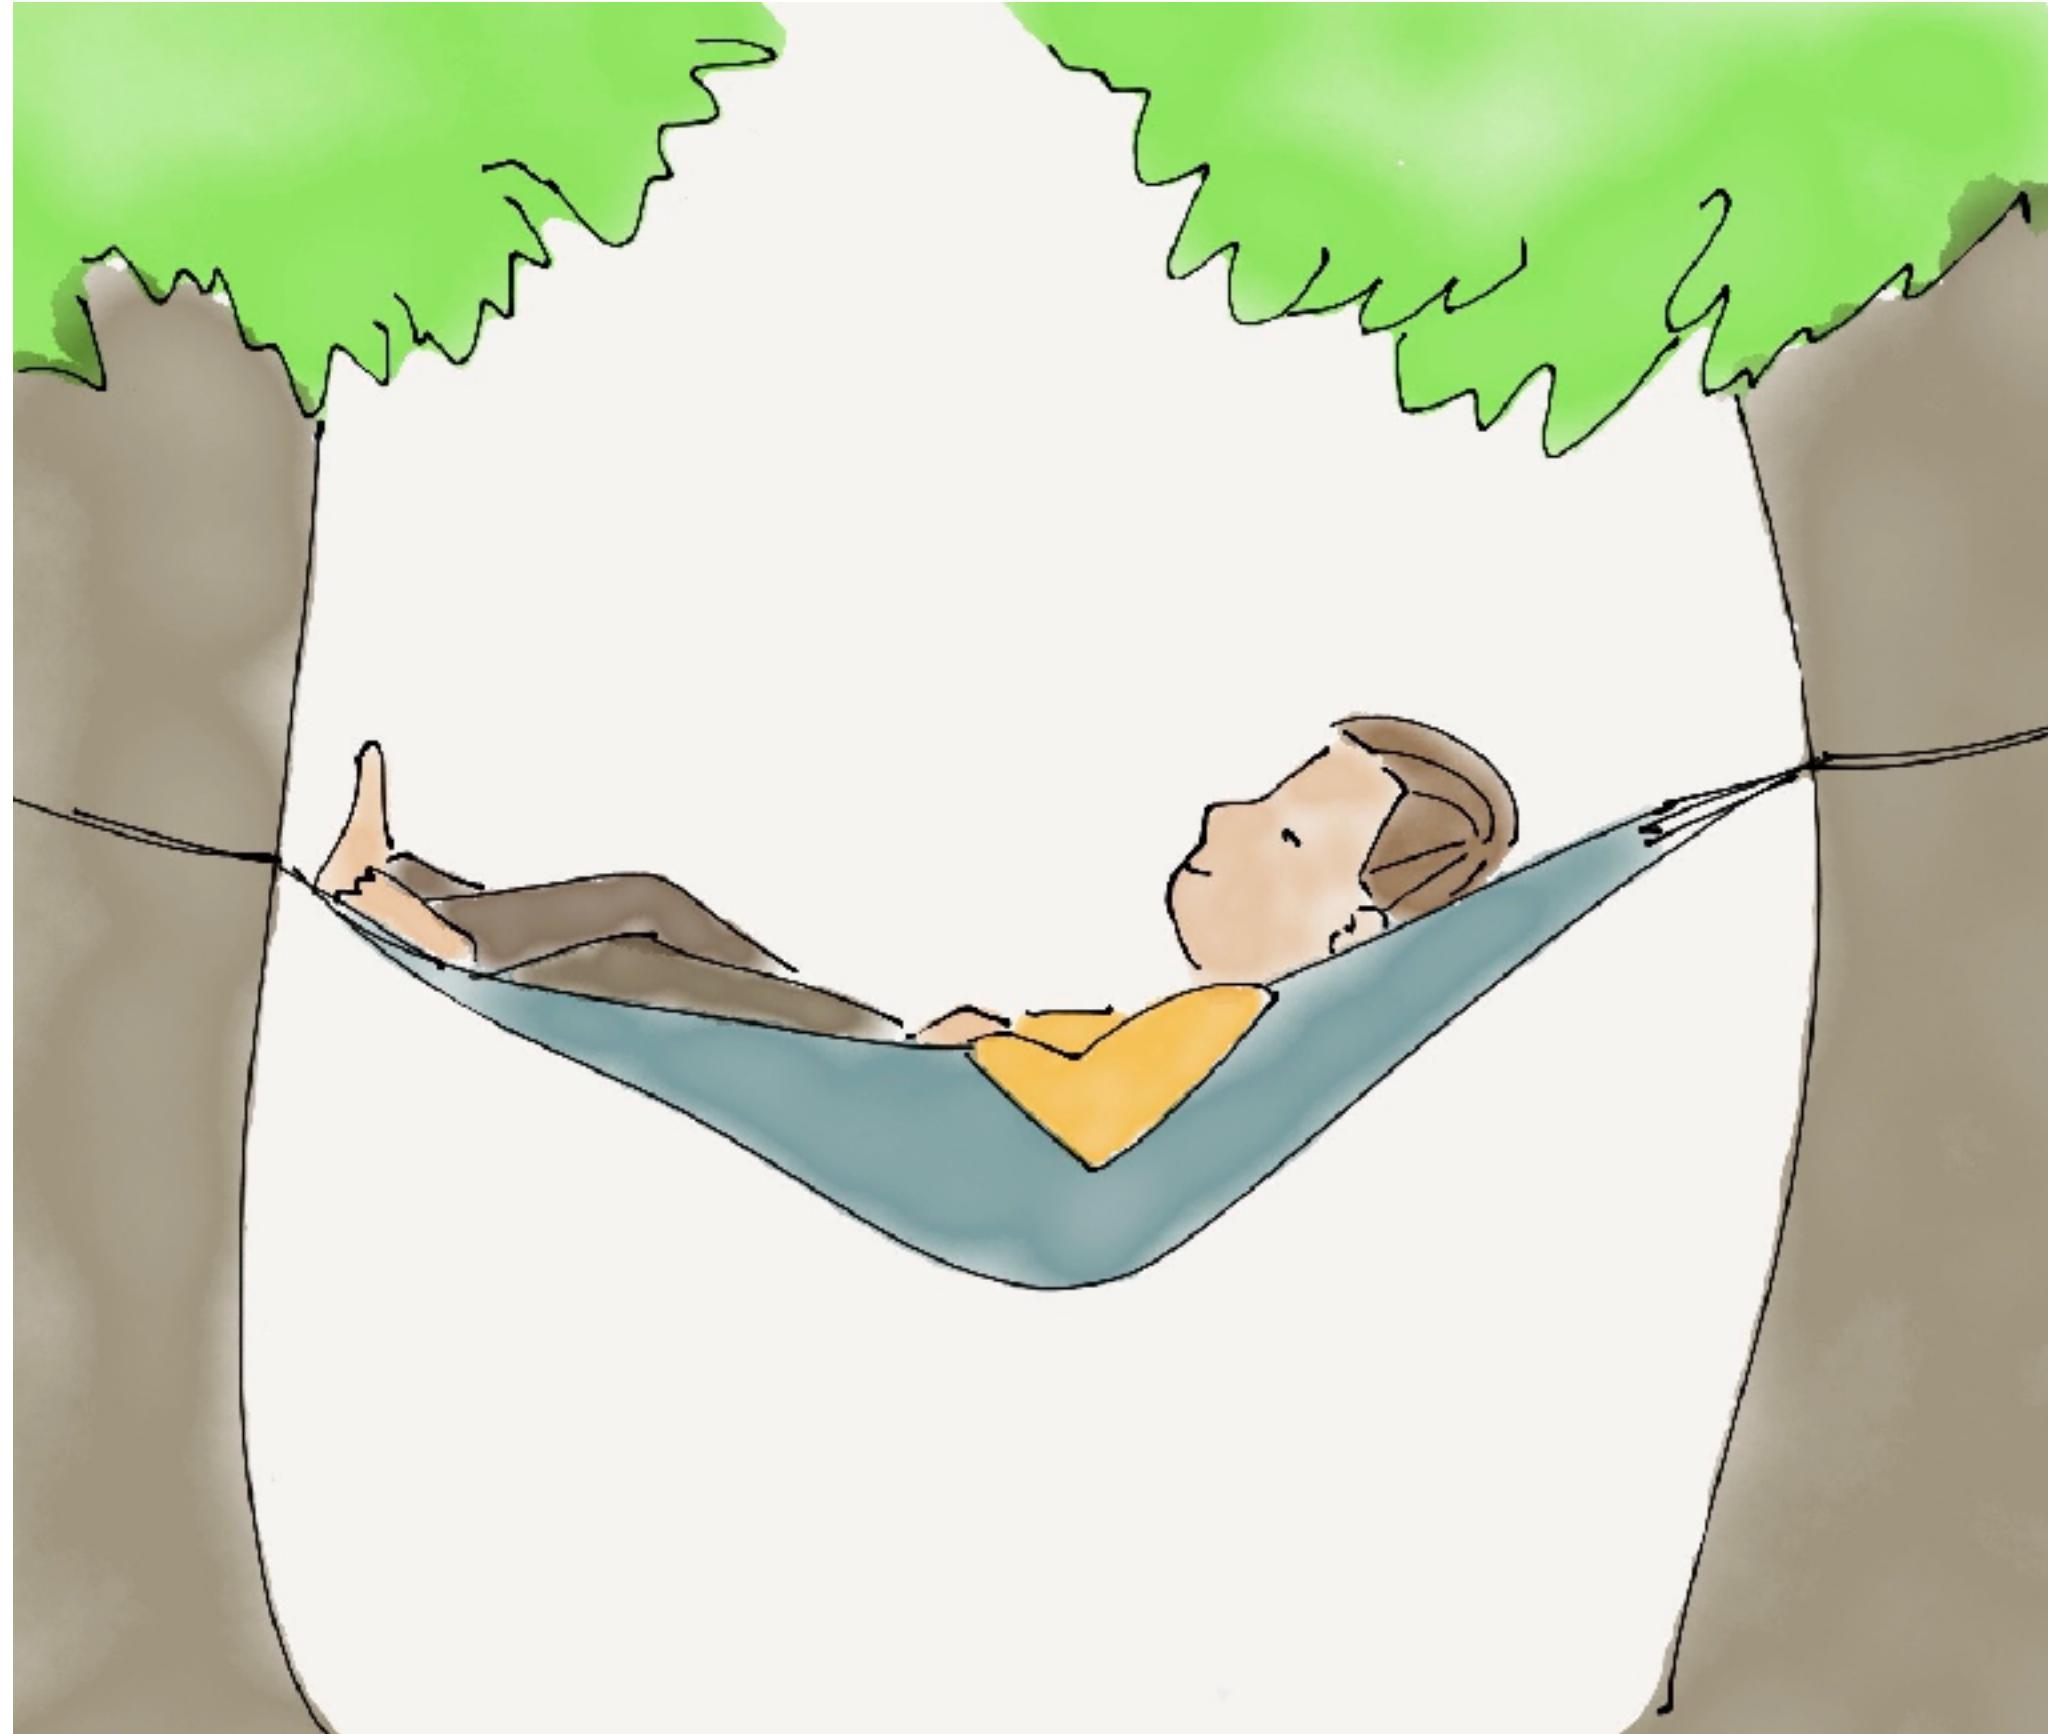
# ksqlDB

```
ksql> CREATE STREAM books (
    title VARCHAR, author VARCHAR, year_published
INT
) WITH (
    KAFKA_TOPIC='booklist', VALUE_FORMAT='AVRO'
);
```

```
ksql> SELECT * FROM books WHERE title LIKE '%kafka%' EMIT CHANGES;
```

TITLE	AUTHOR	YEAR
Kafka: The Definitive Guide	Shapira, Narkhede, Palino	2017
Kafka Streams in Action	Bill Bejeck	2018
Kafka in Action	Dylan Scott	2021
Apache Kafka Cookbook	Raul Estrada	2017

# Confluent REST API



```
curl -X POST -H "Content-Type: application/vnd.kafka.json.v2+json" \
-H "Accept: application/vnd.kafka.v2+json" \
--data '{"records":[{"value":{"foo":"bar"} }]}' \
"http://localhost:8082/topics/jsontest"
```

# Command Line Tools

## Kafka Shell Scripts

- Start and stop brokers
- Create / manage topics
- Produce / consume events

Confluent CLI - <https://docs.confluent.io/confluent-cli/current/overview.html>

- Start / stop services in Confluent Platform / Cloud
- Manage access control and secrets



kcat (formerly Kafkacat) - <https://github.com/edenhill/kcat>

- Produce / consume events in style

kcctl (Casey Cuddle) - <https://github.com/kcctl/kcctl>

- Command-line client for Kafka Connect

## Cluster overview

Data flow

Last 7 days

Topics

Connectors

Consumers

ksqldb

API access

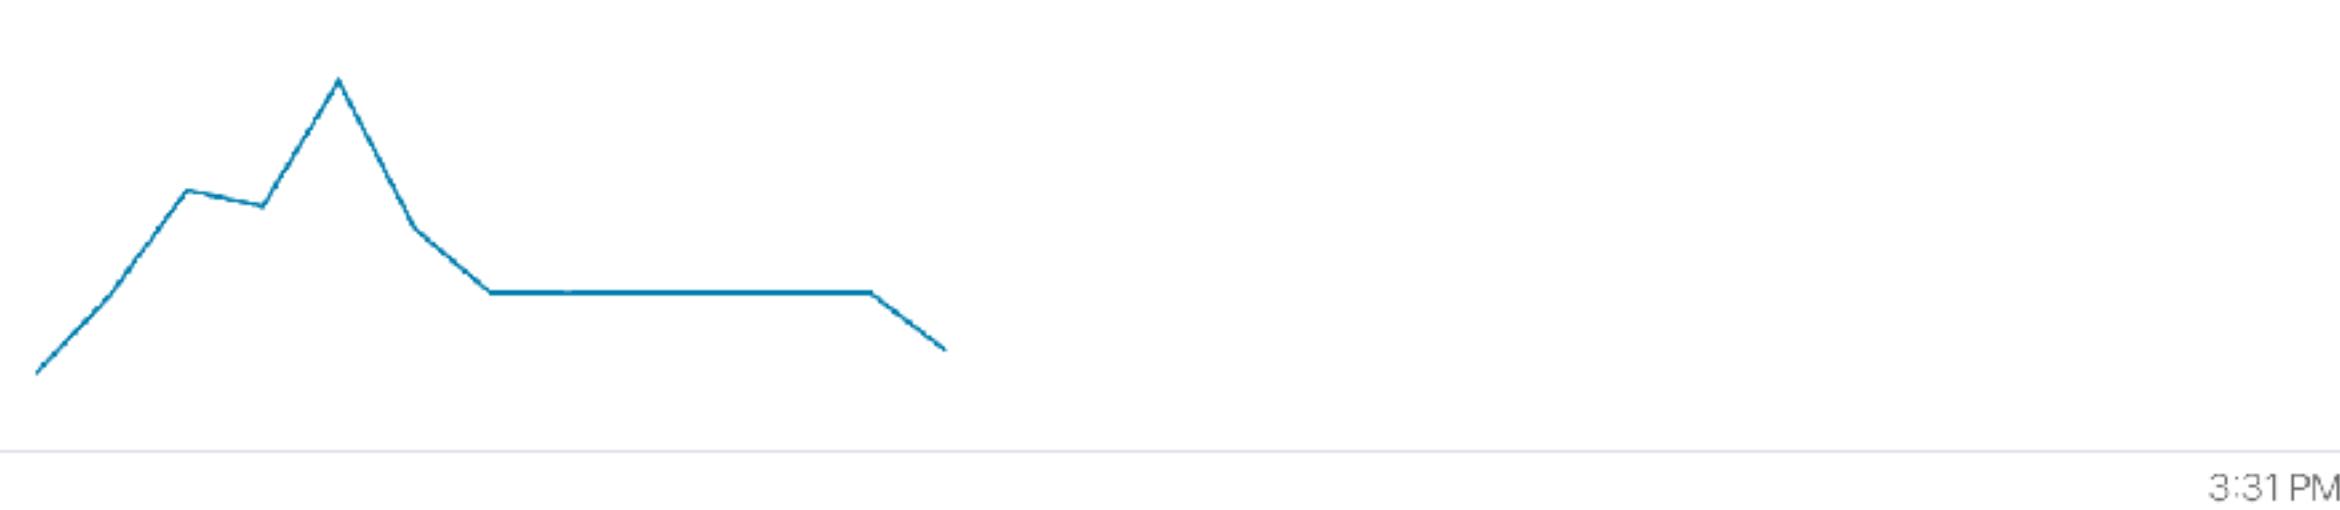
Clients

CLI and Tools

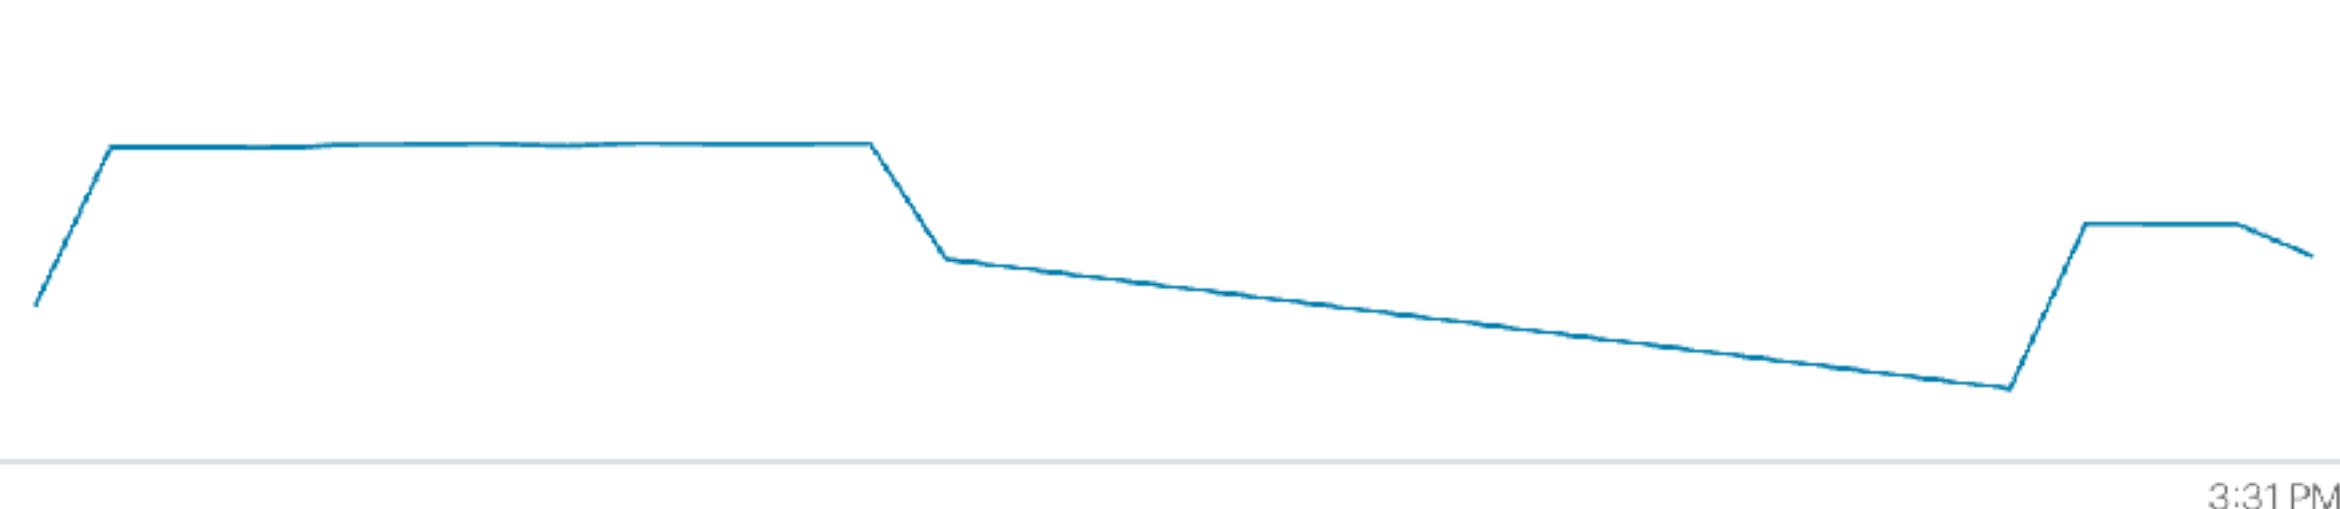
Cluster settings

## Throughput

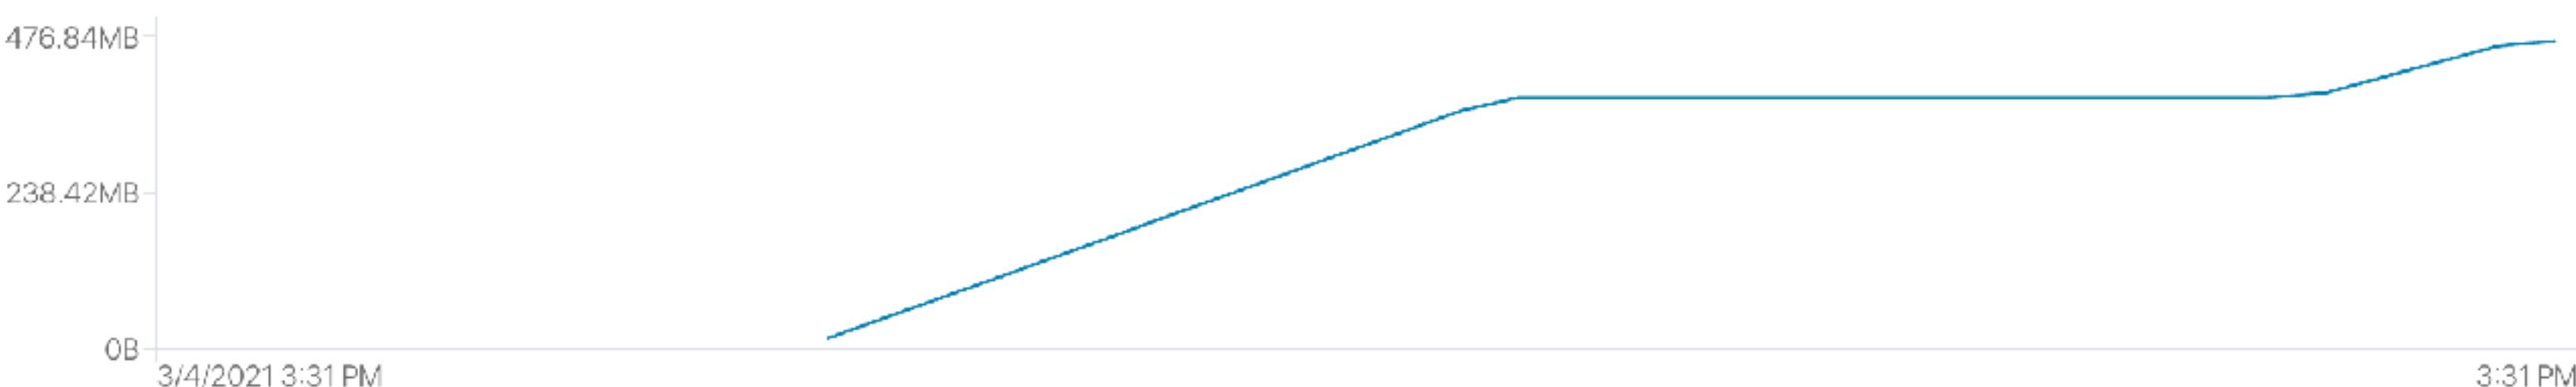
## Consumption (bytes/sec)



## Production (bytes/sec)



## Storage



Cluster overview

Data flow

Topics

Connectors

Consumers

ksqlDB

API access

**Clients**

CLI and Tools

Cluster settings

## ▼ C#

CONFLUENT SUPPORTED

## ▼ C/C++

CONFLUENT SUPPORTED

## ▼ Clojure

## ▼ Go

CONFLUENT SUPPORTED

## ▼ Groovy

## ^ Java

CONFLUENT SUPPORTED

[See example](#) [Copy](#)

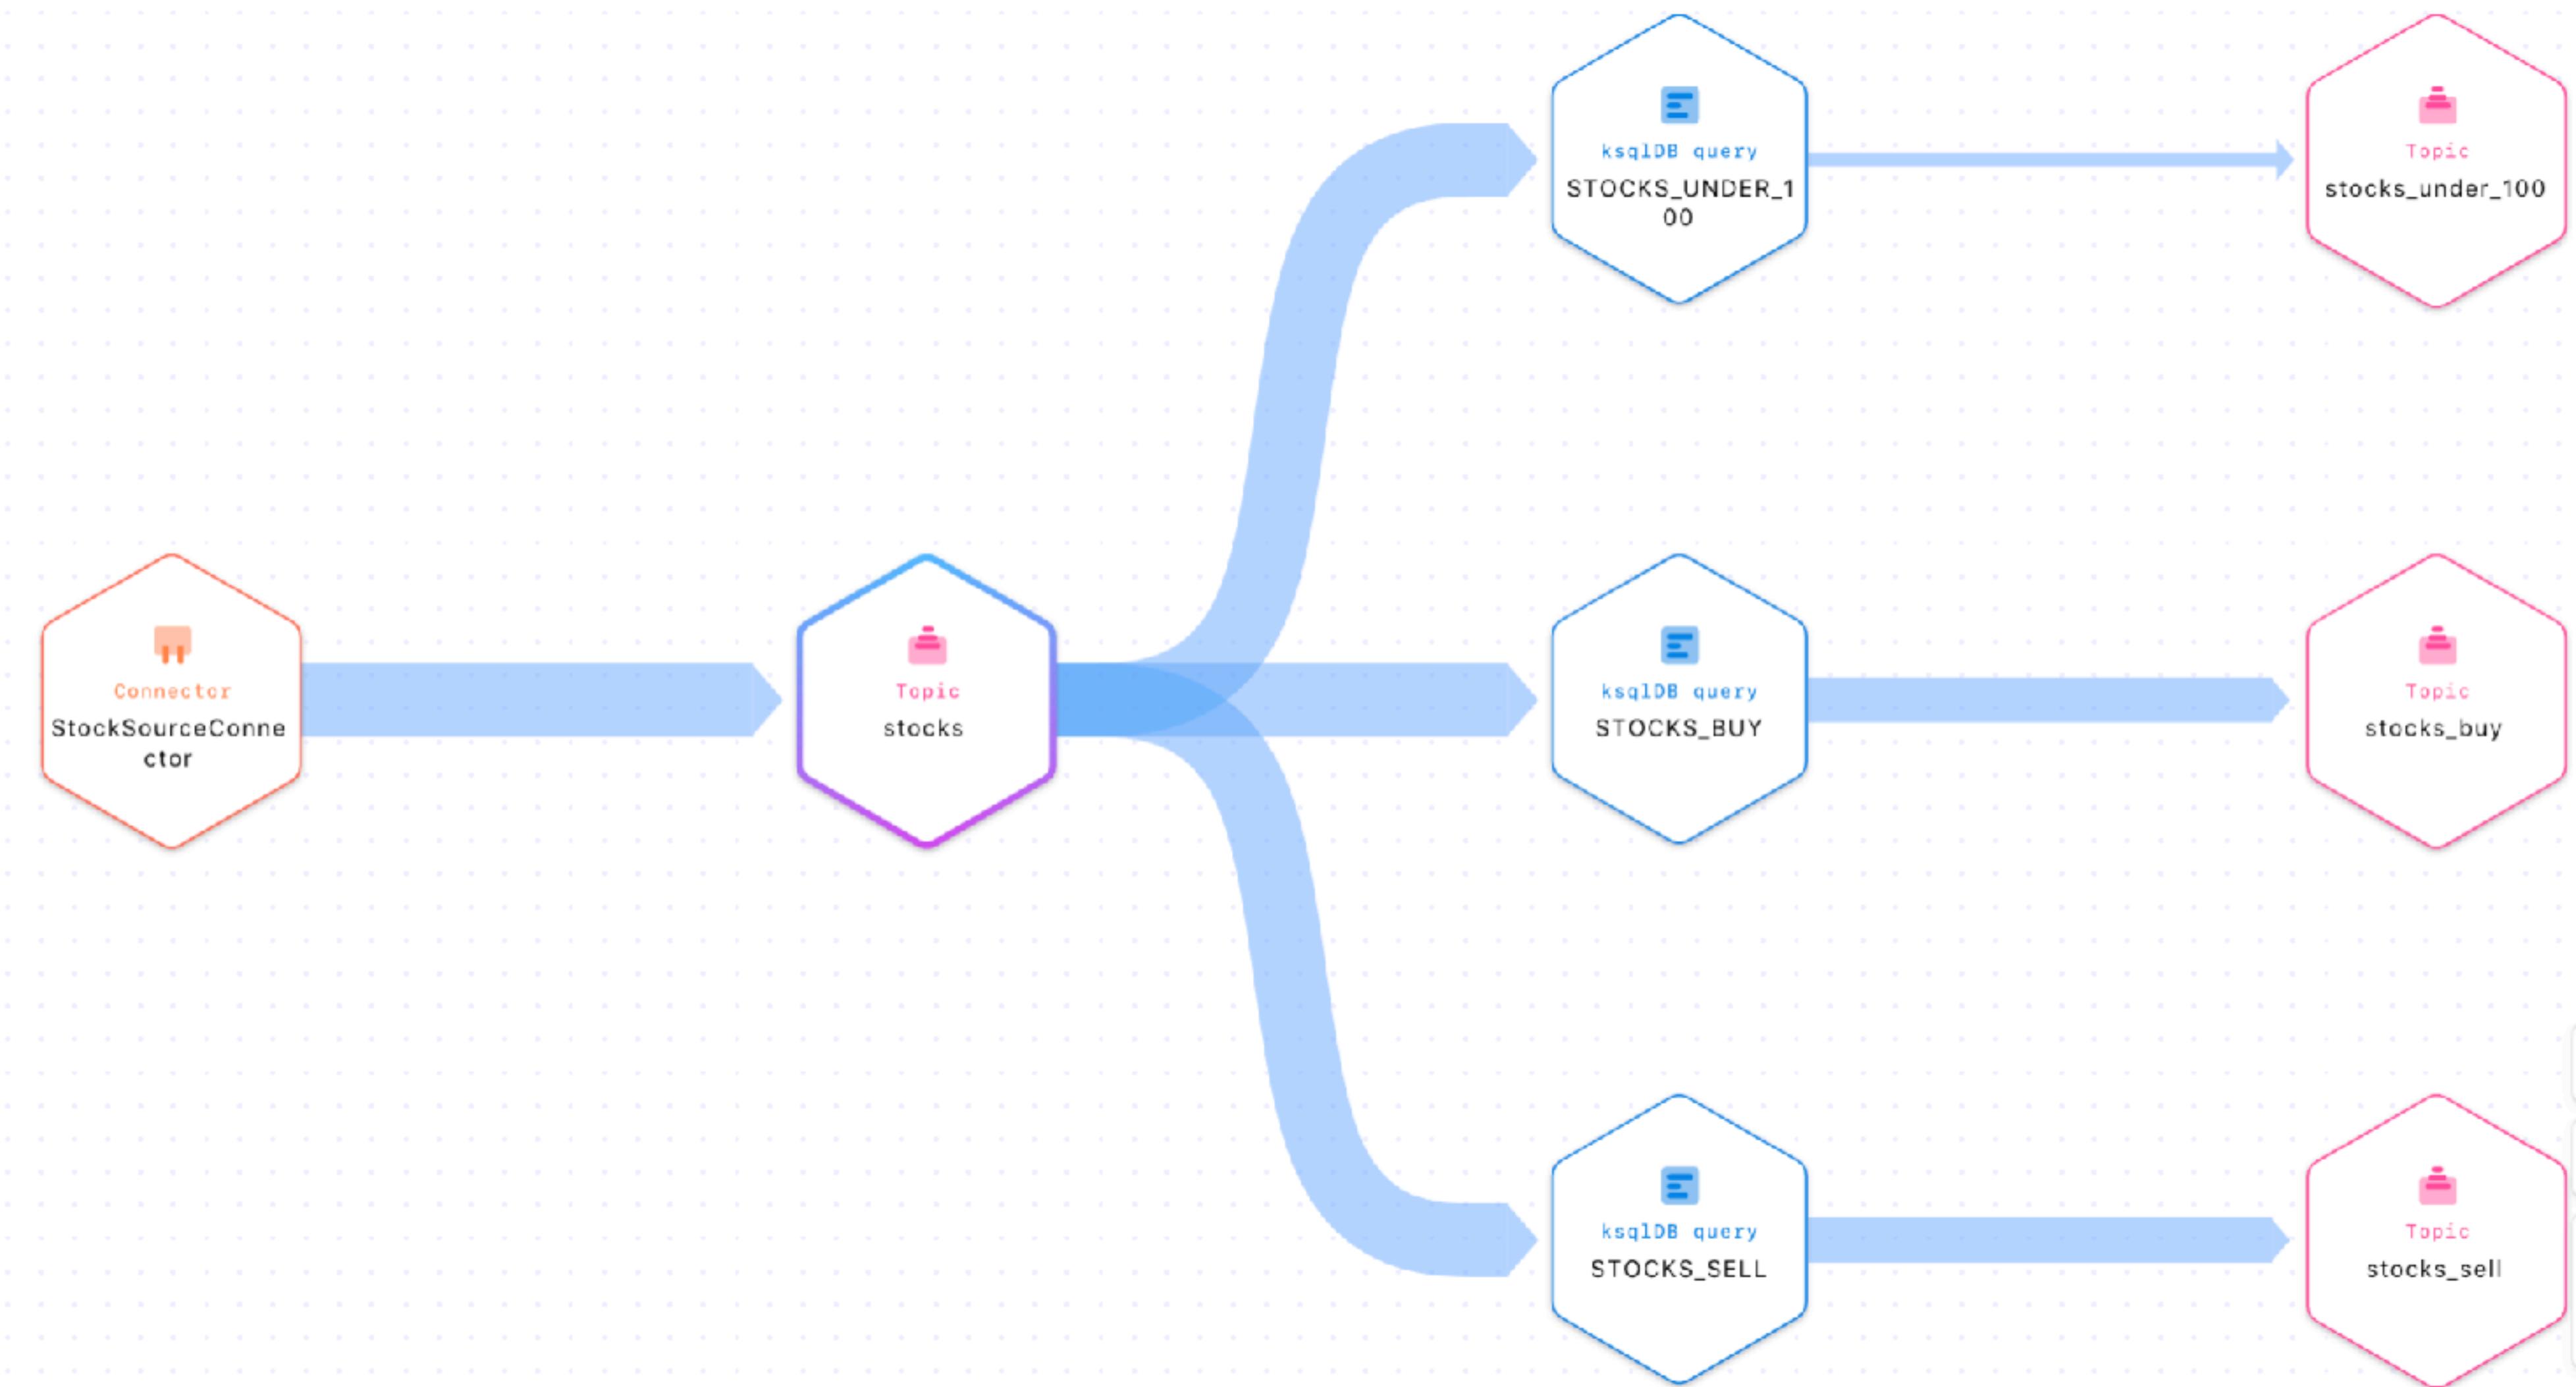
```
1 # Required connection configs for Kafka producer, consumer, and admin
2 bootstrap.servers=pkc-43n10.us-central1.gcp.confluent.cloud:9092
3 security.protocol=SASL_SSL
4 sasl.jaas.config=org.apache.kafka.common.security.plain.PlainLoginModule required username='{{ CLUSTER_API_KEY }}' password='{{ CLUSTER_API_SECRET }}';
5 sasl.mechanism=PLAIN
6 # Required for correctness in Apache Kafka clients prior to 2.6
7 client.dns.lookup=use_all_dns_ips
```

## ▼ Kotlin

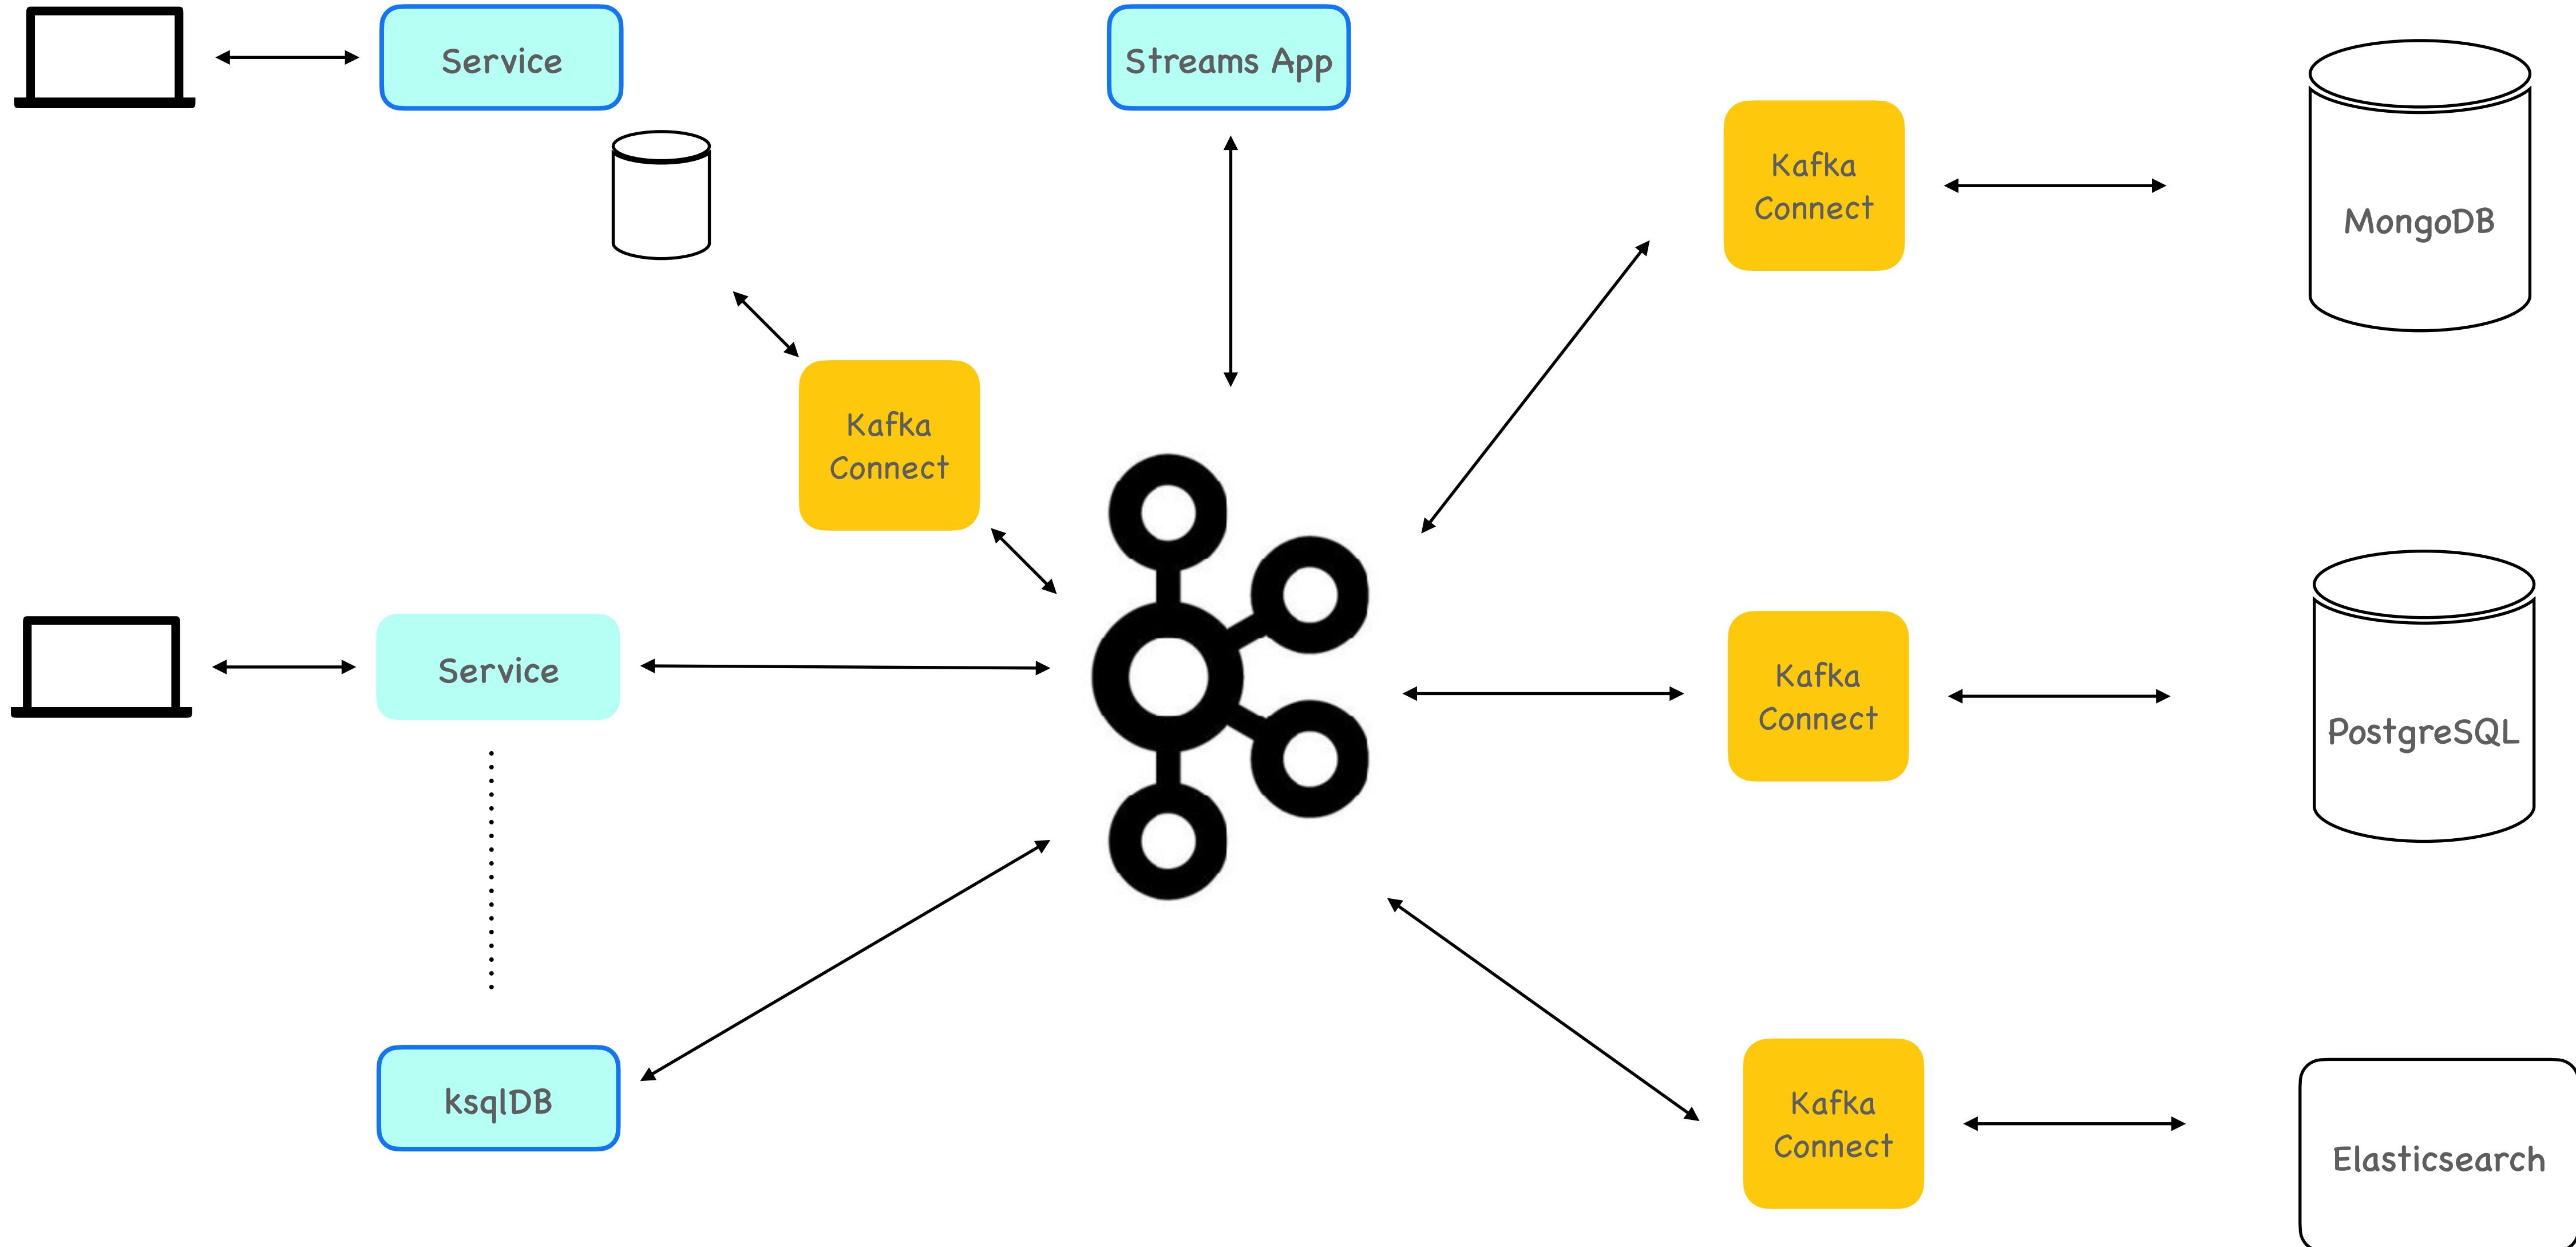
## ▼ Node.js

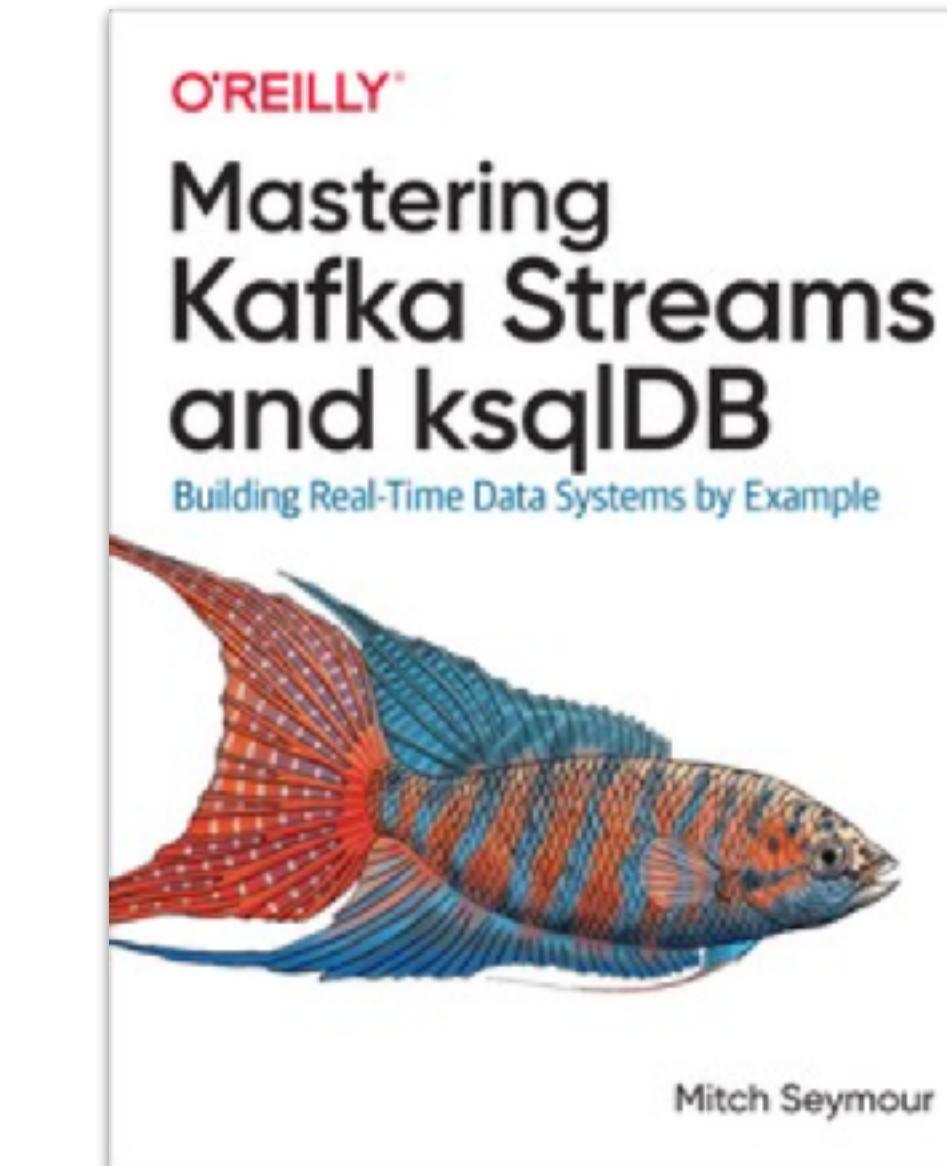
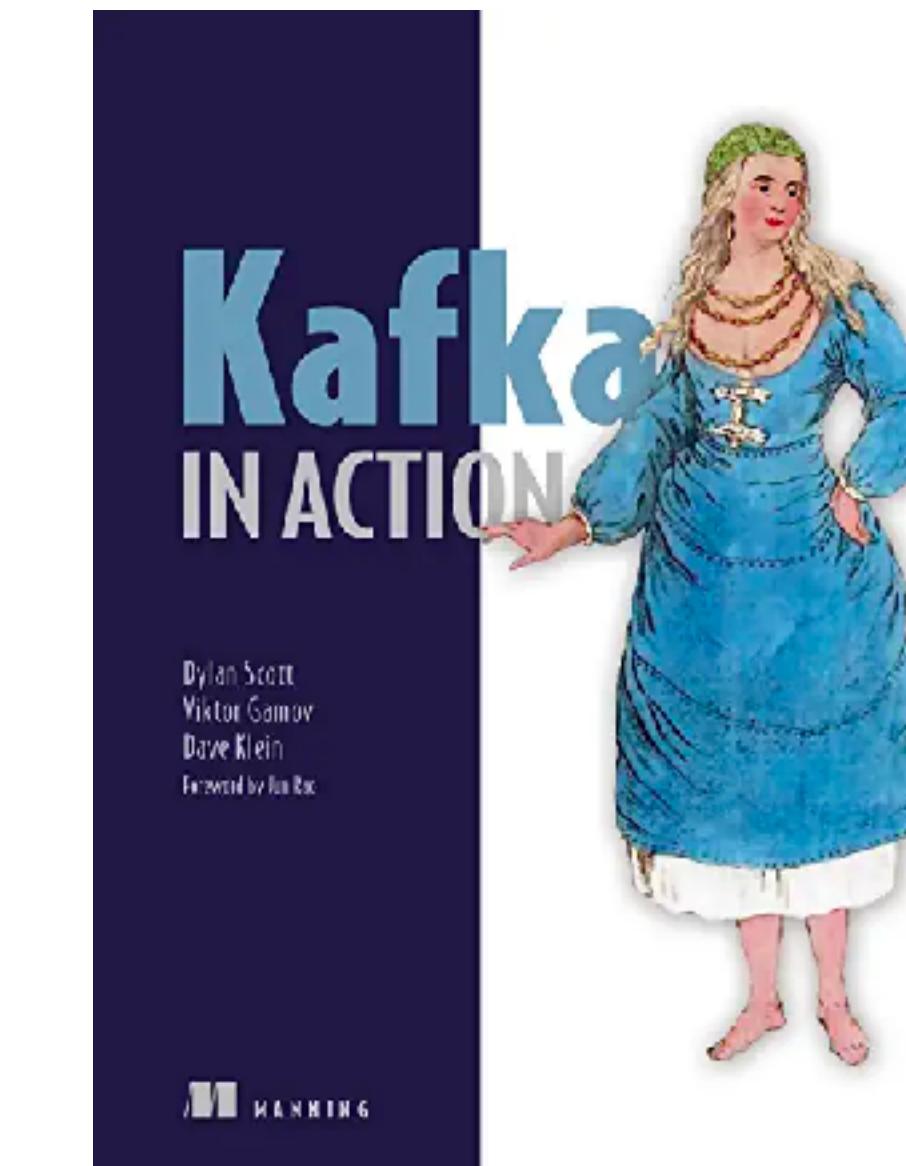
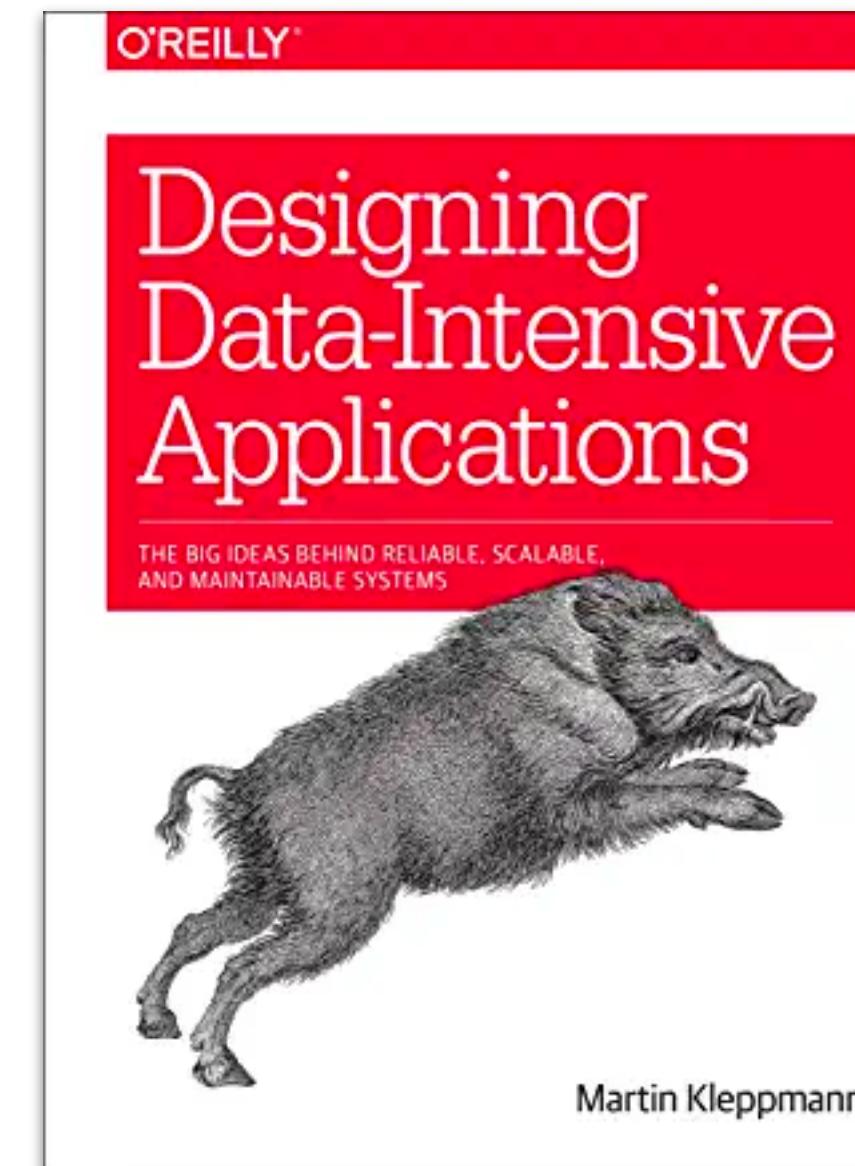
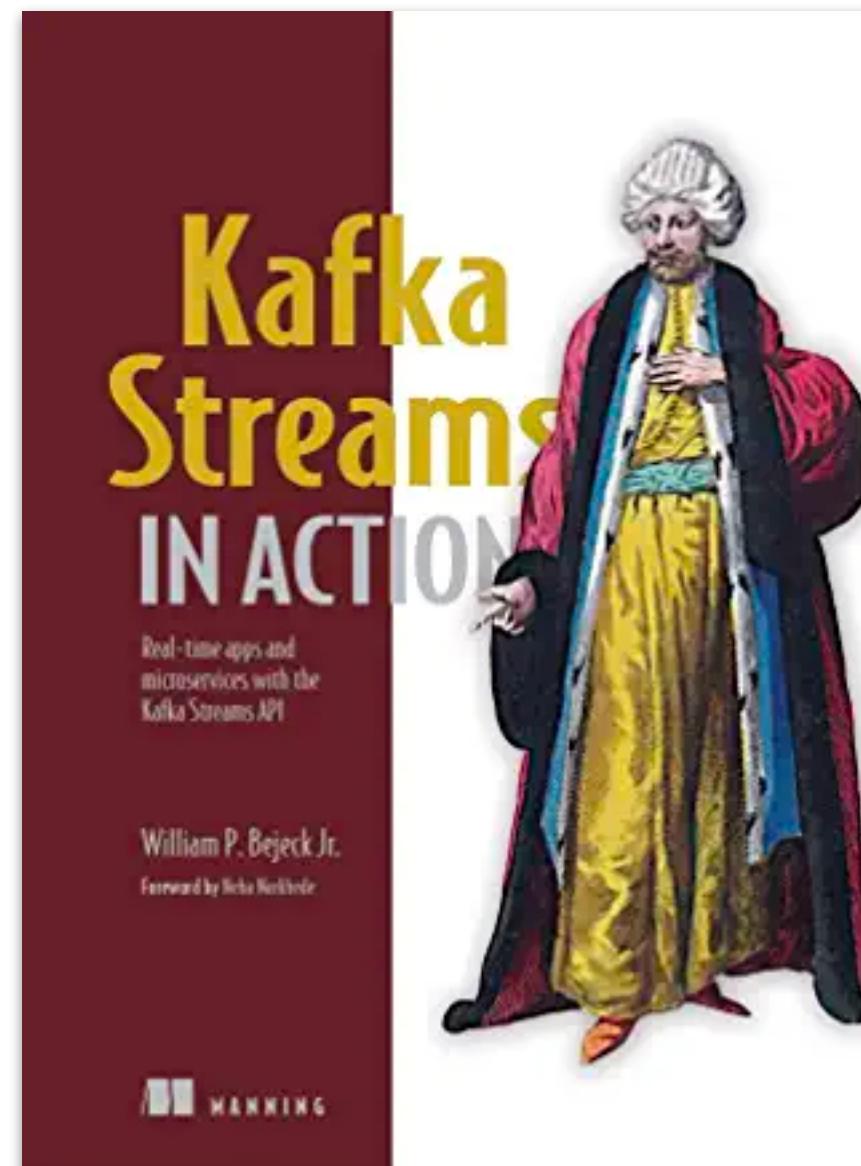
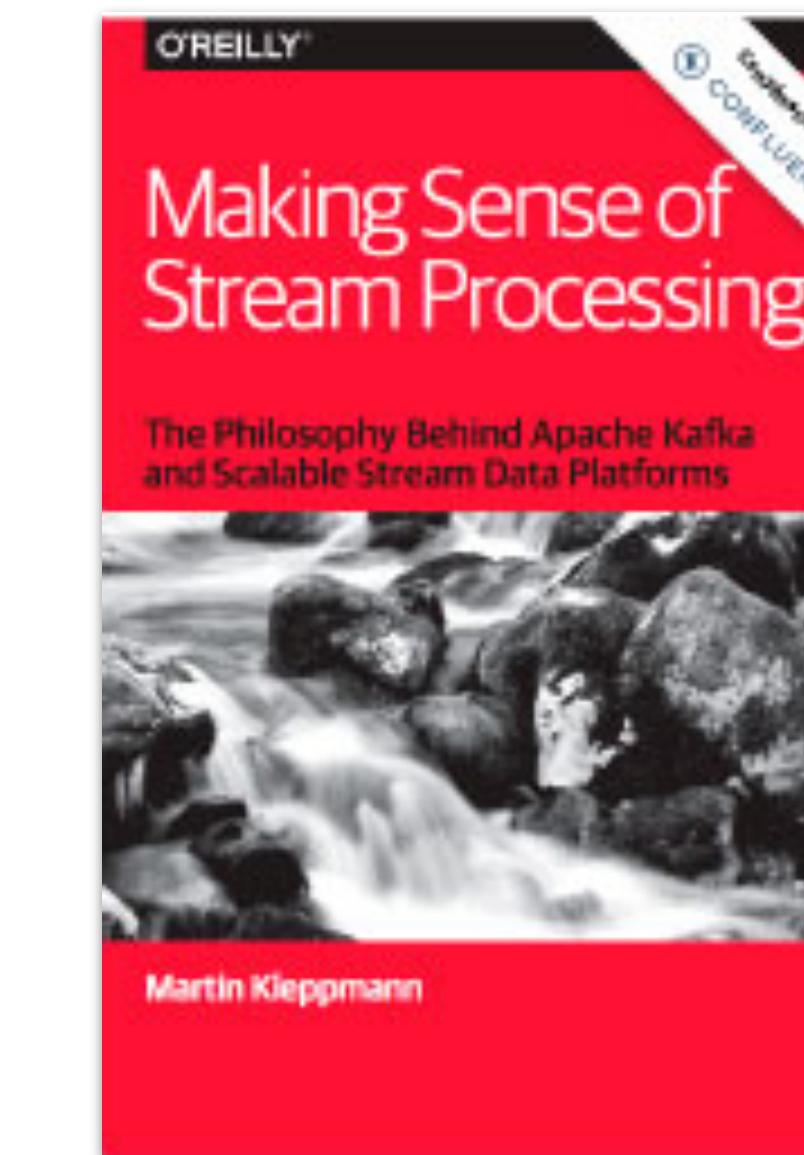
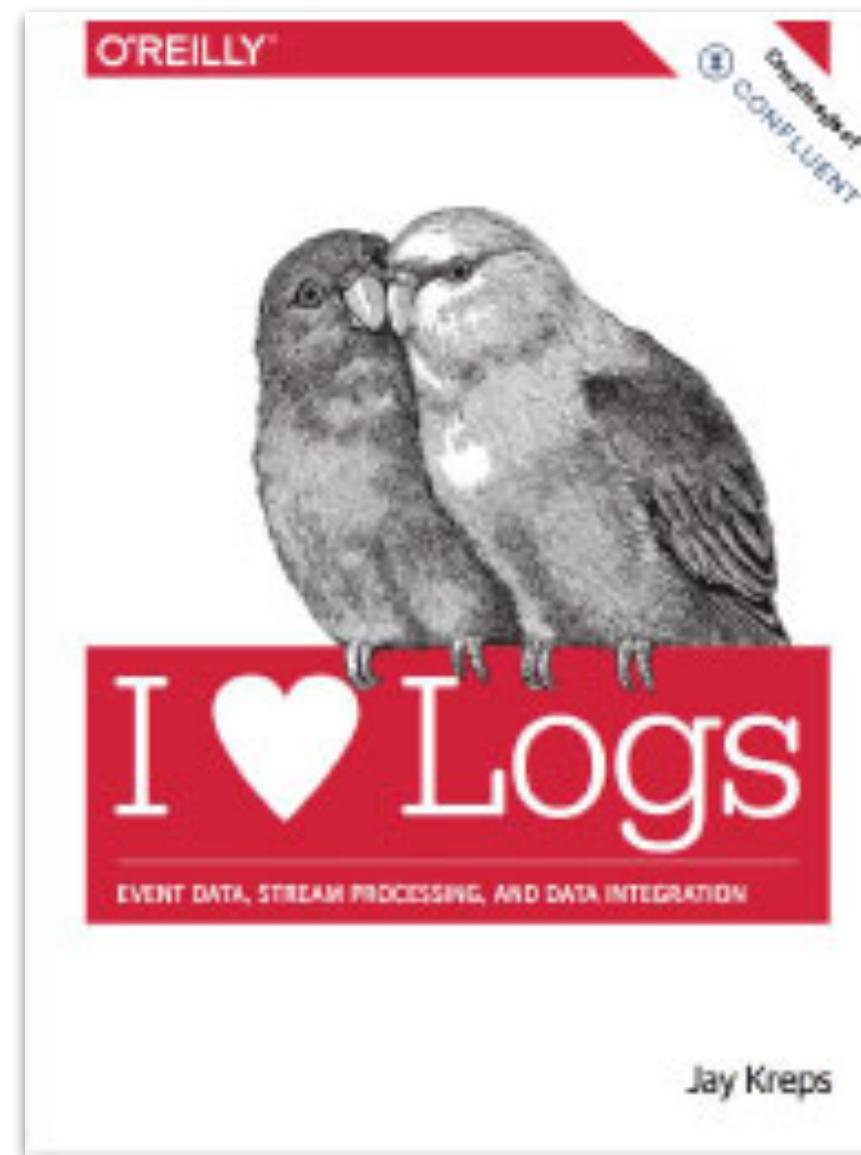
## ▼ Python

CONFLUENT SUPPORTED

Hide internal topics Use consumer groups\*

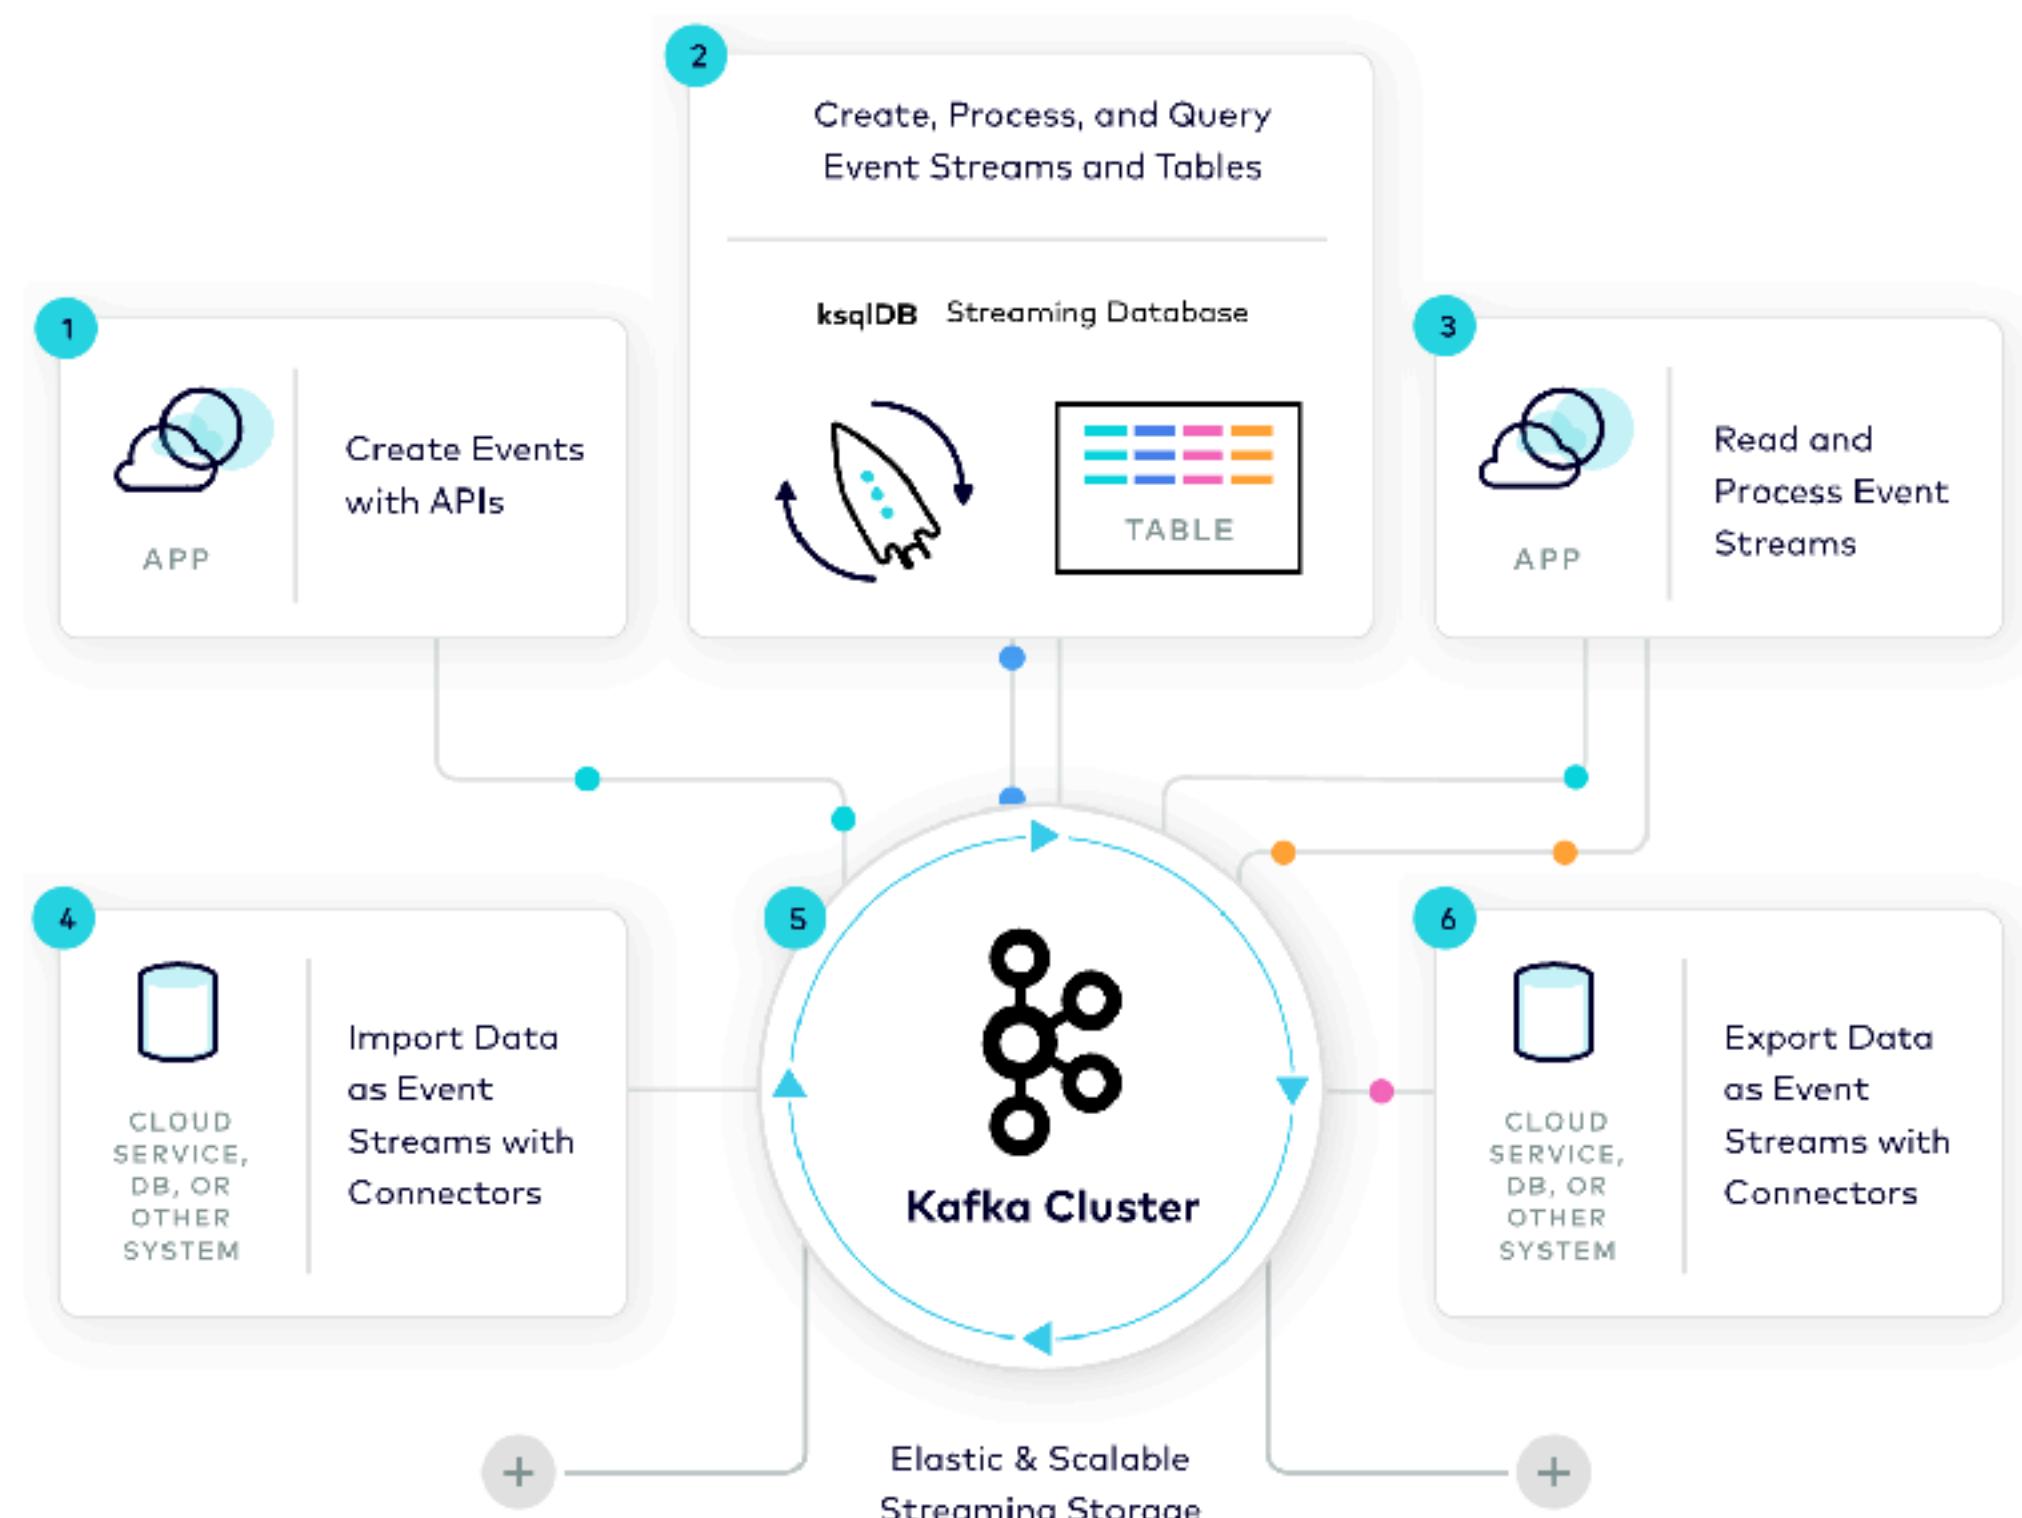
# Event Streaming Platform





# What Does Kafka Do?

Learn about the fundamentals of Kafka, event streaming, and the surrounding ecosystem. Click on an element to find out more.



## Event Streaming

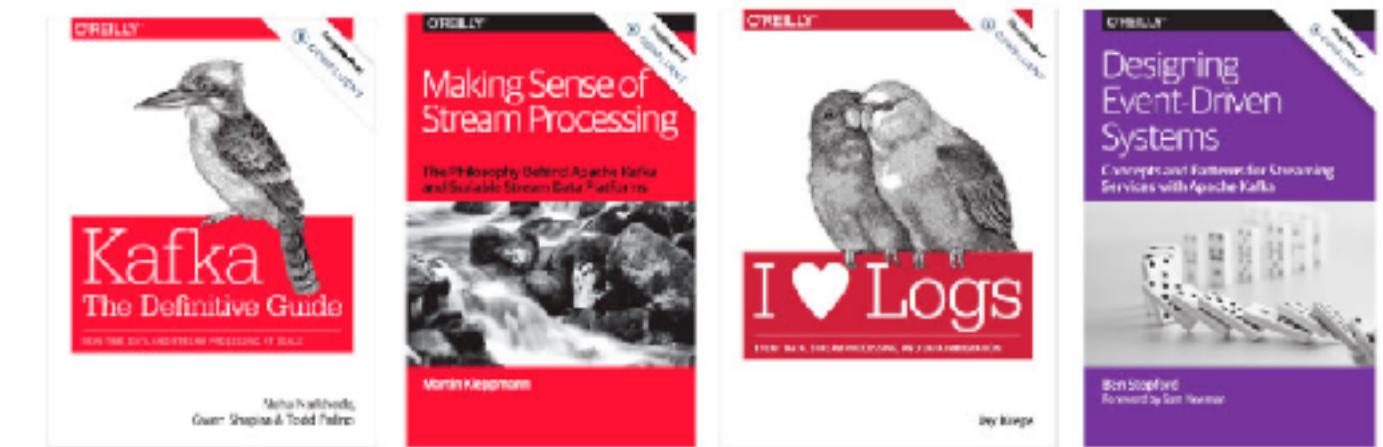
Create, import, share streams of events like payments, orders, and database changes in milliseconds, at scale. Store them as long as you need. Replay and reprocess historical data like a time machine.

## Continuous Real-time Processing

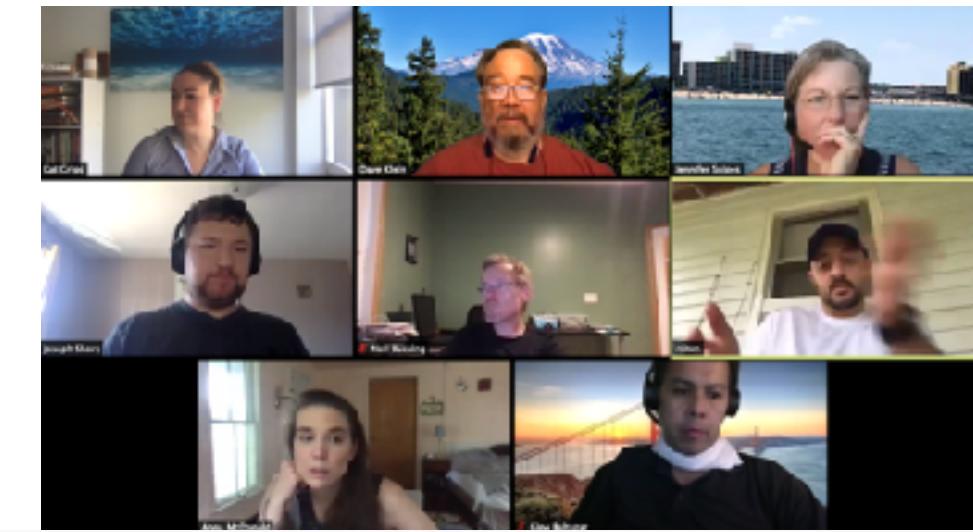
Build lightweight, elastic applications and microservices that respond immediately to events and that scale during live operations. Process, join, and analyze streams and tables of data in real-time, 24x7.

## Your Data, Everywhere

Use one platform to set data in motion across your entire enterprise. Connect systems, data centers, and clouds—all with the same trusted



<https://cnfl.io/book-bundle>



<https://cnfl.io/meetup-hub>



## Community

<https://cnfl.io/community>

[twitter.com/daveklein](https://twitter.com/daveklein)

[dklein@confluent.io](mailto:dklein@confluent.io)

<https://developer.confluent.io>



## Kafka 101



## Spring & Kafka



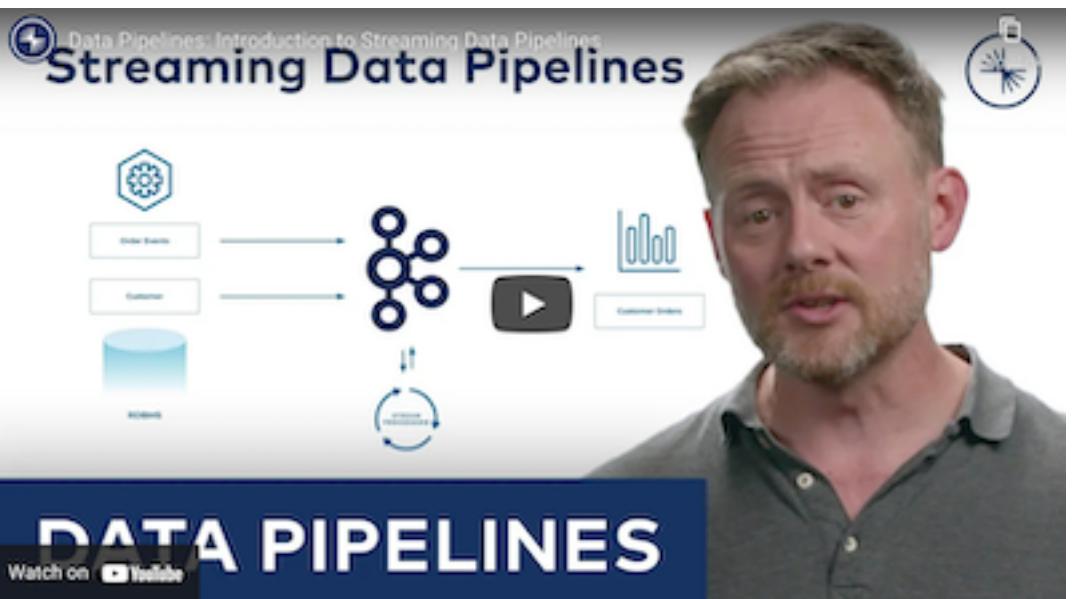
## Kafka Connect



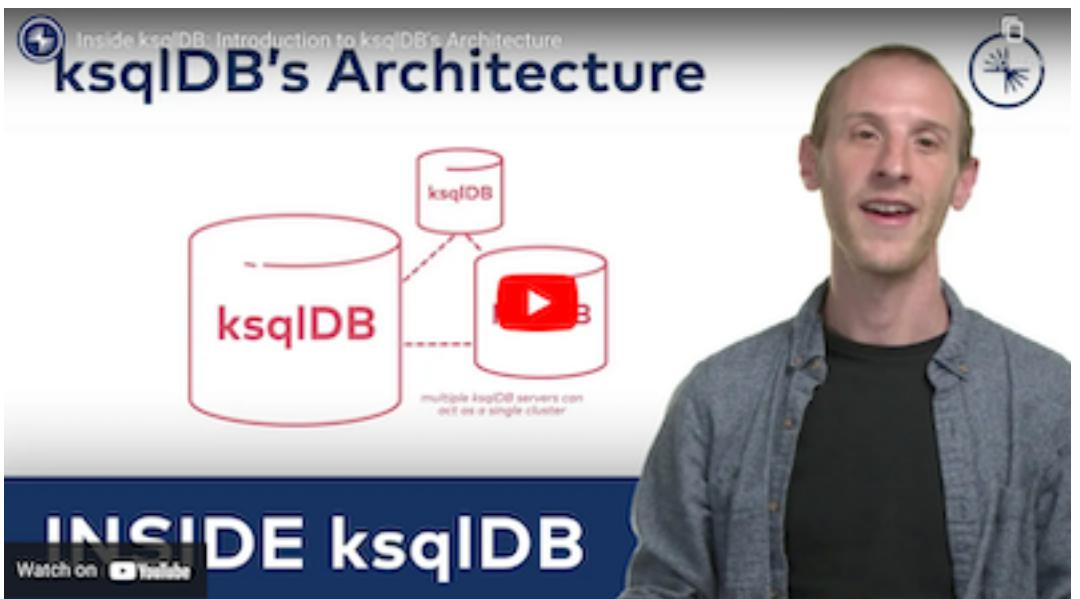
## Kafka Internals



## Kafka Streams 101



## Data Pipelines



## Inside ksqlDB



## Kafka Security



## ksqlDB 101



## Event Sourcing



## Data Mesh 101



## Schema Registry



Apache Kafka® 101: Introduction

# Events

- Internet of Things
- Business process change
- User interaction
- Microservice output



## KAFKA 101

Watch on YouTube

[Back to Courses](#)

COURSE: APACHE KAFKA® 101

### Introduction 5 MIN

**Tim Berglund**

Sr. Director, Developer Advocacy (Course Presenter)

## What Is Apache Kafka®?

Apache Kafka is an event streaming platform used to collect, process, store, and integrate data at scale. It has numerous use cases including distributed logging, stream processing, data integration, and pub/sub messaging.

In order to make complete sense of what Kafka does, we'll delve into what an "event streaming platform" is and how

### Modules 13 MODULES

1. Introduction	5 MIN
2. Hands On: Get Started with Kaf...	7 MIN
3. Topics	6 MIN
4. Partitioning	5 MIN
5. Brokers	2 MIN
6. Replication	2 MIN
7. Producers	5 MIN

### Related Resources

#### BLOG

[Intro to Apache Kafka: How Kafka Works](#)

#### TUTORIAL

[Apache Kafka Quick Start](#)

#### CONFERENCE TALK

[Kafka Summit: Welcome to Kafka; We're Glad You're Here](#)

#### E-BOOK

[Kafka: The Definitive Guide](#)